

IMAGE GRADIENTS AND EDGES



Class 3: Artificial Vision

Reminder: Linear filter

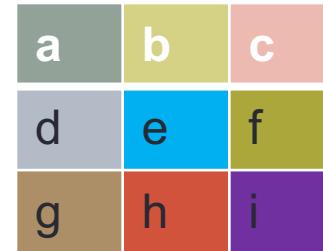
$$F[x, y]$$



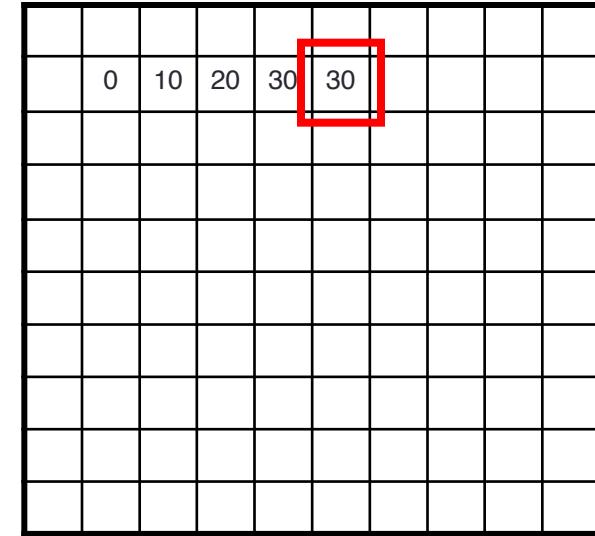
$$H[u, v]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



“box filter”



$$\begin{aligned}
 df[i, j] = & a * f[i - 1, j - 1] + b * f[i - 1, j] + c * f[i - 1, j + 1] + \\
 & d * f[i, j - 1] + e * f[i, j] + f * f[i, j + 1] + \\
 & g * f[i + 1, j - 1] + h * f[i + 1, j] + i * f[i + 1, j + 1]
 \end{aligned}$$

Last lecture

Answers next page

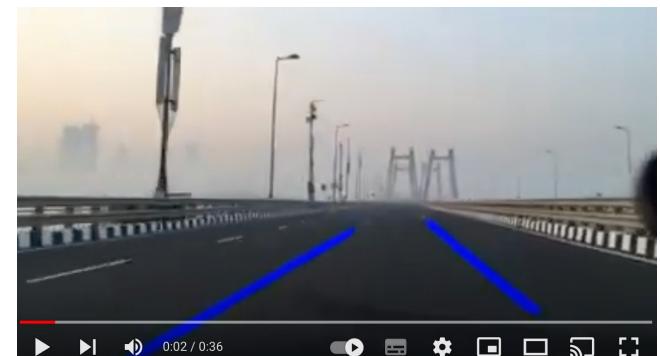
- Linear filters and convolution useful for
 - Image smoothing, removing noise
 - ...box.....filter $[1, 4, 16, \dots]$ → más importancia al pixel que estamos estudiando y a los vecinos.
 - Alternative (ponderate....) filter
 - Impact of scale / width (.....) of the smoothing filter
- Mean filter:
 - Smoothing pixel values by them
 - Uniform average vs. average
- Median filter:
 - a linear or non-linear filter, ?!
 - Does it assure no averaging grey levels?
 - Does it assure no inventing new grey-levels?
 - Is it edge-preserving?

Last lecture

- Linear filters and convolution useful for
 - Image smoothing, removing noise
 - Box filter
 - Alternative Gaussian filter
 - Impact of scale / width (sigma) of the smoothing filter
- Mean filter:
 - Smoothing pixel values by averaging them
 - Uniform average vs. weighted average
- Median filter:
 - a non-linear filter
 - it assures no averaging grey levels
 - it assures no inventing new grey-levels
 - It is edge-preserving

Today

- What is an image derivative and its relation to the edges
 - How to implement different edge detectors and what is their difference?
 - First and second derivatives, image gradient and Laplacian
- How to get image derivatives with a Gaussian function?
- What is a Canny edge detector?
- Application: hybrid images



Review



original image



filtered

Which filter is applied:

vertical \rightarrow $f = 1/9 \times [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]^T$

horizontal \rightarrow or $f = 1/9 \times [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]^T$!

horizontal haciendo suavización horizontal

Nos podemos salir del rango de la imagen

What happens if we have a smoothing filter that is *unnormalized* (does not sum to one)?

Recall: image filtering

- Compute a function of the local neighborhood at each pixel in the image
 - Function specified by a “filter” or mask saying how to combine values from neighbors.
- Uses of filtering:
 - Enhance an image (denoise, etc)
 - Extract information (edges, etc) –
 - Detect patterns (template matching) – to see later

Edge detection

- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.
- **Why?**

edge \Rightarrow cambio brusco de valores de gris
contorno \Rightarrow límites de los objetos.

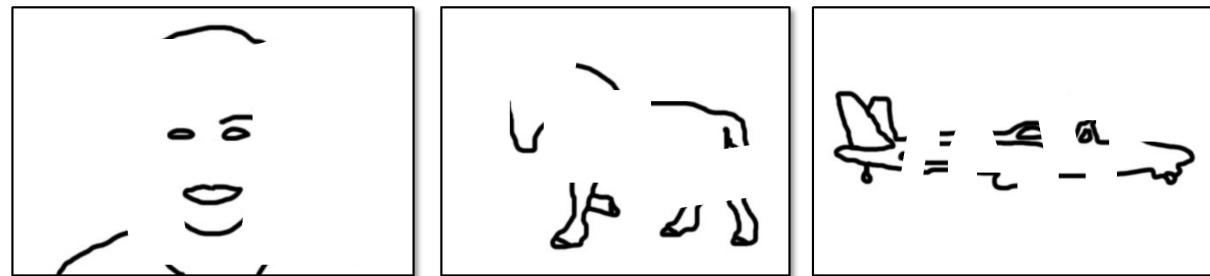


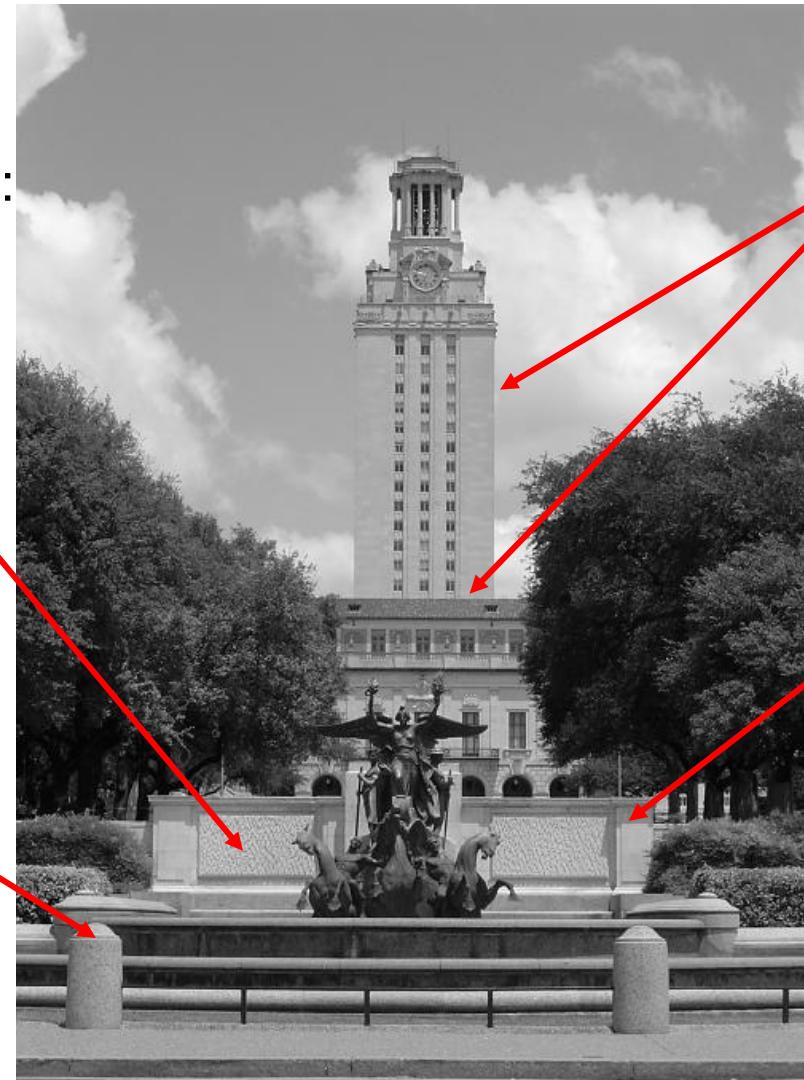
Figure from J. Shotton et al., PAMI 2007

- **Main idea:** look for strong gradients, post-process

What does cause an edge?

Reflectance change:
appearance
information, texture

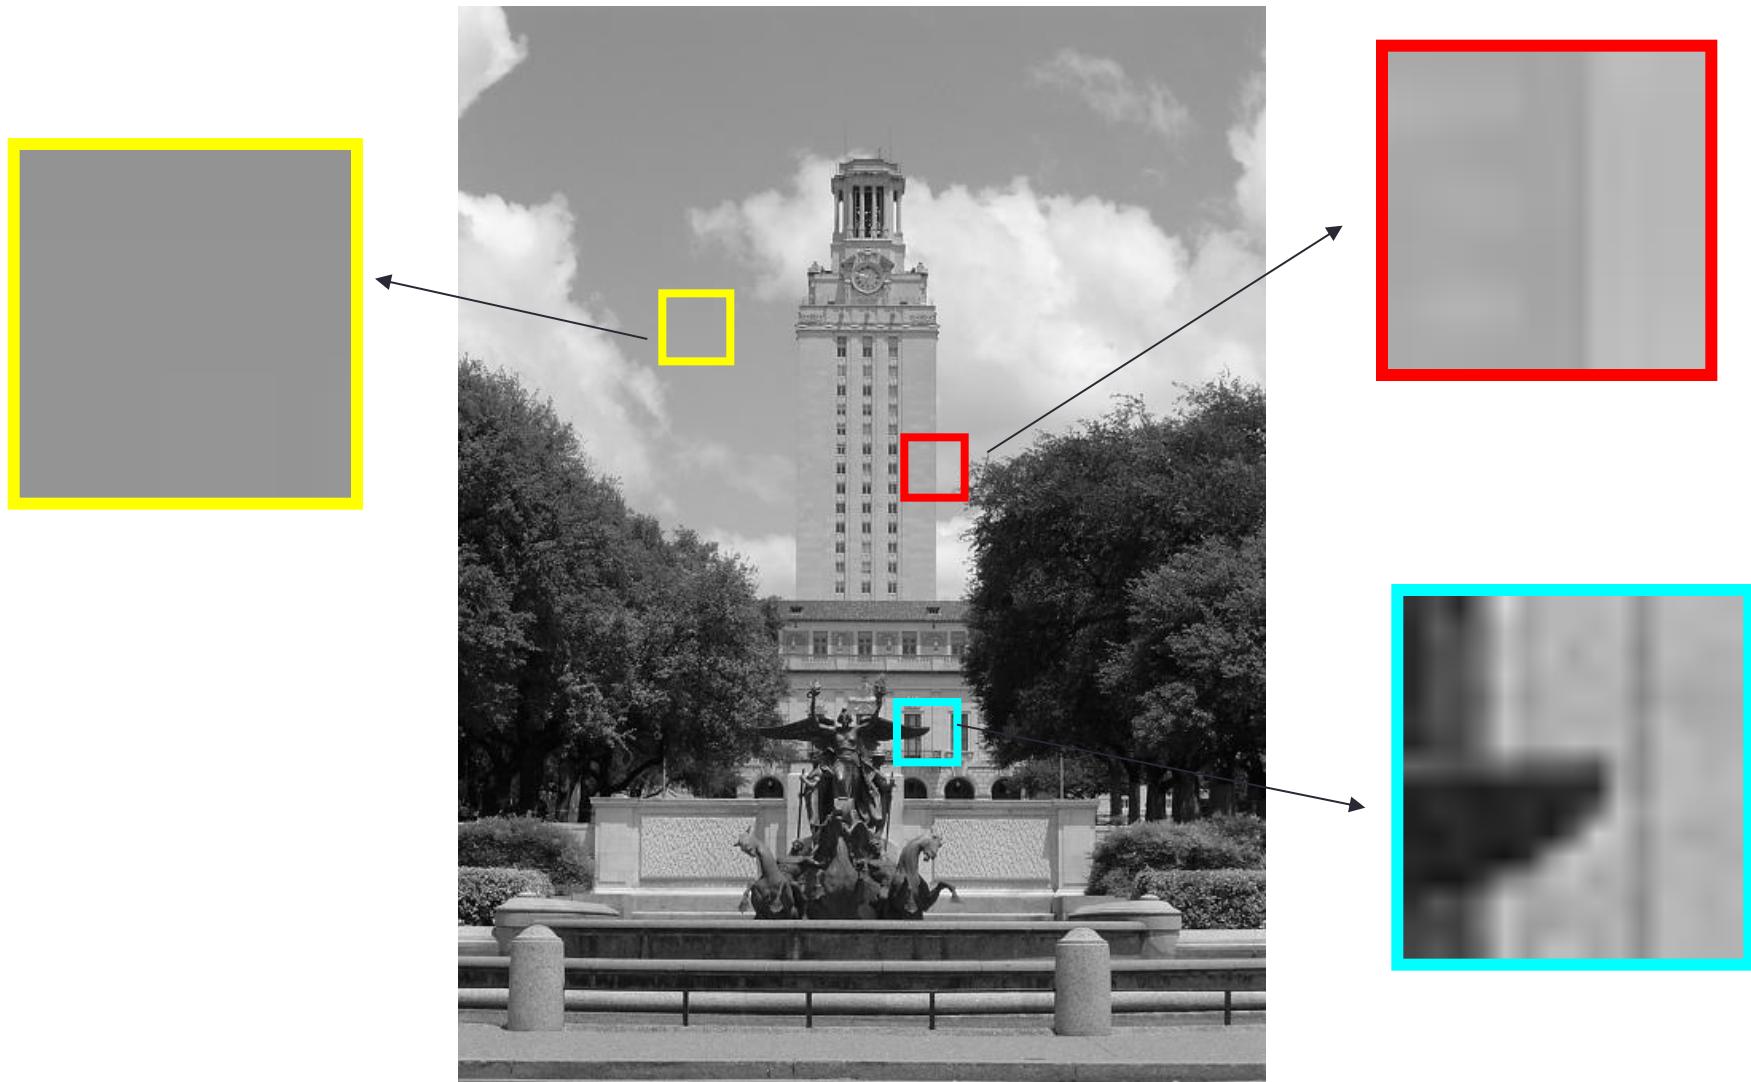
Change in surface
orientation: shape



Depth discontinuity:
object boundary

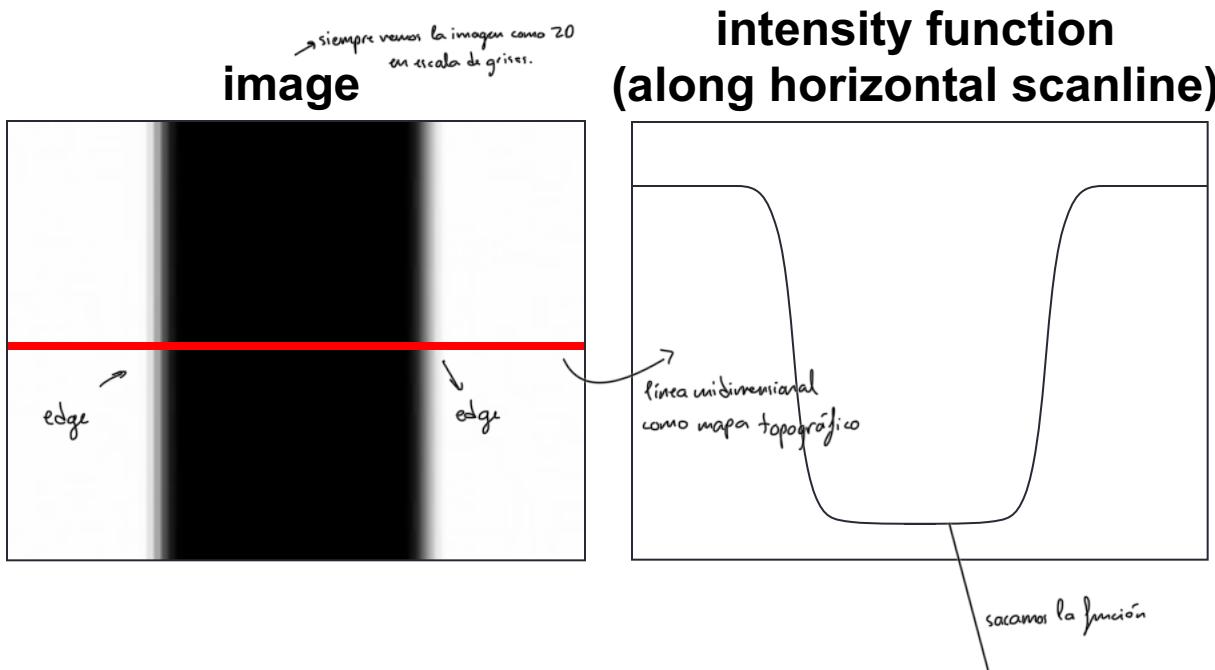
Cast shadows

Edges/gradients and invariance



Derivatives and edges

An edge is a place of rapid change in the image intensity function.



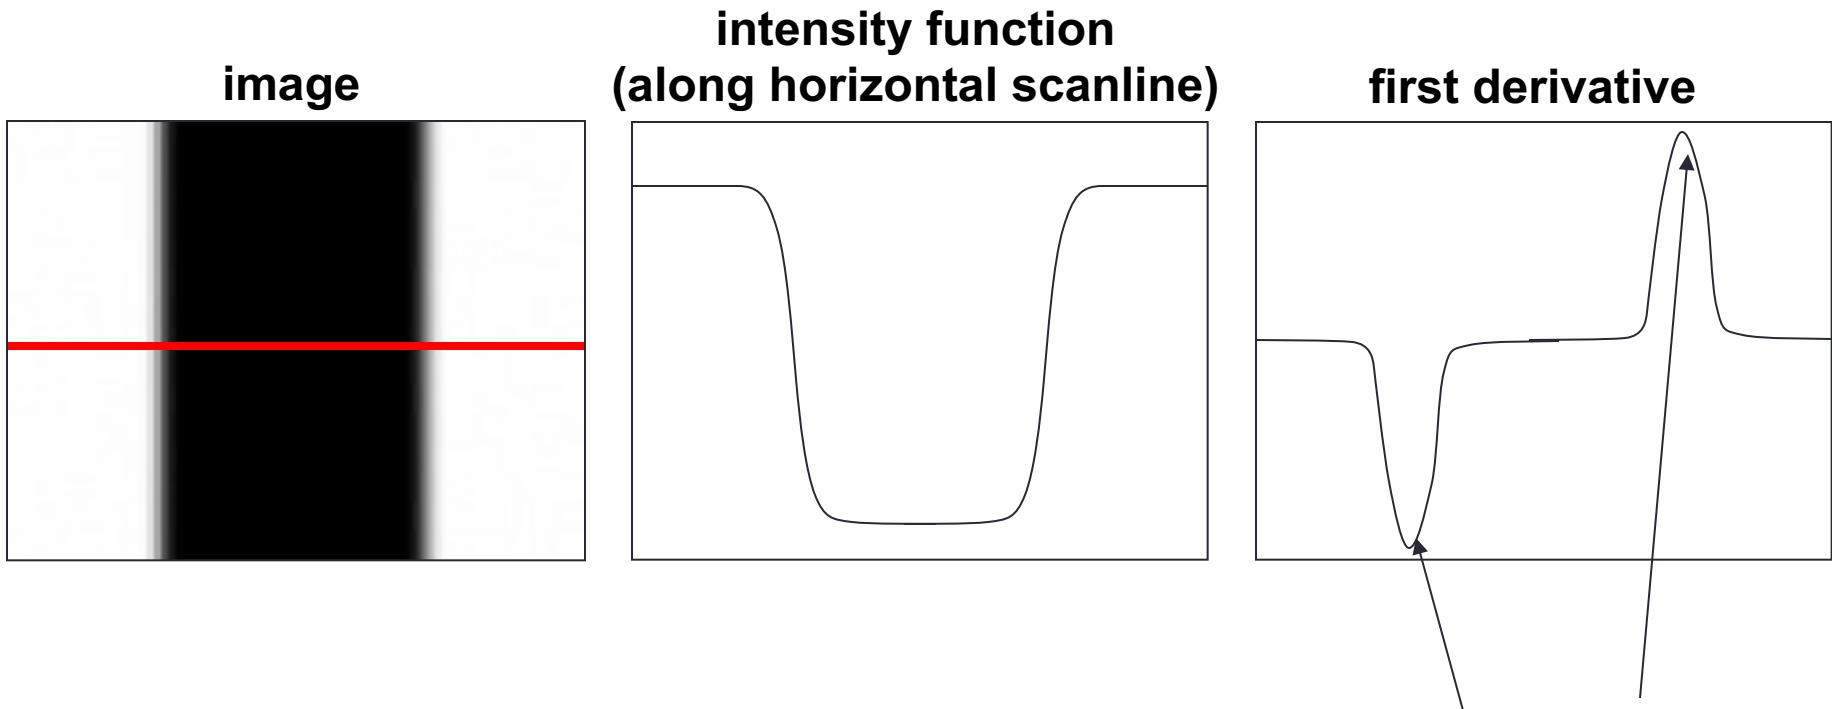
How to detect the points (pixels) of maximal change?

$$f'(a) = \lim_{h \rightarrow 0} \frac{f(a + h) - f(a)}{h}$$

siguiente anterior

Derivatives and edges

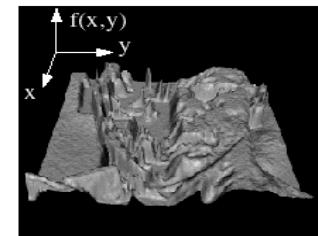
An edge is a place of rapid change in the image intensity function.



Edges correspond to extremes of derivative (max by absolute value)

Derivatives with convolution

But... the image is 2D!



For 2D function, $f(x,y)$, the partial derivative in x is:

$$\frac{\partial f(x, y)}{\partial x} \square \lim_{\varepsilon \rightarrow 0} \frac{f(x \square \varepsilon, y) - f(x, y)}{\varepsilon}$$

↓ diferencia entre 2 puntos
 ↓ pixel actual
 ↓ como mínimo 1 (diferencia de distancia entre 2
 pixels vecinos).

→ No tiene sentido porque estamos hablando de 3 colores diferentes. Calcular la derivada respecto a los canales no tiene sentido (podemos tener RGB en vez de RGB).

Does it make sense to get the derivative wrt 3rd dimension / channel? → RGB

Derivatives with convolution

But... the image is a discrete matrix!

$$\frac{\partial f(x, y)}{\partial x} \square \lim_{\varepsilon \rightarrow 0} \frac{f(x \square \varepsilon, y) - f(x, y)}{\varepsilon}$$

↓
como mínimo 1

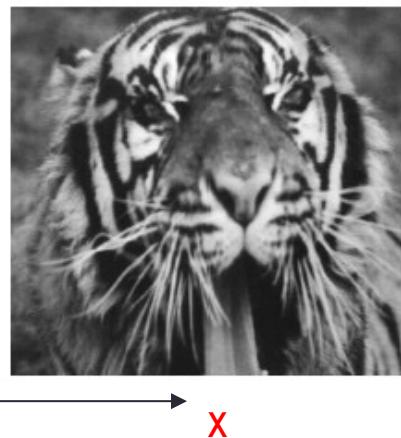
We can approximate it using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x \square \overset{+}{1}, y) - f(x, y)}{1}$$

→ derivada de la imagen

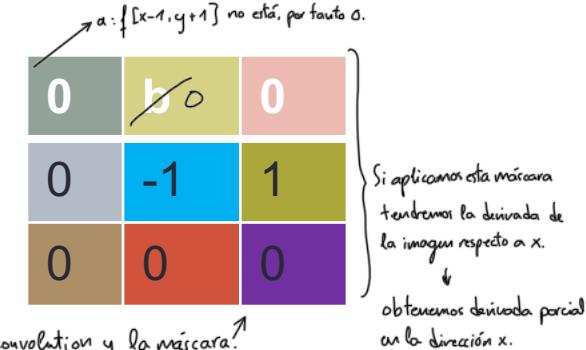
How to compute it on the image?

Partial derivatives of an image



What would be the associated convolution filter?

a	b	c
d	e	f
g	h	i



↓ queremos obtener este resultado mediante convolution y la máscara!

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

→ lo tiene la e, pero la tenemos en negativo

Recall:

→ Por cada pixel multiplicamos sus vecinos por $[a, b, c, d, e, f, g, h, i]$

$$df[x, y] = a * f[x-1, y+1] + b * f[x, y+1] + c * f[x+1, y+1] +$$

$$d * f[x-1, y] + e * f[x, y] + f * f[x+1, y] +$$

$$g * f[x-1, y-1] + h * f[x, y-1] + i * f[x+1, y-1]$$

Partial derivatives of an image



a	b	c
d	e	f
g	h	i

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

-1	1
----	---

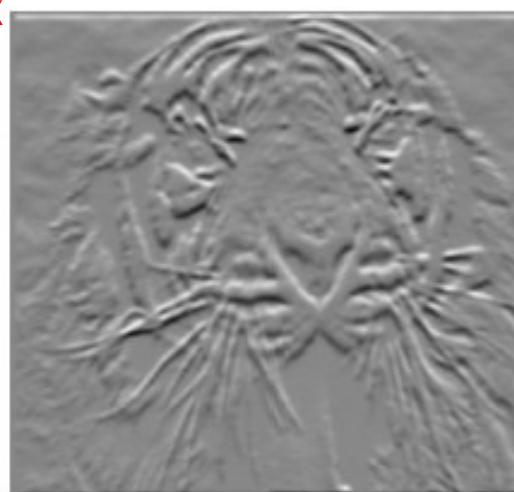
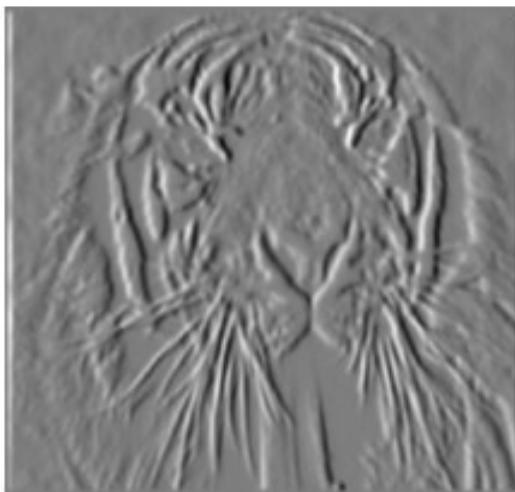
Partial derivatives of an image



$$\text{respecto } y \frac{\delta f(x,y)}{\delta y} = \frac{f(x,y+1) - f(x,y)}{1}$$

$$\frac{\partial f(x, y)}{\partial x}$$

-1 1



o desfoco contornos verticales o horizontales

Which derivative does show changes with respect to x?

First derivatives: discrete operators.

Given that the function f to derive our image I is:

$$\frac{\partial f}{\partial y} = I_{i+1,j} - I_{i,j}$$


$$\frac{\partial f}{\partial x} = I_{i,j+1} - I_{i,j}$$


Masks

$I_{i+1,j-1}$	$I_{i+1,j}$	$I_{i+1,j+1}$
$I_{i,j-1}$	$I_{i,j}$	$I_{i,j+1}$
$I_{i-1,j-1}$	$I_{i-1,j}$	$I_{i-1,j+1}$

y ↑ x →

Exercise: Which derivative will give maximal values on the image?

→ derivada horizontal → resalta contornos verticales
y viceversa

0 10 10
0 10 10
0 10 10

First derivatives: discrete operators.

Given that the function f to derive is our image I :

$$\frac{\partial f}{\partial y} = I_{i+1,j} - I_{i,j}$$

$$\frac{\partial f}{\partial x} = I_{i,j+1} - I_{i,j}$$

Hacemos el promedio de las derivadas para suavizar.

$I_{i+1,j-1}$	$I_{i+1,j}$	$I_{i+1,j+1}$
$I_{i,j-1}$	$I_{i,j}$	$I_{i,j+1}$
$I_{i-1,j-1}$	$I_{i-1,j}$	$I_{i-1,j+1}$



What happens when there is noise in the image?

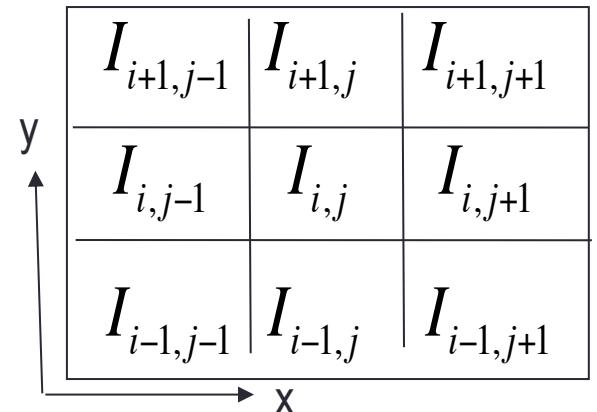
$$\begin{matrix}
 0 & 0 & 0 \\
 & \diagdown & \\
 0 & 25 & 0 \\
 & \diagup & \\
 0 & 0 & 0
 \end{matrix}$$

valores muy diferentes, seguramente sea ruido y no lo queremos.

First derivatives: discrete operators.

How to be less sensitive to noise?

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 25 & 0 \\ 0 & 0 & 0 \end{matrix}$$



Alternative: getting the average of the 3 derivatives in order to be less sensitive to noise:

$$\frac{\partial I}{\partial y} = (I_{i+1,j+1} - I_{i,j+1}) + (I_{i+1,j} - I_{i,j}) + (I_{i+1,j-1} - I_{i,j-1})$$

$$\frac{\partial I}{\partial x} = (I_{i+1,j+1} - I_{i+1,j}) + (I_{i,j+1} - I_{i,j}) + (I_{i-1,j+1} - I_{i-1,j})$$

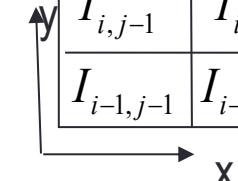
First derivatives: discrete operators.

If we consider symmetric finite differences:

$$\frac{\partial I}{\partial y} = (I_{i+1,j+1} - I_{i-1,j+1}) + (I_{i+1,j} - I_{i-1,j}) + (I_{i+1,j-1} - I_{i-1,j-1})$$

$$\frac{\partial I}{\partial x} = (I_{i+1,j+1} - I_{i+1,j-1}) + (I_{i,j+1} - I_{i,j-1}) + (I_{i-1,j+1} - I_{i-1,j-1})$$

$I_{i-1,j-1}$	$I_{i,j-1}$	$I_{i+1,j-1}$
$I_{i,j-1}$	$I_{i,j}$	$I_{i+1,j-1}$
$I_{i-1,j-1}$	$I_{i-1,j}$	$I_{i-1,j+1}$



What would be the convolution mask to compute the derivatives?

$M_x =$

-1	0	1
-1	0	1
-1	0	1

misma máscara que antes pero mejorada. Entrada en el pixel y haciendo el promedio con los pixeles vecinos.

$M_y =$

1	1	1
0	0	0
-1	-1	-1

Prewitt masks

Assorted finite difference filters

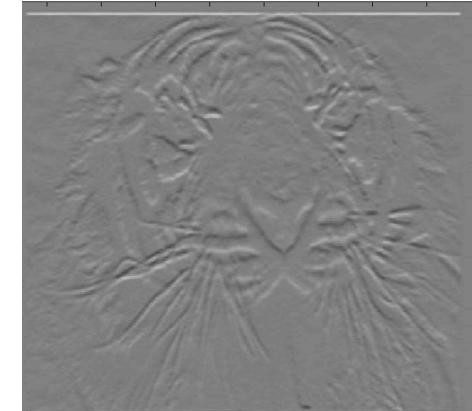
Valores ponderados para dar más prioridad al pixel.

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$	$; \quad M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$
Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$	$; \quad M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$
	→ Funciona mejor porque da más prioridad al pixel y menos al vecino.
	→ resalta contornos horizontales

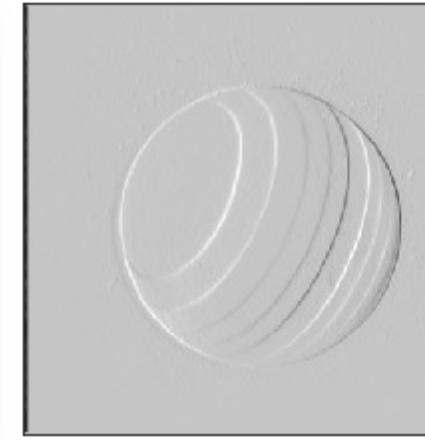
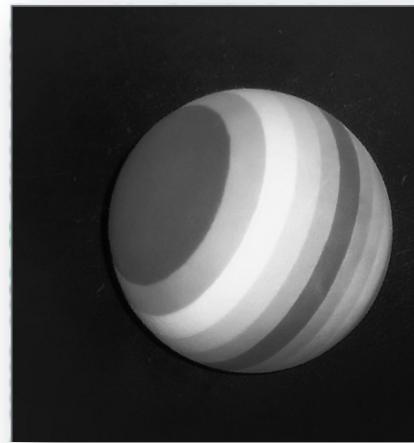
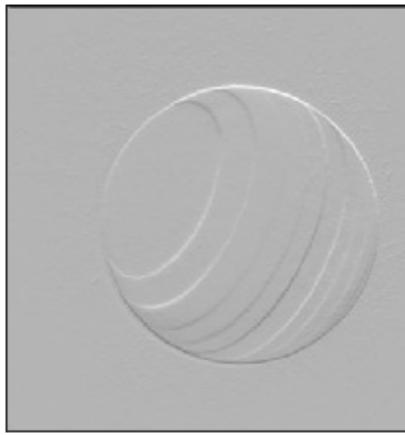
What is the difference between the Prewitt and Sobel operators?

- Sobel, by using non-uniform weights, gives more importance to the closest neighbors.

Which do you expect to be better?



Assorted finite difference filters



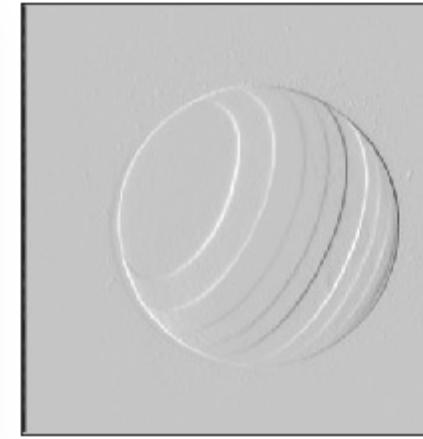
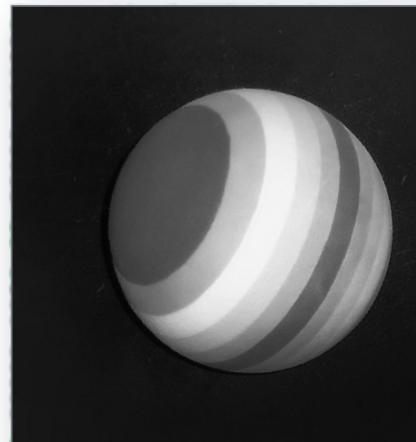
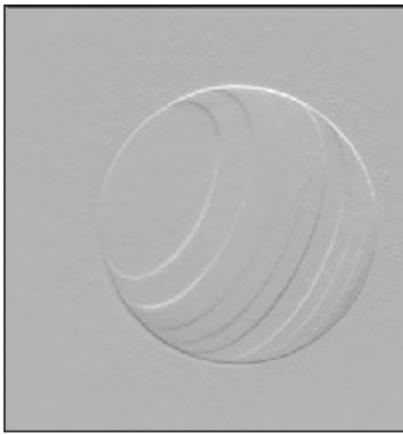
Determine to which Sobel masks (wrt x and y) do these images correspond to?!

How to apply Sobel in skimage to detect the image edges?

Skimage:

```
>>from skimage import filters  
  
>>im_vh=filters.sobel_v(im)  
# Image should be float and gray value!
```

Assorted finite difference filters



Alternative:

```
>>from scipy.ndimage import  
>>convolveim_sv=convolve(im,mask)
```

How to combine both derivatives (in x and y)?

```
>>im_s=filters.sobel(im)
```

↓ ya está implementado en sobel.

↓ la combinación horizontal, vertical, diagonal...

Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

vector compuesto por la derivada parcial en x y la derivada parcial en y.

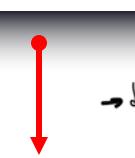
The gradient points in the direction of most rapid change in intensity



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

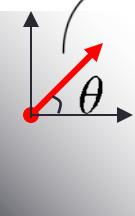
El gradiante tiene siempre una dirección (la de máximo cambio).
orientado a esa dirección

máximo cambio en la dirección horizontal



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

dirección vertical
derivada horizontal=0



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

cambio en las dos direcciones.

Image gradient

The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

↓
es un vector

mayor contraste \Rightarrow mayor magnitud de gradiente

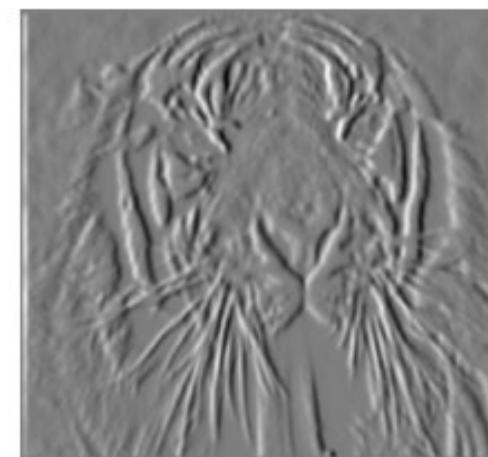


Image gradient

The **gradient direction** (orientation of edge normal) is given by:

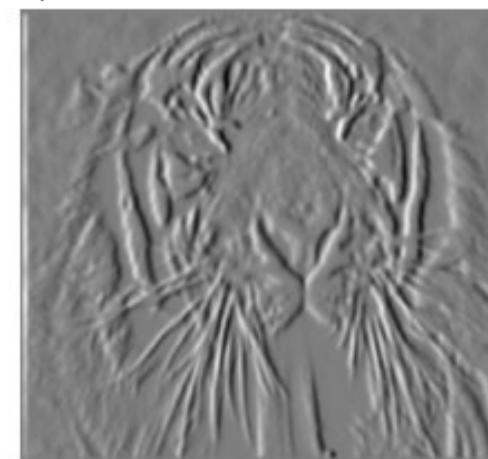
$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

arctan *división*

Which aspect (magnitude or direction) is invariant to image contrast?

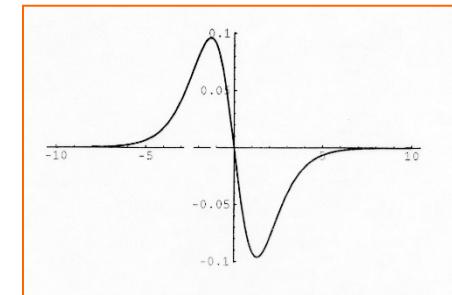
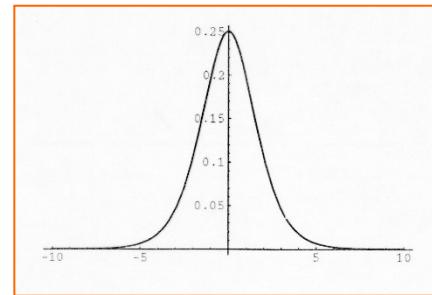
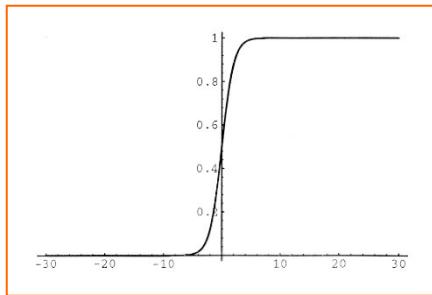
Si aumentamos el contraste \Rightarrow la magnitud del gradiente será mayor
Si oscurezco las zonas claras, más claras; y las zonas oscuras más oscuras no cambia la dirección del gradiente.

} cambiamos magnitud pero no dirección.



Discrete operators: second derivatives

What do you expect about the second derivatives?

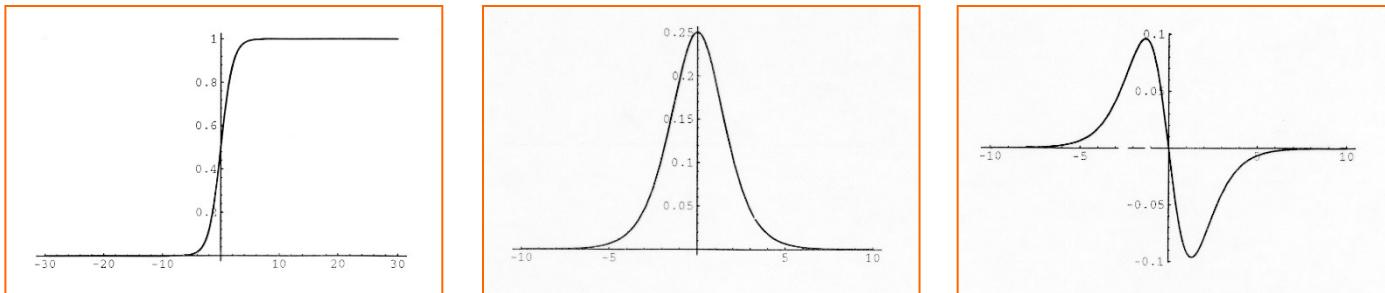


$$f'(a) = \lim_{h \rightarrow 0} \frac{f(a + h) - f(a)}{h}$$

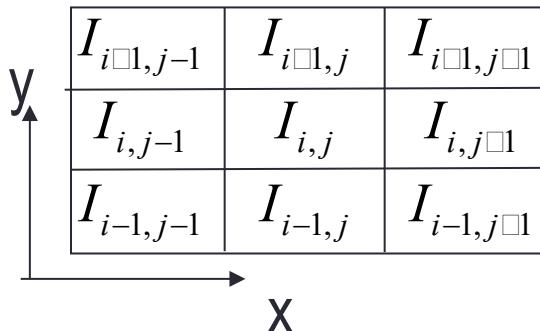
Given an edge, the first derivative has an extreme, and the second one has a zero-crossing in the edge and extremes before and after the edge.

Discrete operators: second derivatives

What do you expect about the second derivatives?



Given an edge, the first derivative has an extreme, and the second one has a zero-crossing in the edge and extremes before and after the edge.



The second derivative is the derivative of the first derivative:

$$\frac{\partial^2 I}{\partial y^2} = (I_{i+1,j} - I_{i,j}) - (I_{i,j} - I_{i-1,j}) = (I_{i-1,j} - 2I_{i,j} + I_{i+1,j})$$

$$\frac{\partial^2 I}{\partial x^2} = (I_{i,j-1} - 2I_{i,j} + I_{i,j+1}) \longrightarrow$$

1 -2 1

Discrete operators: Laplacian

Let's define:

$$\Delta I(x, y) = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2} = (I_{i-1,j} + I_{i,j-1} + I_{i+1,j} + I_{i,j+1}) - 4I_{i,j}$$



y ↑

$I_{i-1,j-1}$	$I_{i-1,j}$	$I_{i-1,j+1}$
$I_{i,j-1}$	$I_{i,j}$	$I_{i,j+1}$
$I_{i-1,j-1}$	$I_{i-1,j}$	$I_{i-1,j+1}$

x →

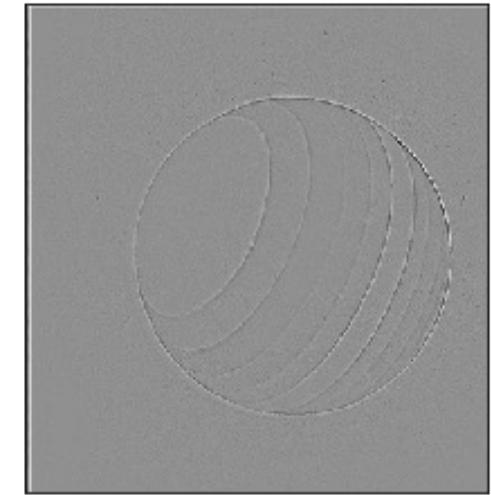
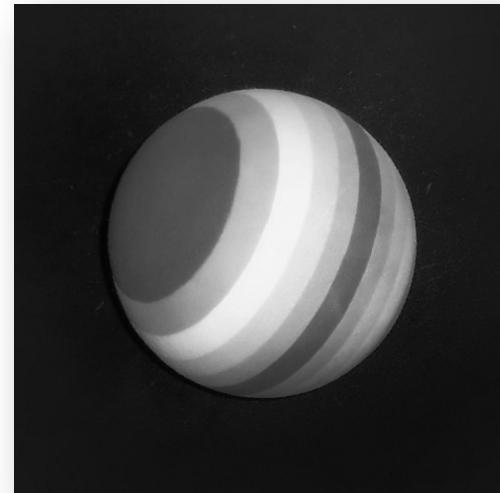
Laplacian mask:

a	b	c	0	1	0
d	e	f	1	-4	1
g	h	i	0	1	0

Discrete operators: Laplacian

Laplacian mask:

$$\begin{matrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{matrix}$$

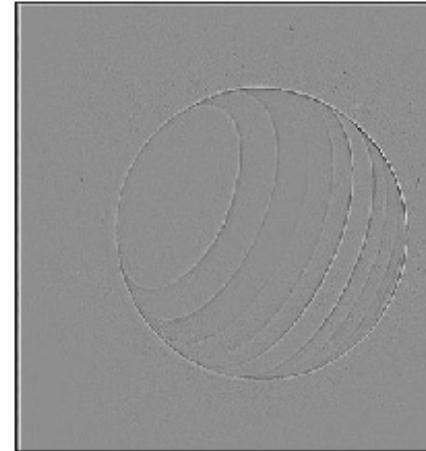
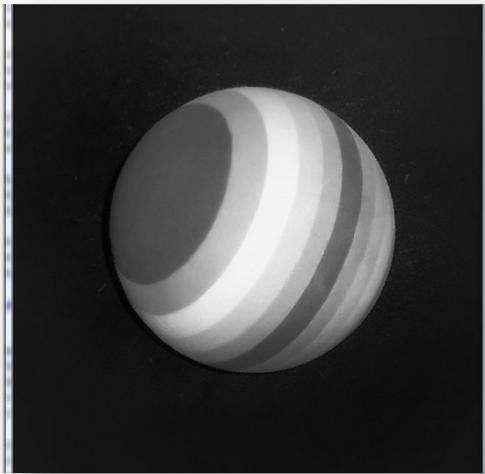


Skimage:

```
>> mask=numpy.array([[0,1,0],[1,-4,1],[0,1,0]], dtype='float')  
  
>> im_11= scipy.ndimage.convolve(skimage.img_as_float(im),mask)
```

Discrete operators: Laplacian

Including the diagonal neighbours:



Laplacian:

$$\begin{matrix} 1 & 4 & 1 \\ 4 & -20 & 4 \\ 1 & 4 & 1 \end{matrix}$$

Mask properties

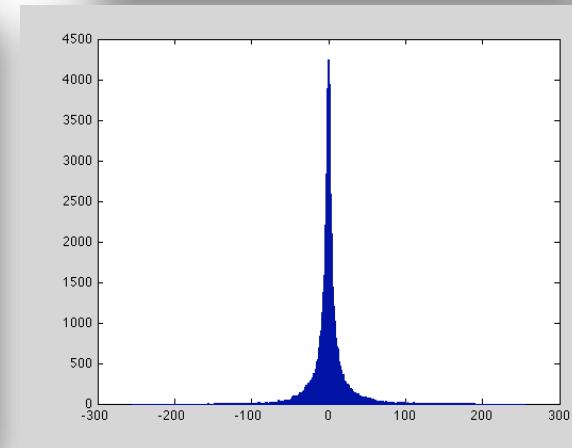
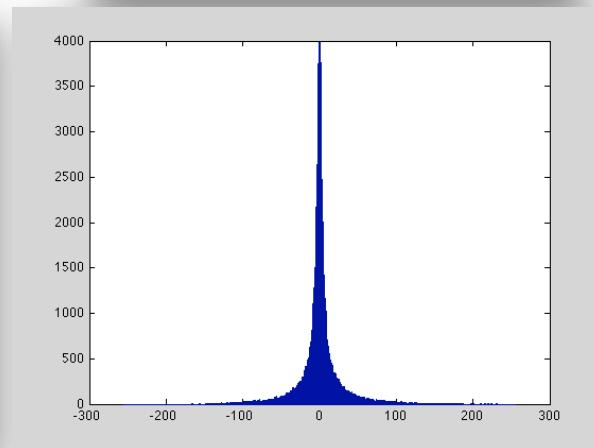
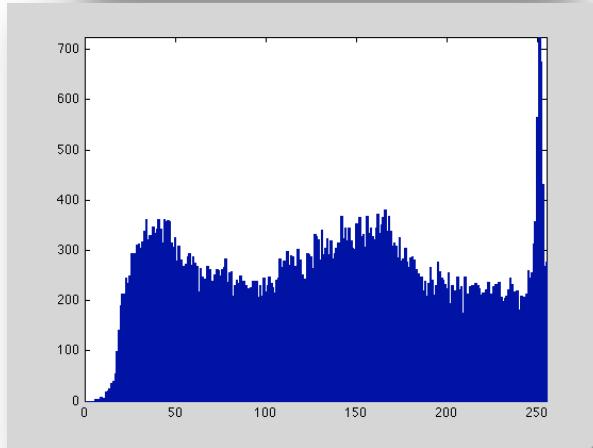
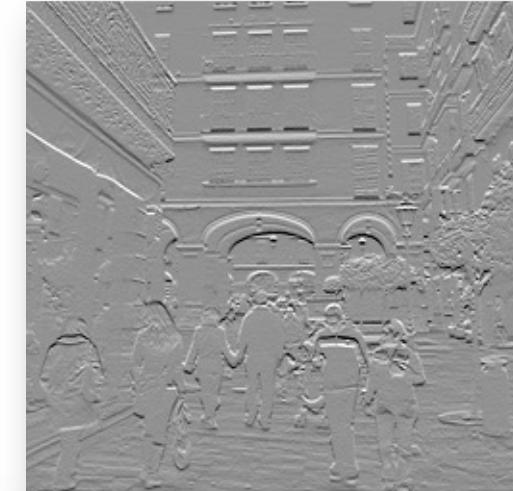
- Test on smoothing

- Values should be _____
- Sum to ___ → constant regions same as input
- Amount of smoothing _____ to mask size
- Remove “_____ -frequency” components; “_____ -pass” filter

- Test on derivatives

- _____ signs used to get high response in regions of high contrast
- Sum to ___ → no response in constant regions
- _____ value at points of high contrast

Explain the histogram of x and y edge maps



Horizontal and vertical derivatives distribution: why are they positive and negative?
Why the maximum is in 0? Can I store them in a uint8 type image?.

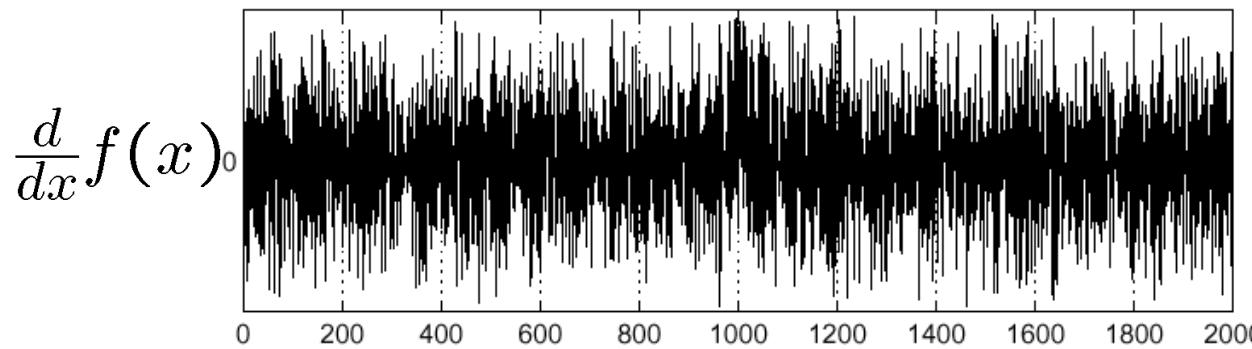
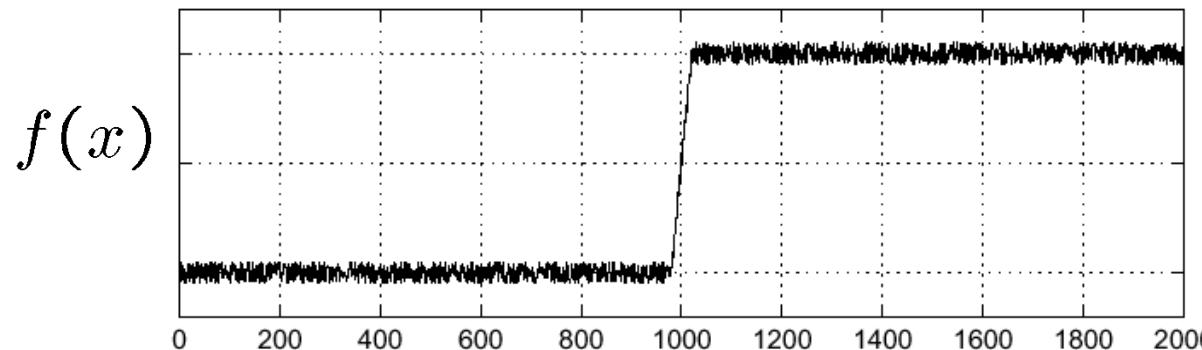
Today

- What is an image derivative and its relation to the edges
 - How to implement different edge detectors and what is their difference?
 - First and second derivatives, image gradient and Laplacian
- How to get image derivatives with a Gaussian
- Canny edge detector
- Application: hybrid images

Effects of noise

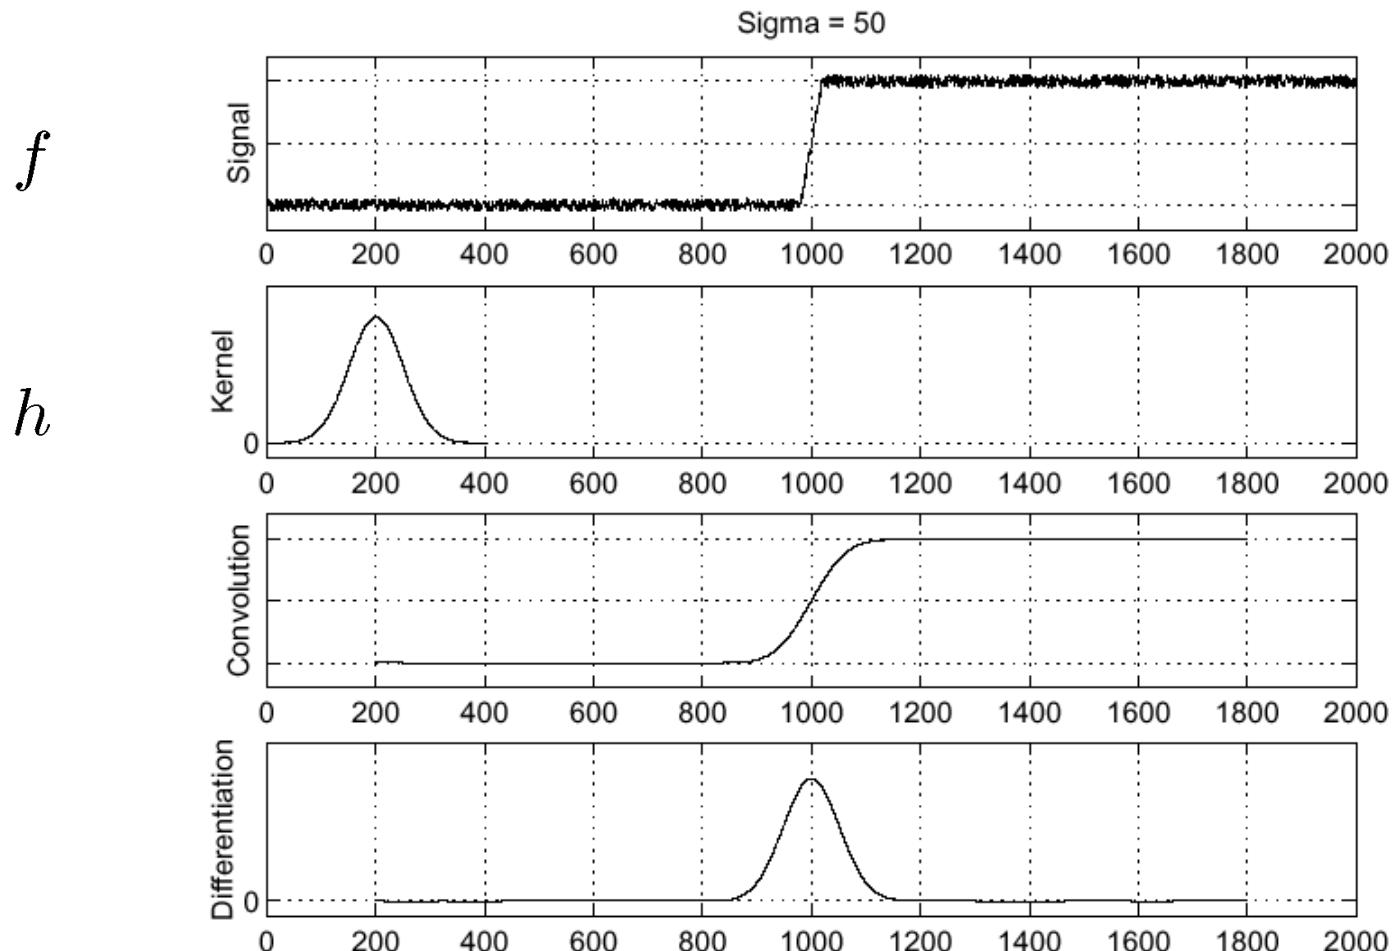
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

Solution: smooth first



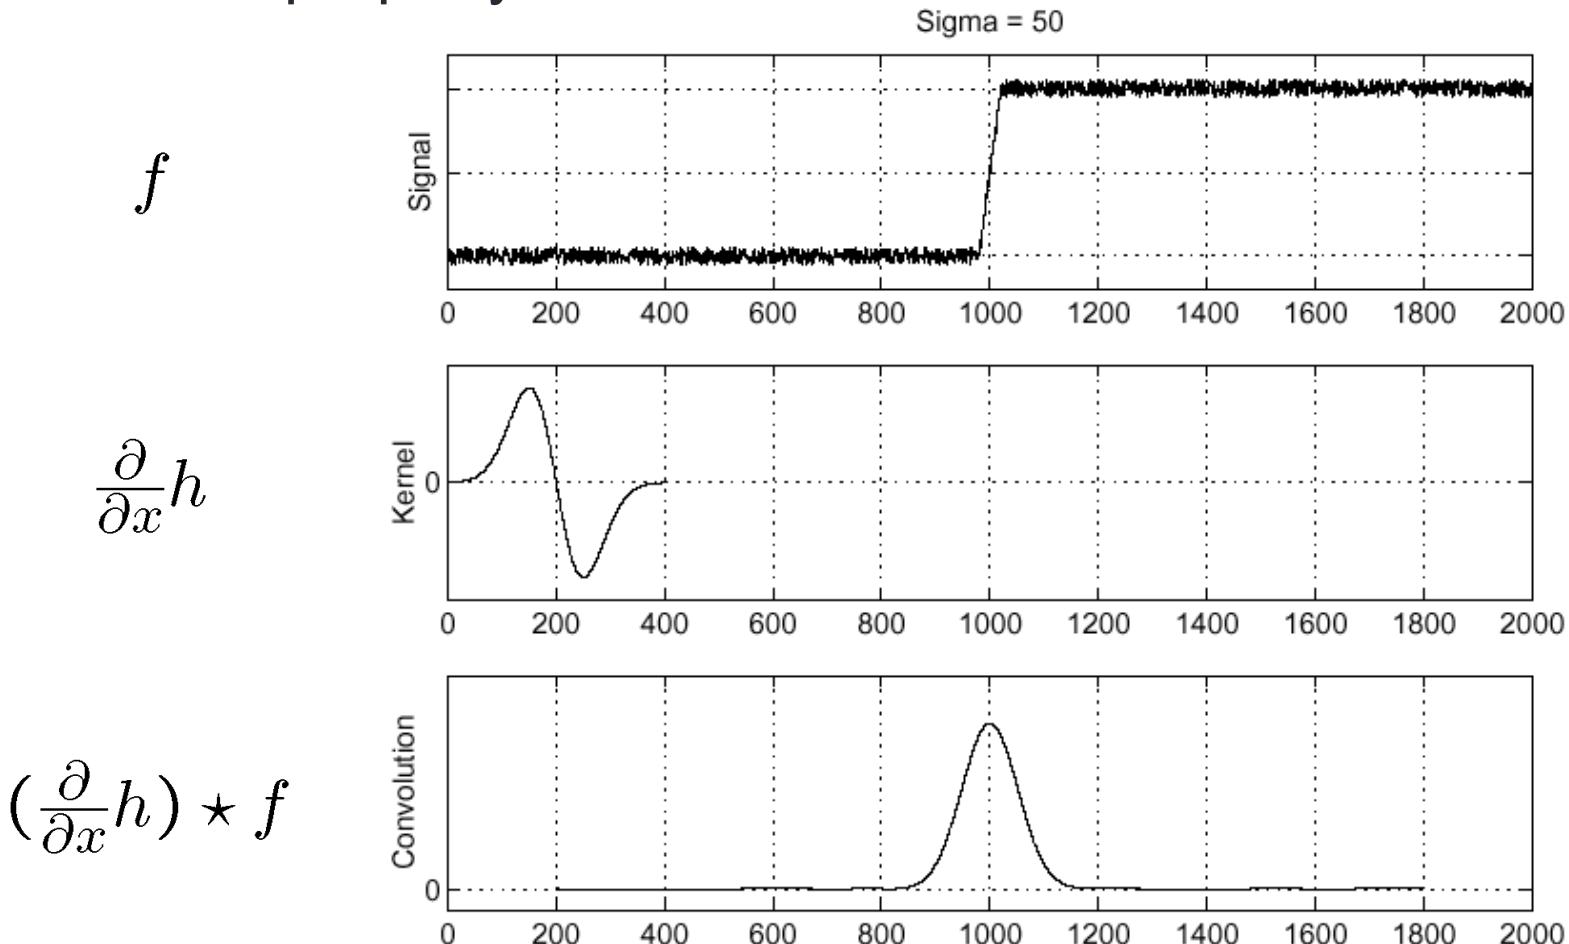
Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

Differentiation property of convolution.



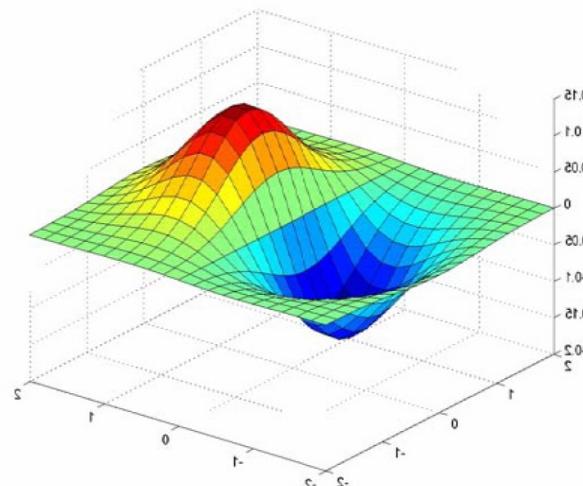
Derivative of Gaussian filters

$$(I \otimes g) \otimes h \quad \square \quad I \otimes (g \otimes h)$$

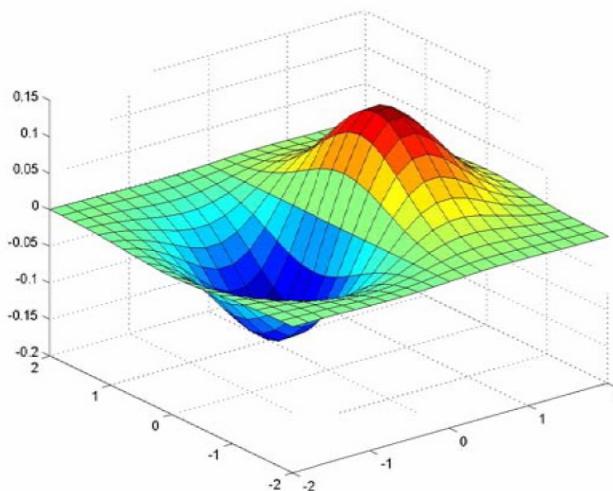
$$\begin{bmatrix} & & & & \\ & \otimes & & & \\ & & 1 & -1 & \end{bmatrix}$$

0.0030	0.0133	0.0219	0.0133	0.0030
0.0133	0.0596	0.0983	0.0596	0.0133
0.0219	0.0983	0.1621	0.0983	0.0219
0.0133	0.0596	0.0983	0.0596	0.0133
-0.0030	0.0133	0.0219	0.0133	0.0030

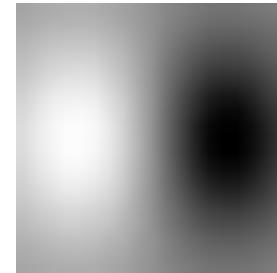
$$\begin{bmatrix} & & & & \\ & \otimes & & & \\ & & 1 & -1 & \end{bmatrix}$$



Derivative of Gaussian filters

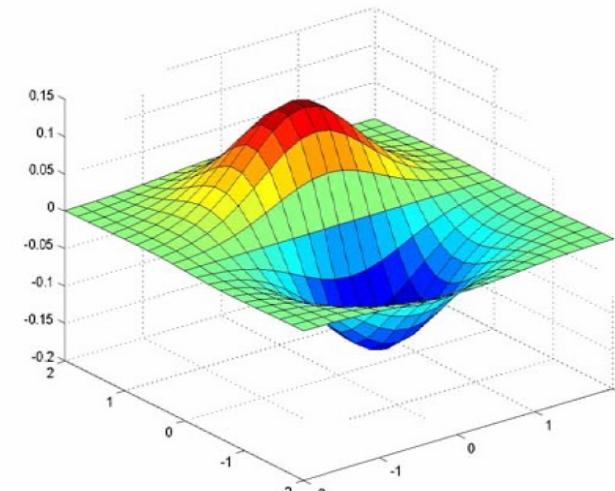


x-direction

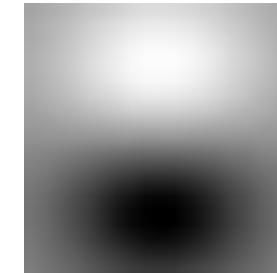


Gaussian

$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



y-direction



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Implementation in Skimage

Gaussian

$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

In skimage:

```
>>im_gaus_der11=scipy.ndimage.filters.gaussian_filter(im, 1, order=[1,1])  
  
>>im_gaus_der10=scipy.ndimage.filters.gaussian_filter(im, 1, order=[1,0])  
  
>>im_gaus_der01=scipy.ndimage.filters.gaussian_filter(im, 1, order=[0,1])
```

It can be shown that **convolving with a 2D Gaussian** is the same that **convolving with a 1D Gaussian** in x direction followed by convolving in y direction!

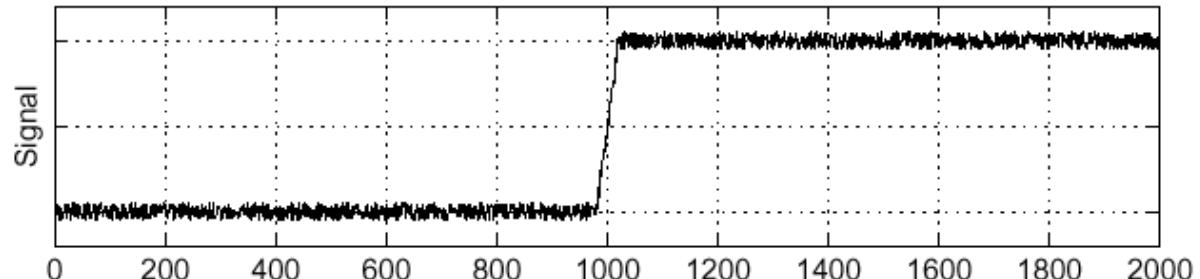
Laplacian of Gaussian

Consider

$$\frac{\partial^2}{\partial x^2}(h \star f)$$

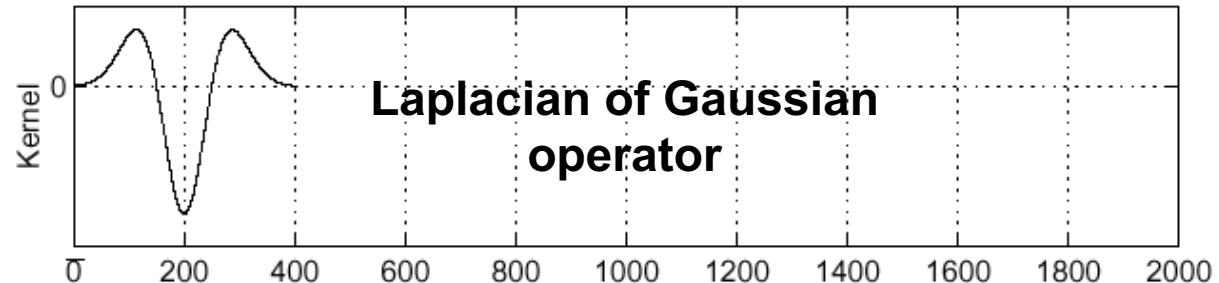
f

Sigma = 50

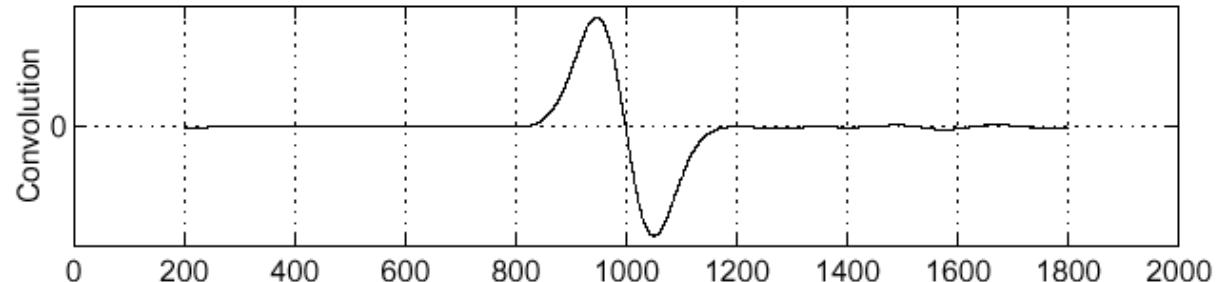


$$\frac{\partial^2}{\partial x^2} h$$

Laplacian of Gaussian operator



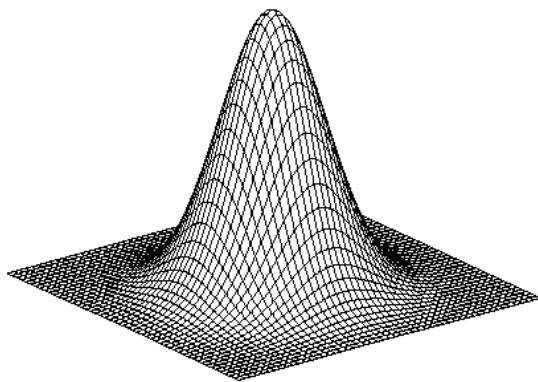
$$\left(\frac{\partial^2}{\partial x^2} h\right) \star f$$



Where is the edge?

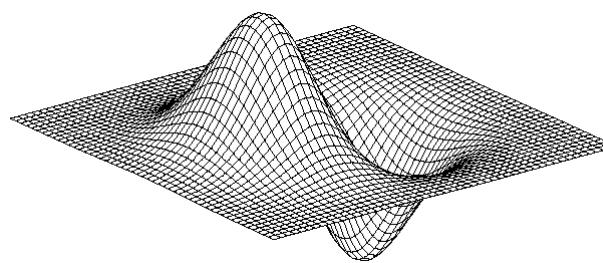
Zero-crossings of bottom graph

2D edge detection filters



Gaussian

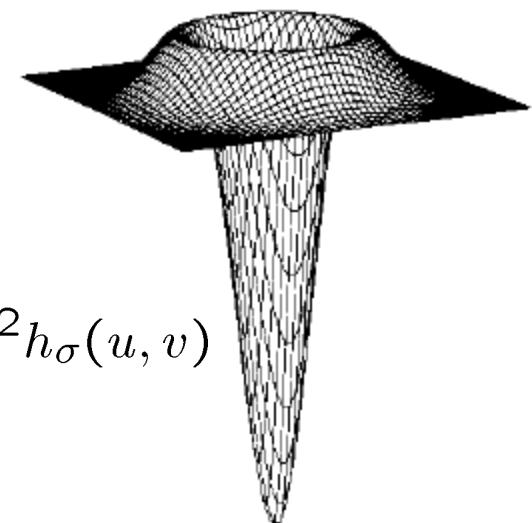
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



- ∇^2 is the Laplacian operator:

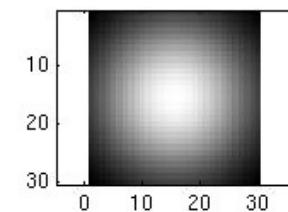
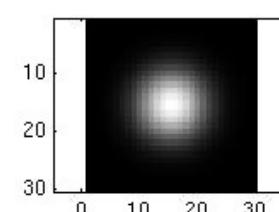
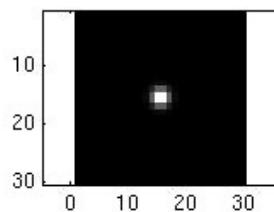
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Smoothing with a Gaussian

Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



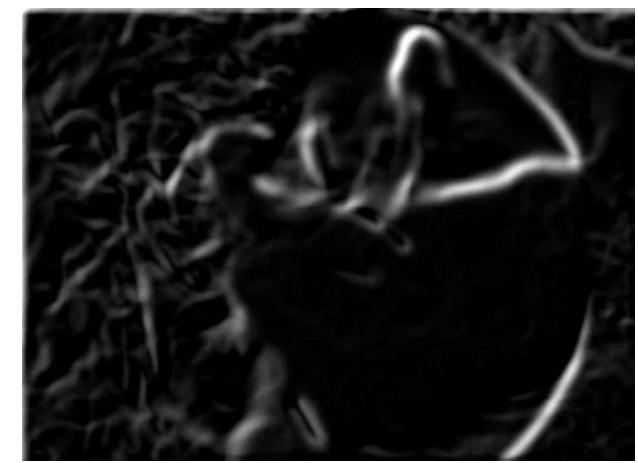
...



Effect of σ on derivatives



$\sigma = 1$ pixel



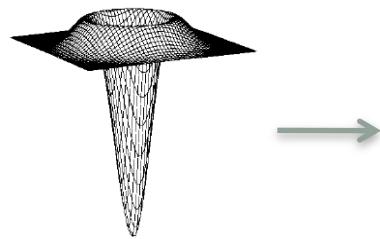
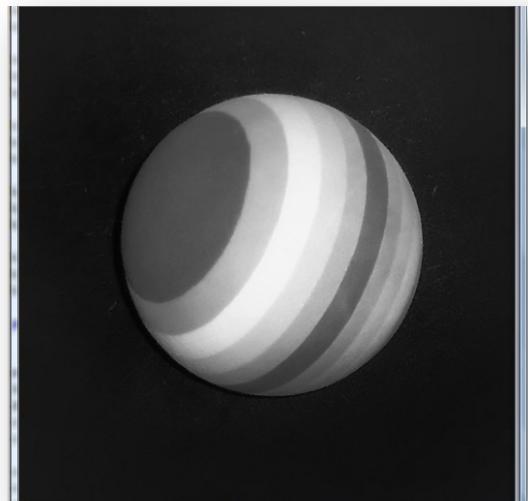
$\sigma = 3$ pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: scale edges detected.

Smaller values:features detected.

Laplacian

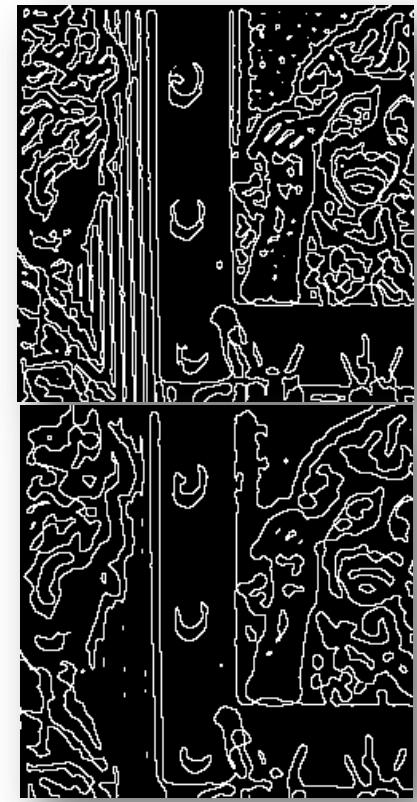


$$\nabla^2 h(x,y) = \left(\frac{x^2 + y^2}{\sigma^4} - \frac{2}{\sigma^2} \right) h(x,y)$$

Skimage:

```
im_l=filters.laplace(im)
```

Laplacian and zero-crossings



Using different sigma values.

Why edges obtained by the zero-crossing of the Laplacian map assure to be continuous?

Although the Laplacian operator convolved with the image leads to contours, they can be too much (depends on sigma).



Student Login

Room Name

A text input field containing the room name "COMPUTERVISION2021". The input field has a light blue border and a white background.

JOIN



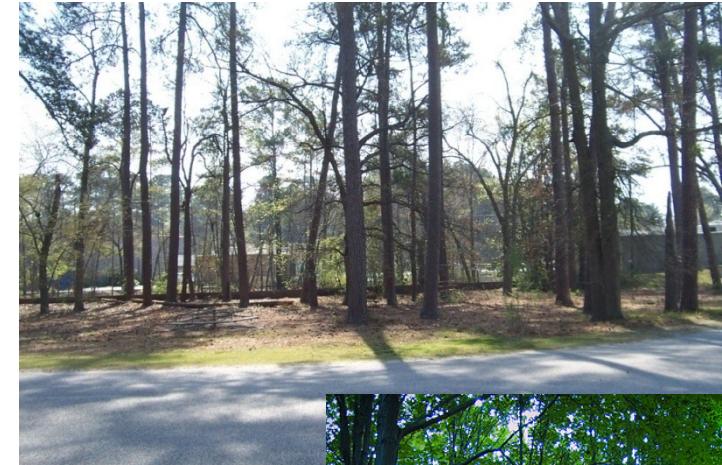
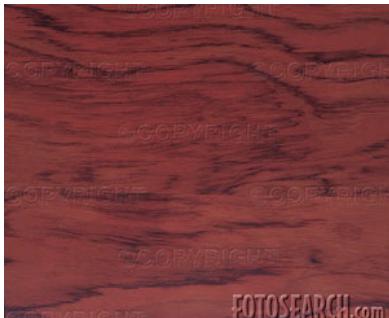
b.socrative.com/login/student

Today

- What is an image derivative and its relation to the edges
 - How to implement different edge detectors and what is their difference?
 - First and second derivatives, image gradient and Laplacian
- How to get image derivatives with a Gaussian
- Canny edge detector
- Application: hybrid images

So, what scale to choose?

It depends what we're looking for.



We need previous knowledge about the right scale.

Or manage different scales. How?

Original image



Gradient magnitude image



Thresholding gradient with a lower threshold



Thresholding gradient with a higher threshold



Canny edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
 *crestas*
- **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

Skimage:

```
>>im_canny=skimage.feature.canny(im, sigma=5)
```

The Canny edge detector



original image (Lena)

The Canny edge detector



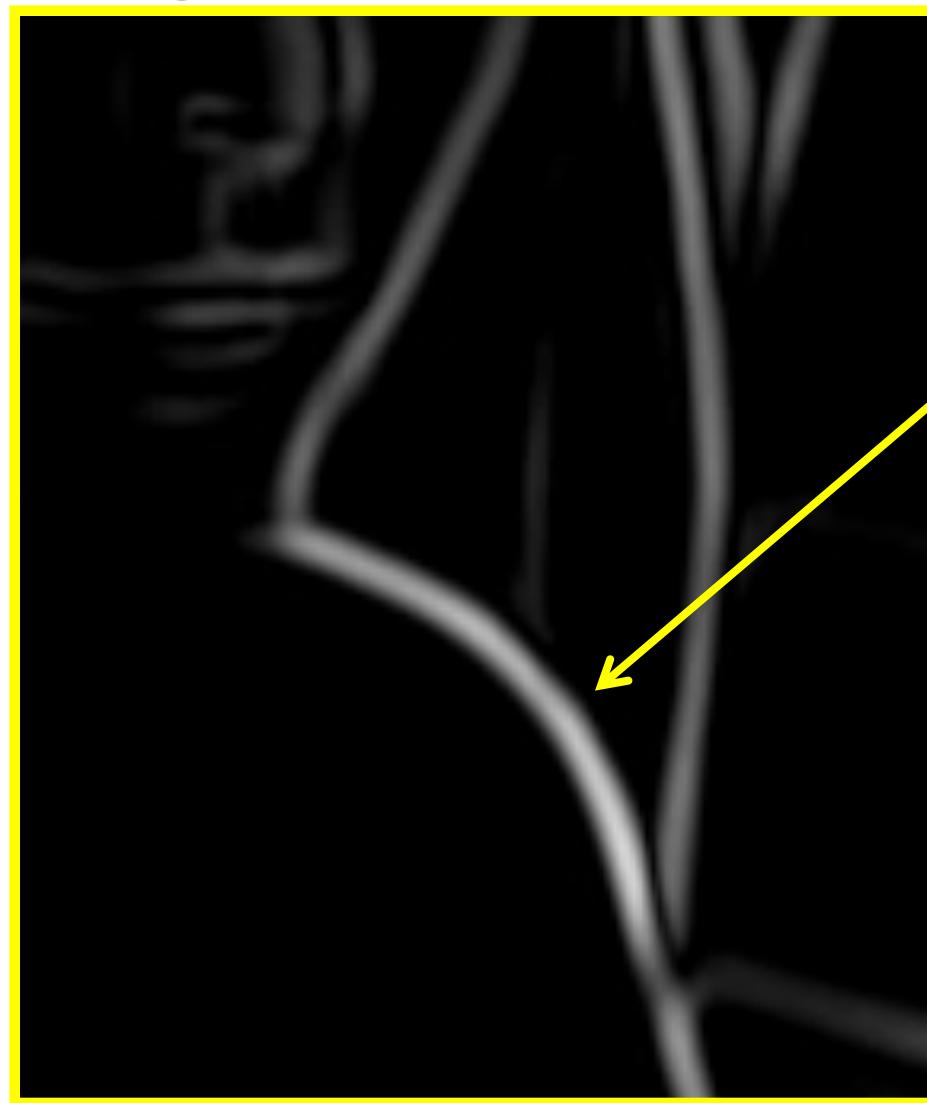
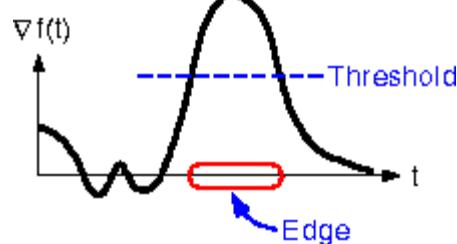
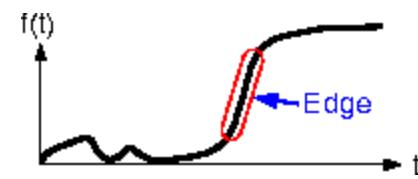
norm of the gradient

The Canny edge detector



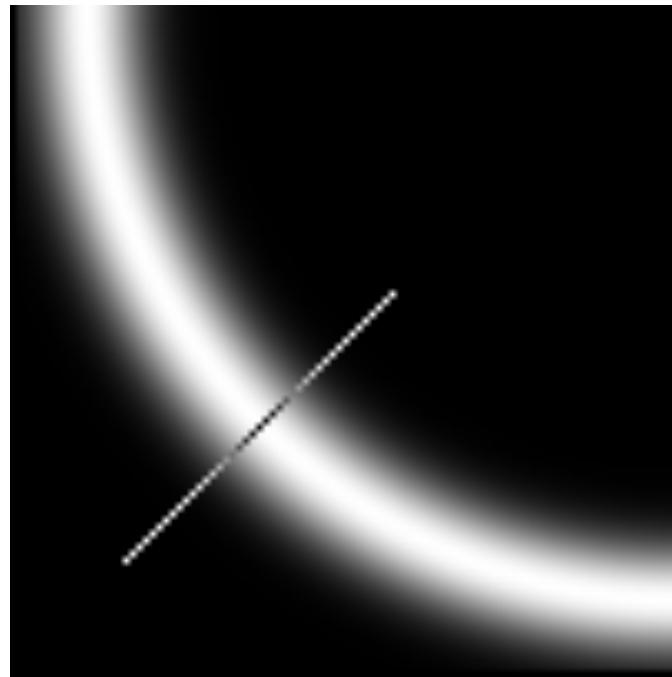
thresholding

The Canny edge detector



How to turn these thick regions of the gradient into curves?

Non-maximum suppression



Check if pixel is local maximum along gradient direction,
select single max across width of the edge

The Canny edge detector



thinning
(non-maximum suppression)

Problem:
pixels along
this edge
didn't
survive the
thresholding

Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



Hysteresis thresholding



original image



high threshold
(strong edges)



low threshold
(weak edges)



hysteresis threshold

Hysteresis thresholding



**high threshold
(strong edges)**



**low threshold
(weak edges)**



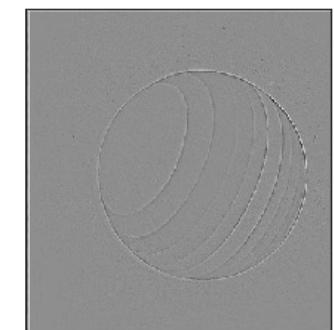
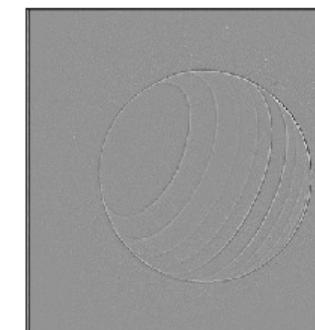
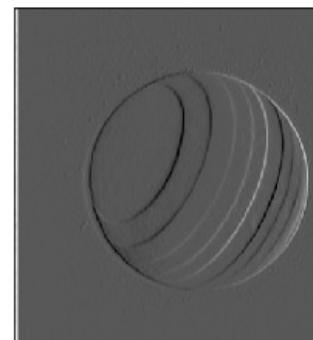
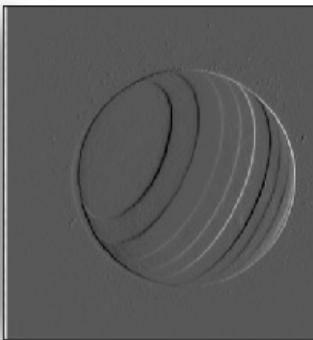
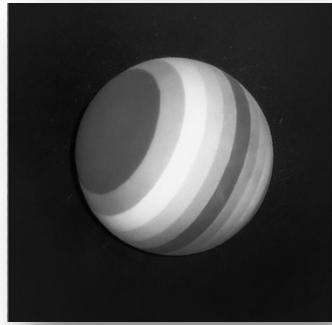
hysteresis threshold

Recap: Canny edge detector

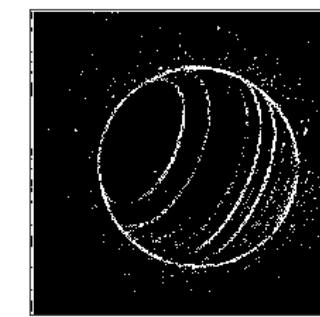
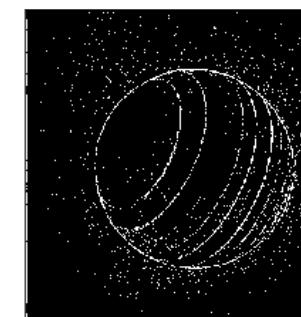
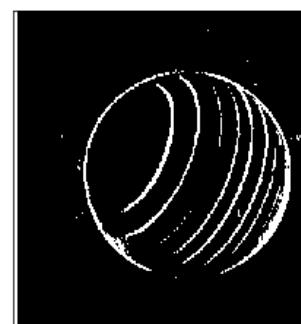
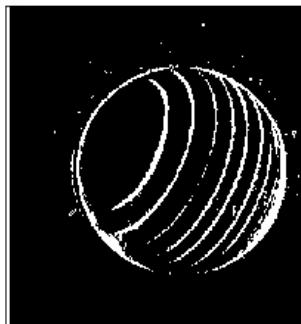
1. Filter image with derivative of Gaussian
 2. Find magnitude and orientation of gradient
 3. **Non-maximum suppression:**
 - Thin wide “ridges” down to single pixel width
 4. **Linking and thresholding (hysteresis):**
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
-
- In skimage:
 - `>>im_canny=skimage.feature.canny(im, sigma=5)`

Exercices (comparison)

- Determine the operator applied:



Original image



Sobel with a mask

```
mask=np.array([[-1,0,1],[-2,0,2],[-1,0,1]], dtype='float')
```

```
im_sobel=convolve(im,mask)
```

Prewitt with a mask

```
mask=np.array([[-1,0,1],[-1,0,1],[-1,0,1]], dtype='float')
```

```
im_prewitt=convolve(im,mask)
```

Laplacian with a mask

```
mask=np.array([[0,-1,0],[-1,4,-1],[0,-1,0]], dtype='float')
```

```
im2=convolve(im,mask)
```

```
fig=plt.imshow(im2<7.5, cmap='gray')
```

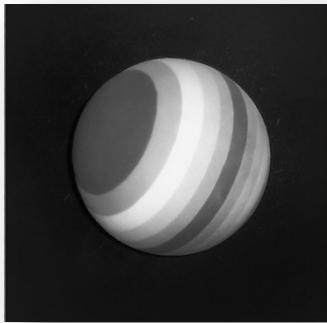
Laplacian with a mask

```
mask=np.array([[1,4,1],[4,-20,4],[1,4,1]], dtype='float')
```

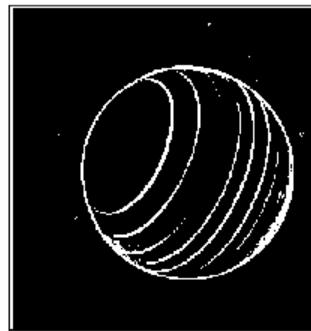
```
im_Lapl=convolve(im,mask)
```

Exercices (comparison)

- Determine the operator applied:



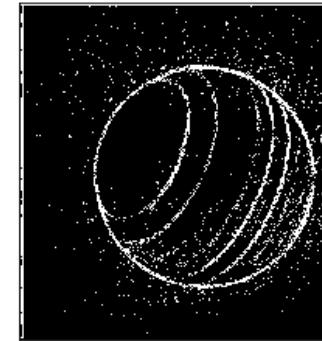
Original image



Sobel

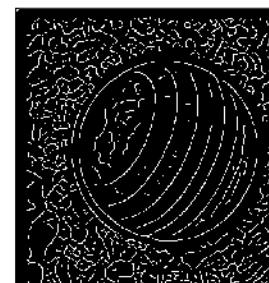
```
im2=filters.sobel(im)
```

```
fig=plt.imshow(abs(im2)>10.5)
```



Laplacian of a Gaussian

```
im2=filters.laplace(im)
```



Canny sigma=4

```
im2=feature.canny(im, sigma=4)
```



Canny sigma=6.5

```
im2=feature.canny(im, sigma=6.5)
```



Student Login

Room Name

JOIN

— — —

b.socrative.com/login/student

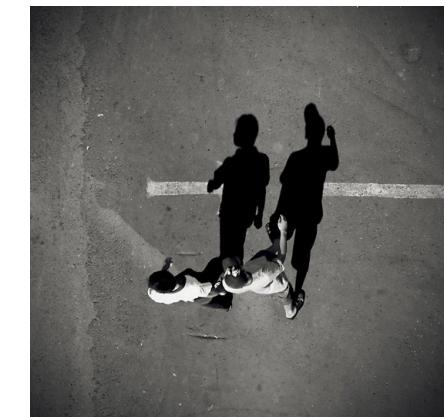
Low-level edges vs. perceived contours



Background



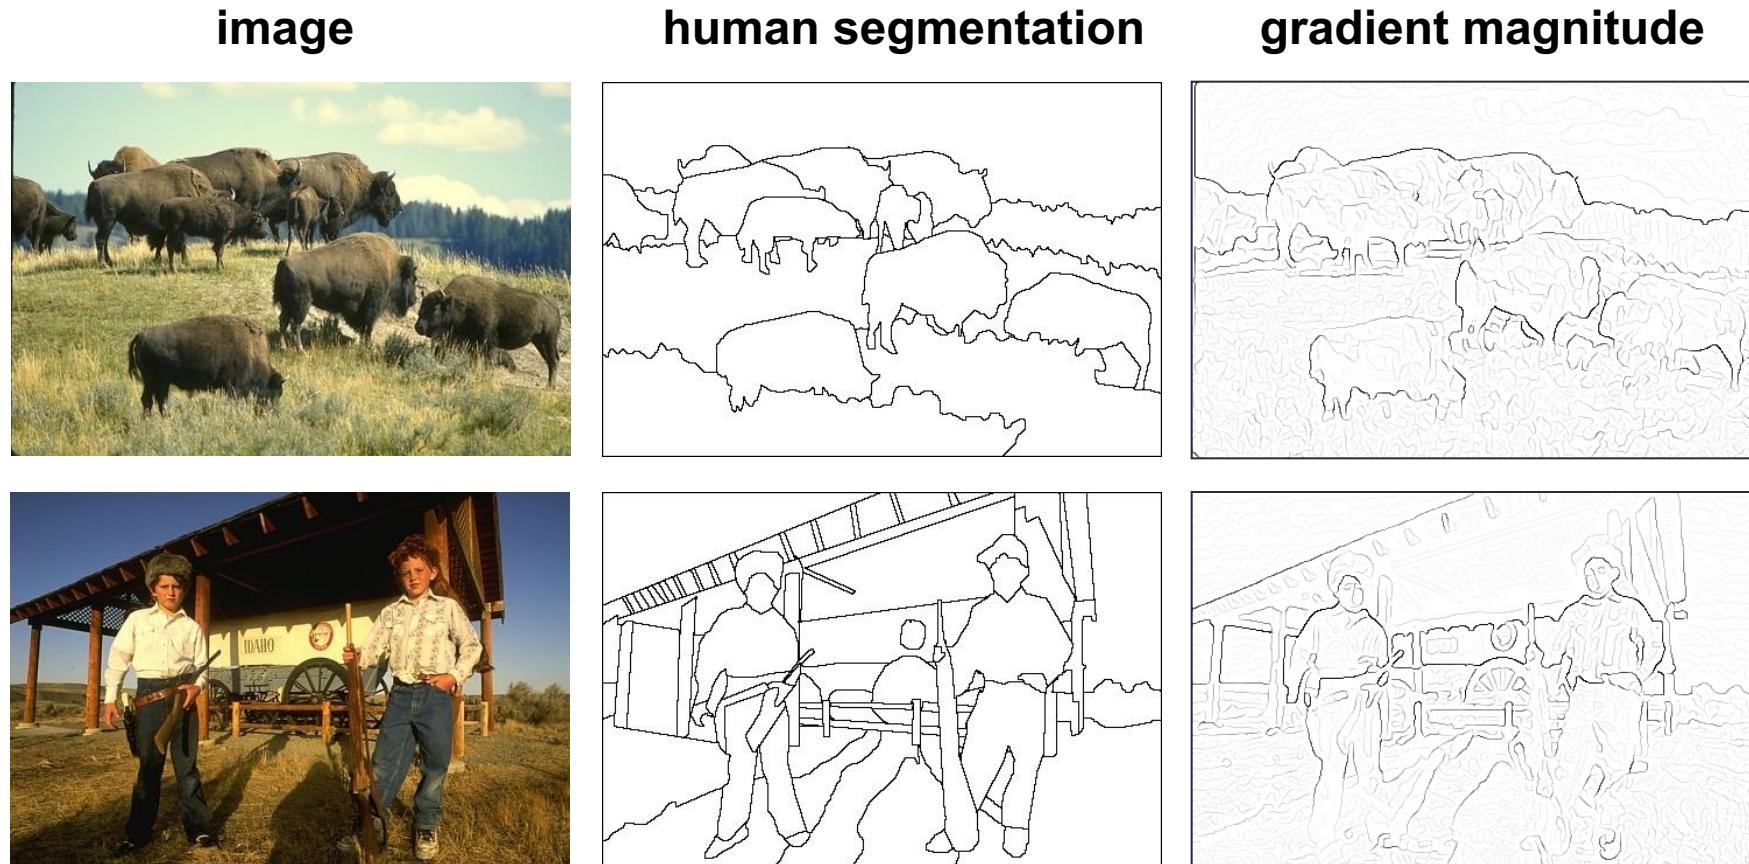
Texture



Shadows

Is Canny edge detector distinguishing these contours?

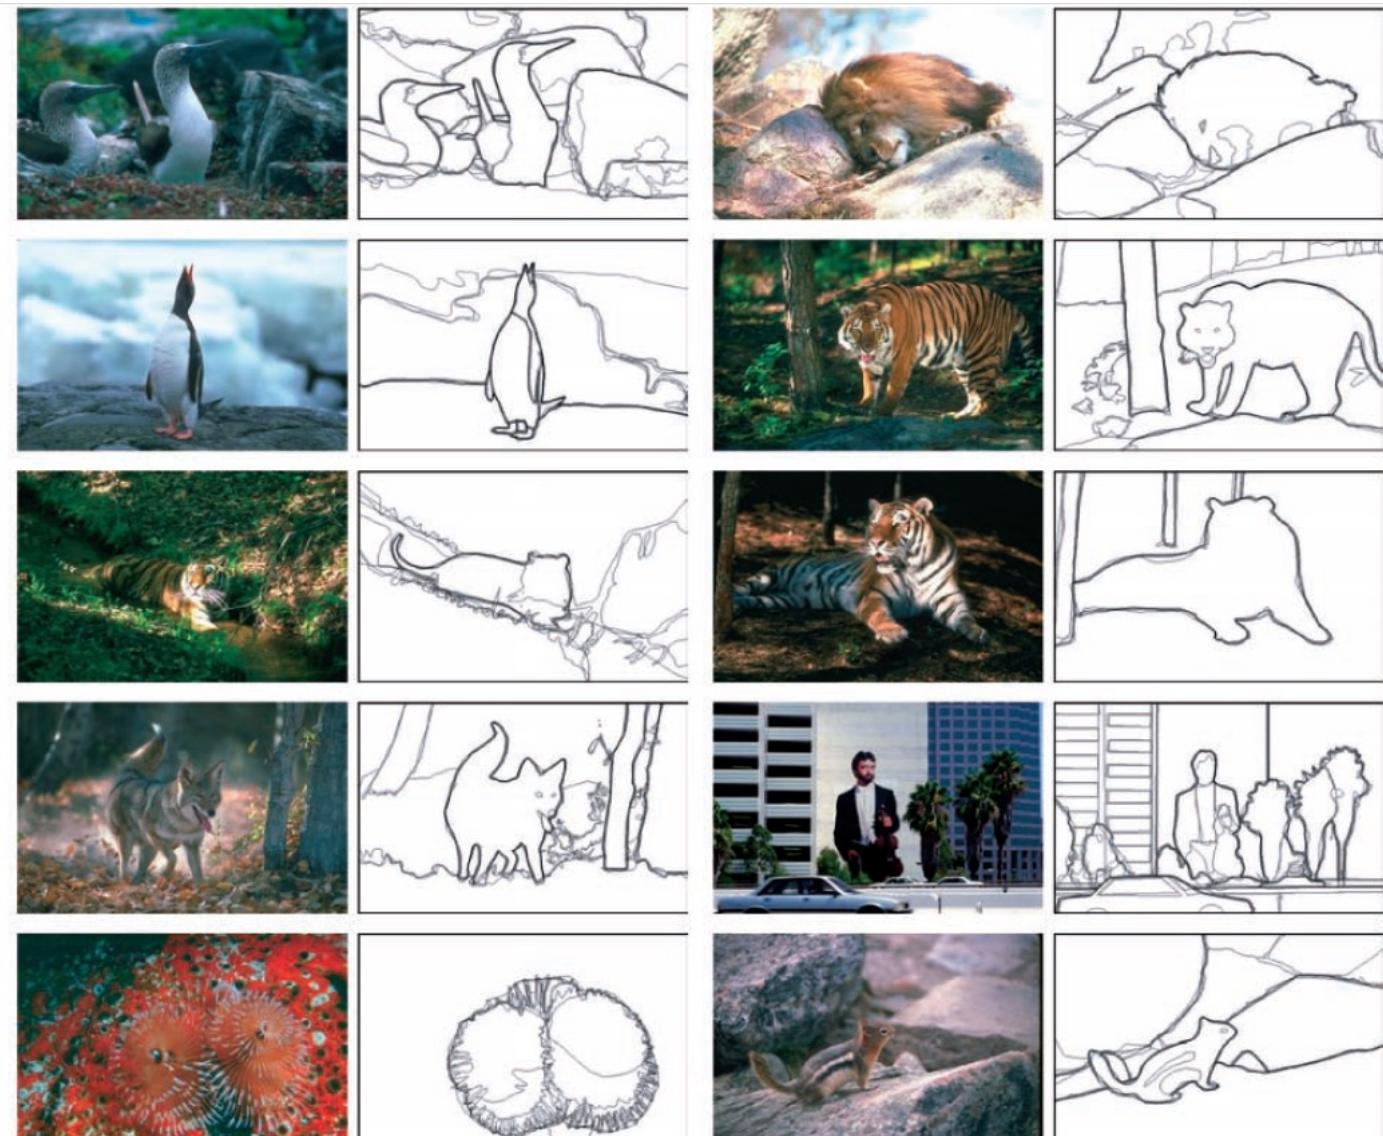
Low-level edges vs. perceived contours



- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

Learn from humans which combination of features is most indicative of a “good” contour?



Today

- What is an image derivative and its relation to the edges
 - How to implement different edge detectors and what is their difference?
 - First and second derivatives, image gradient and Laplacian
- How to get image derivatives with a Gaussian
- Canny edge detector
- Application: hybrid images

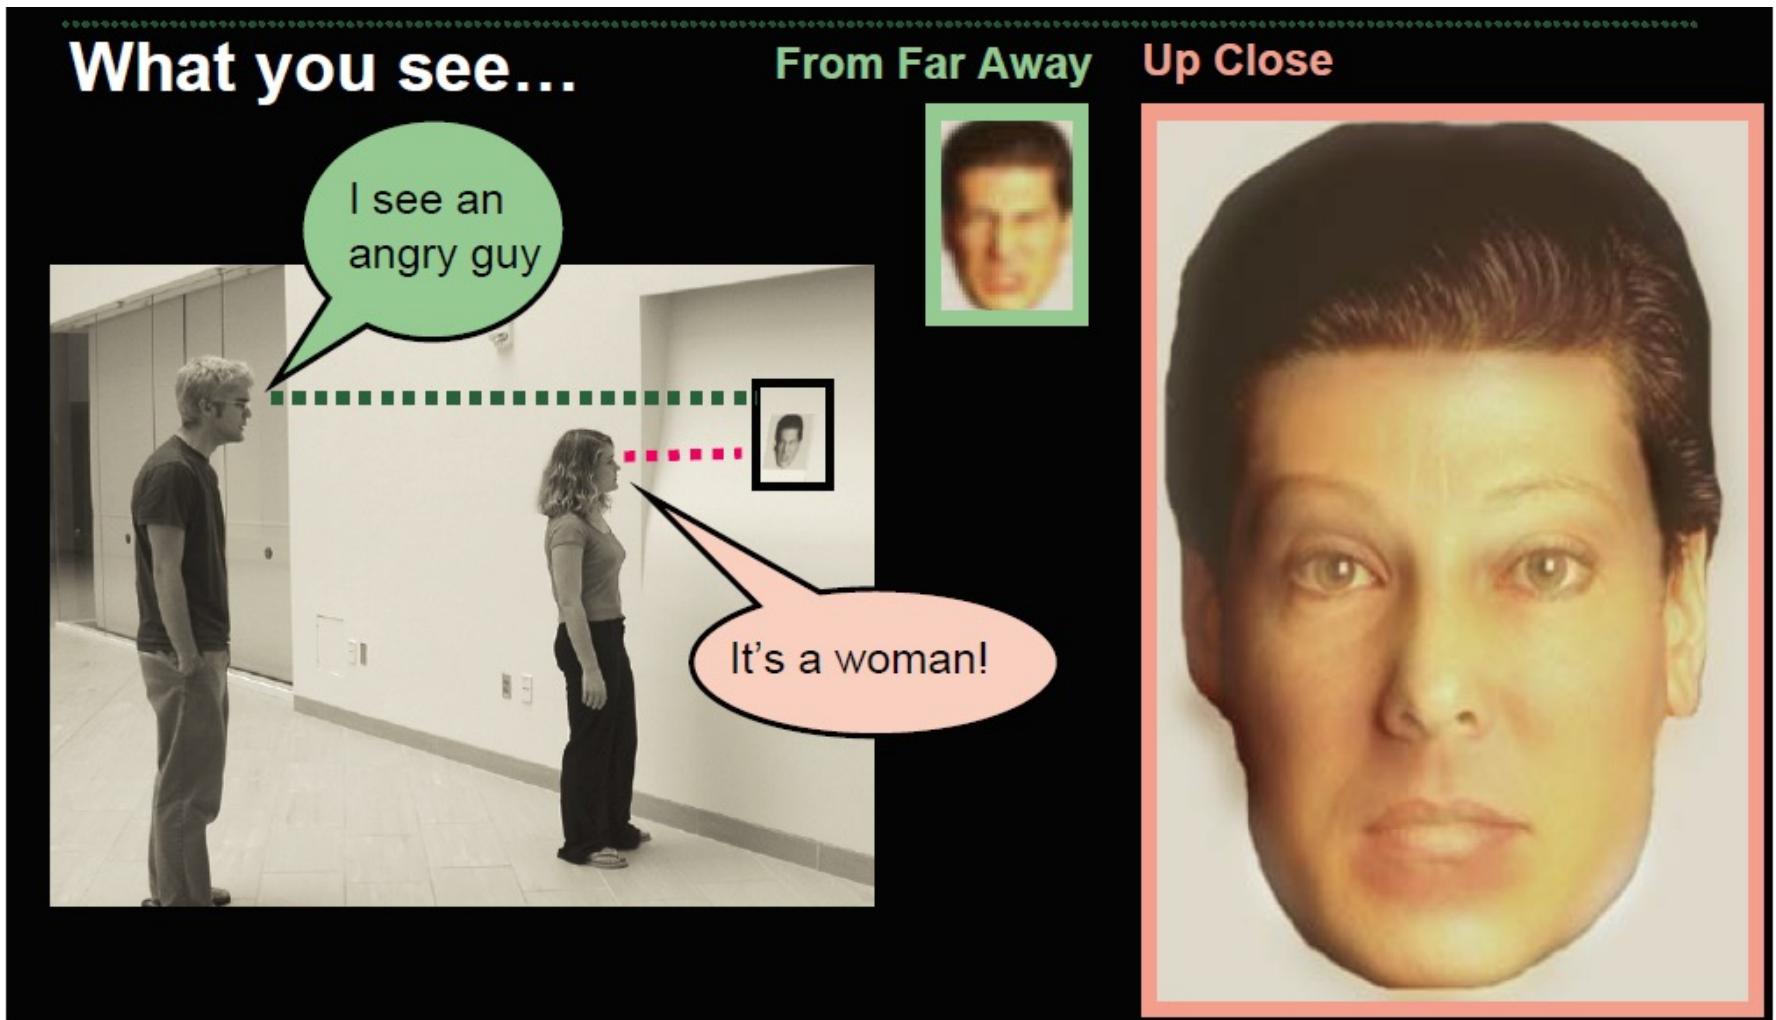
Squeeze

Qué derivada
utilizarías para
encontrar:

- a) La linea
continua
separadora
- b) El final del
carril?



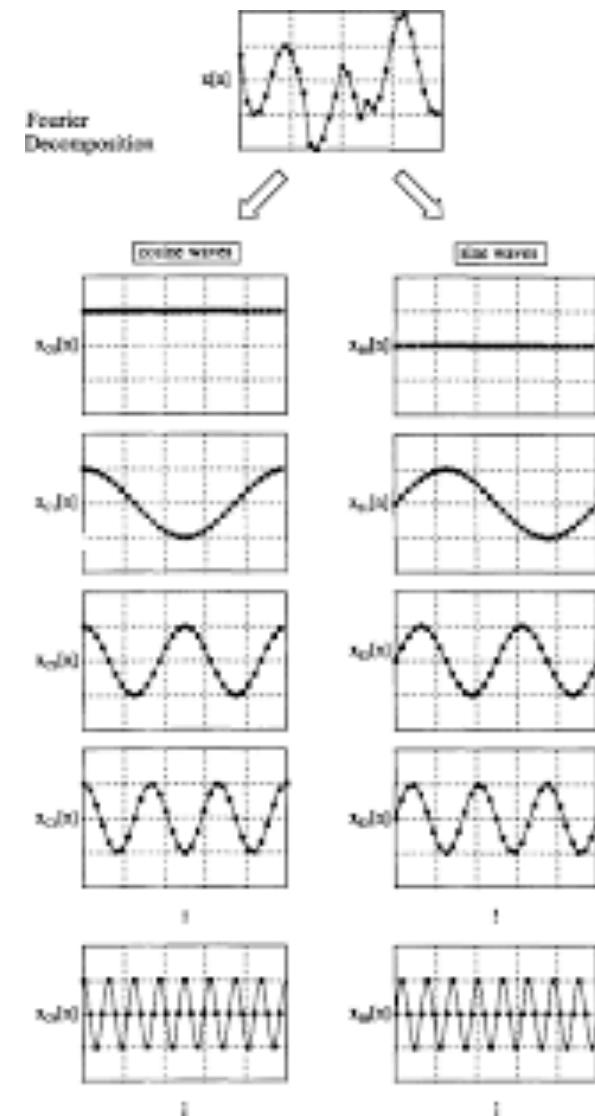
Application: Hybrid Images



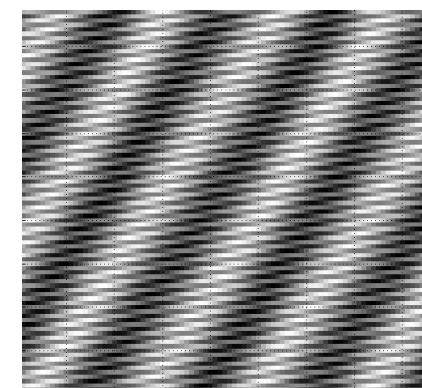
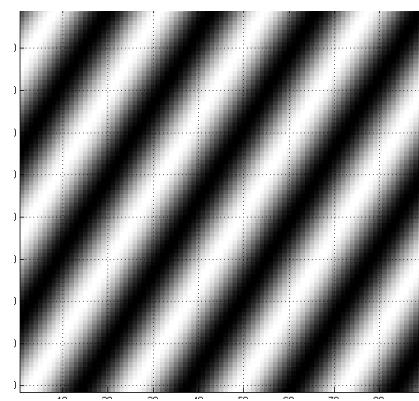
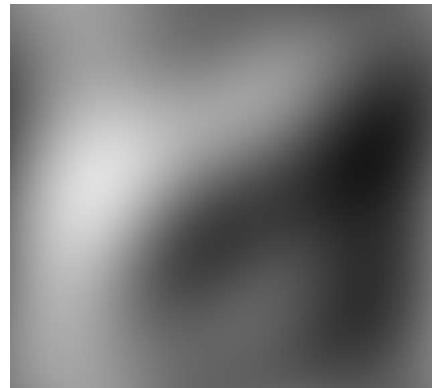
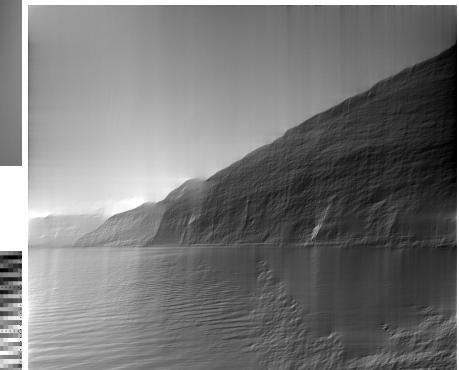
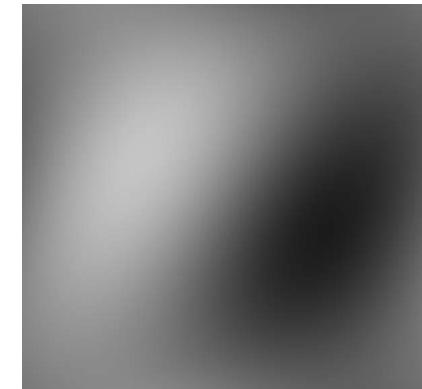
Fourier decomposition



Impulse noise



Intuition about low and high image frequencies



Sum of 50
frequencies

Low frequencies -> pixel values that are changing slowly over space.
High frequency -> pixel values that are rapidly changing in space.

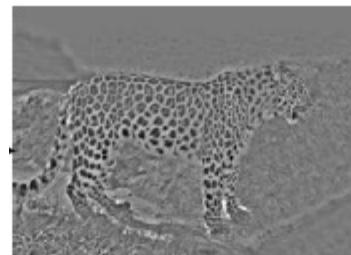
Application: Hybrid Images

Low image frequencies



How to extract them?

High image frequencies



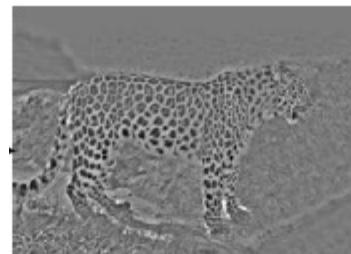
Application: Hybrid Images

Gaussian Filter



Low frequencies:

$$L(I_1) \rightarrow \text{Gaussian smoothing}$$



High frequencies:

$$H(I_2) = I_2 - L(I_2)$$



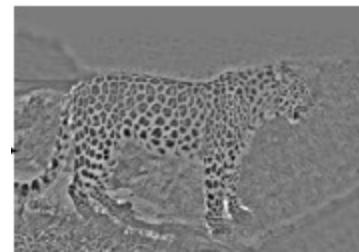
Laplacian Filter or the
difference btw the image and
its low frequencies

Application: Hybrid Images

Low frequencies

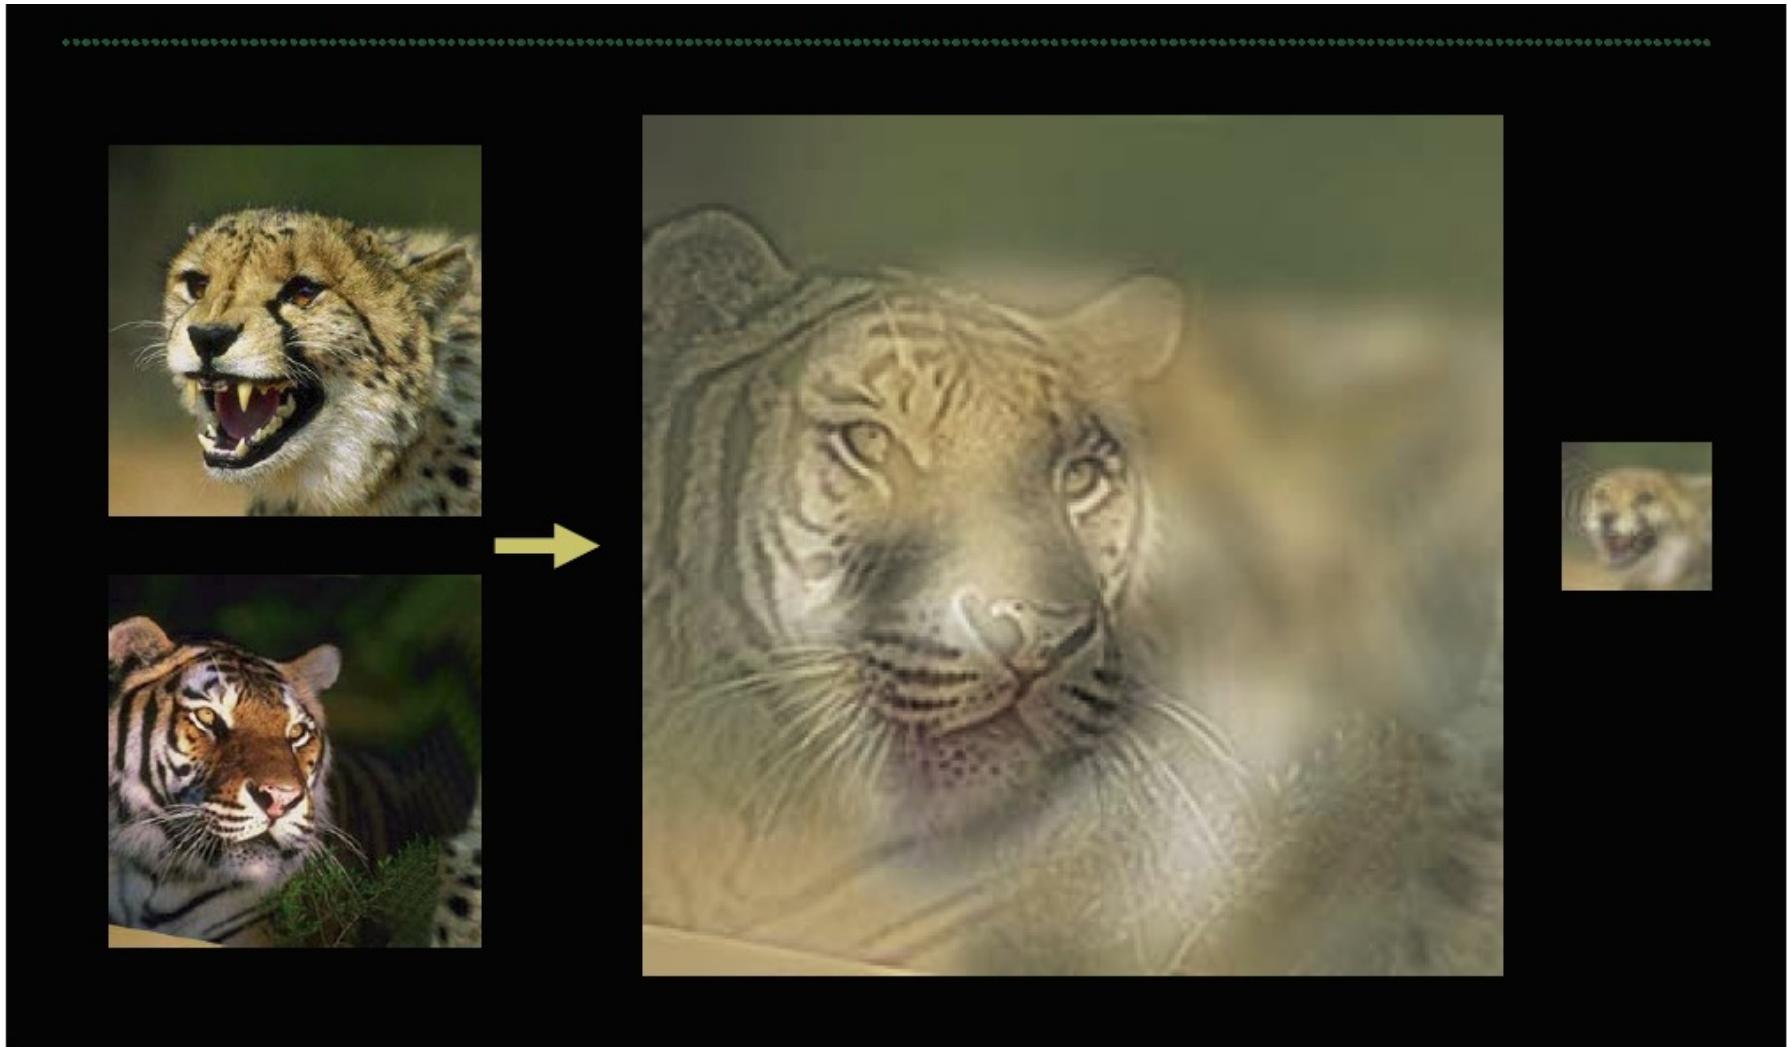


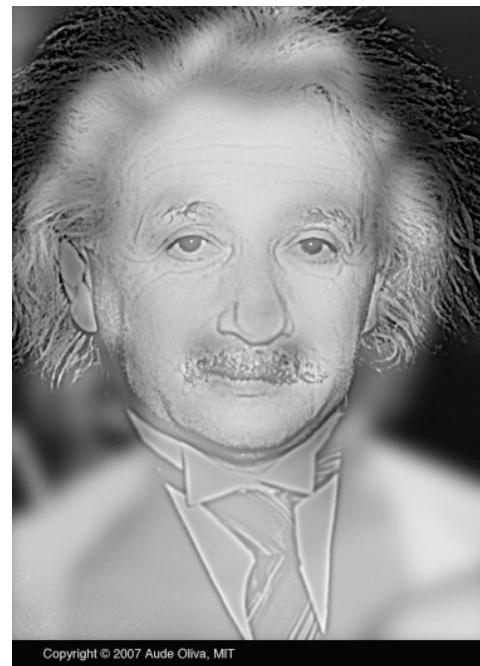
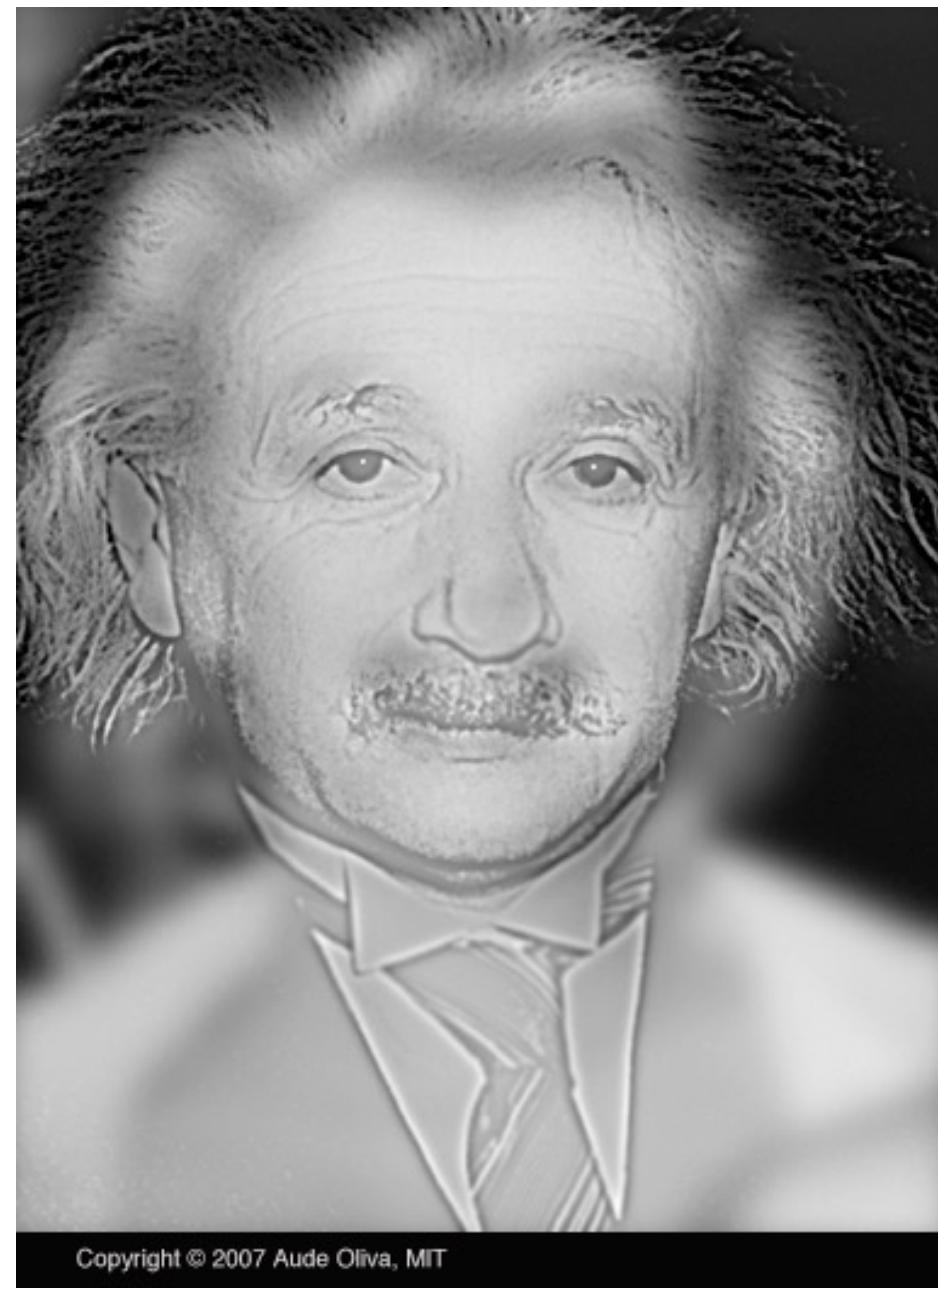
High frequencies



Hybrid image:
 $\text{Hybrid}(I_1, I_2) = L(I_1) + H(I_2)$

Note: In order to emphasize the effect, use different sigma parameters in the Gaussian





http://cvcl.mit.edu/hybrid_gallery/gallery.html

Changing expression



SIGGRAPH2006

Sad



Surprised



Summary

- Filters allow local image neighborhood to influence our description and features
 - Smoothing to reduce noise
 - Derivatives to locate contrast, high image gradient
- Different edge detectors:
 - Sobel & Prewitt – fast but sensitive to noise
 - Convolution with the Gradient of a Gaussian –
 - Less fast but more robust
 - Using different sigma allows to smooth more or less
 - Zero-crossing with a Laplacian – assures closed contours
 - Canny edge detector – **state of the art edge detector**
 - Assures continuous and thin contours due to the hysteresis and the thinning steps but needs parameters
 - Edges used to contain good contours but also many other pixels with high intensity change
- Still the question about the object contours is not solved!