

Synchronization - Part 1

Operating Systems – EDA093/DIT401

Vincenzo Gulisano
vincenzo.gulisano@chalmers.se



UNIVERSITY OF
GOTHENBURG

What to read (Main textbook)

- Both for part 1 and part 2

- Chapter 2.3.1-2.3.6, 2.5.1-2.5.2, 6.1-6.2, 6.5-6.6, 6.7.3-6.7.4;
- Quicker reading, for awareness, of sections 2.3.7-2.3.10, 6.3

(facultative reading: sections 6.1-6.7, 6.9, 7.1-7.5, 7.6-7.8 from OS Concepts by Silberschatz et-al).

Agenda

- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques

Agenda

- **Motivation**
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques

Previously on Operating Systems

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```


```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid, &attr, runner, argv[1]);
/* wait for the thread to exit */
pthread_join(tid, NULL);

printf("sum = %d\n", sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```



On the need for proper synchronization: example

- Setup
 - 2 threads, A and B
 - Each has access to:
 - a **local** variable x
 - a **shared** variable S , initialized with value 20
 - A and B run in parallel, each on its core
 - A and B run 3 instructions to update S (see figure)
- Consider the execution on the right:

Time (in discrete steps) ↓

Thread A	Thread B	Variable S
		20
$x \leftarrow S$		20
$x \leftarrow x + 50$	$x \leftarrow S$	20
$S \leftarrow x$	$x \leftarrow x - 20$	70
	$S \leftarrow x$	0

Agenda

- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques

The Critical-Section (aka Critical Region) Problem

Structure of process/thread P

Repeat

entry section

critical section

exit section

remainder section

Forever

n processes/threads competing

Assumptions:

- atomic read/write
- no process failures
- $n=2$ in the initial examples

Notice: duration of the [remainder section] is unbounded


Problem formulation:

A solution must **provide** the **entry** and **exit** sections

It must **ensure**:

1. **Mutual Exclusion.** **Only one process/thread** at a time is allowed to execute in its **critical section**.
2. **Progress (no deadlock/no livelock).** **If no process/thread is in its critical section** and some is/are **trying**, some must **enter its critical section in bounded time**.
3. **Fairness / Bounded Waiting / no starvation.** Variety of formulations, e.g., FCFS; or bounded #bypasses, ...

Critical section: desirable behavior



Beware!

- Is a solution valid?
- We are looking for formal answers
 - Can we find an example showing it is not? then it is NOT valid.
 - Can we show it is impossible to find such an example? Then it is valid.
- **An example showing it works is not sufficient**

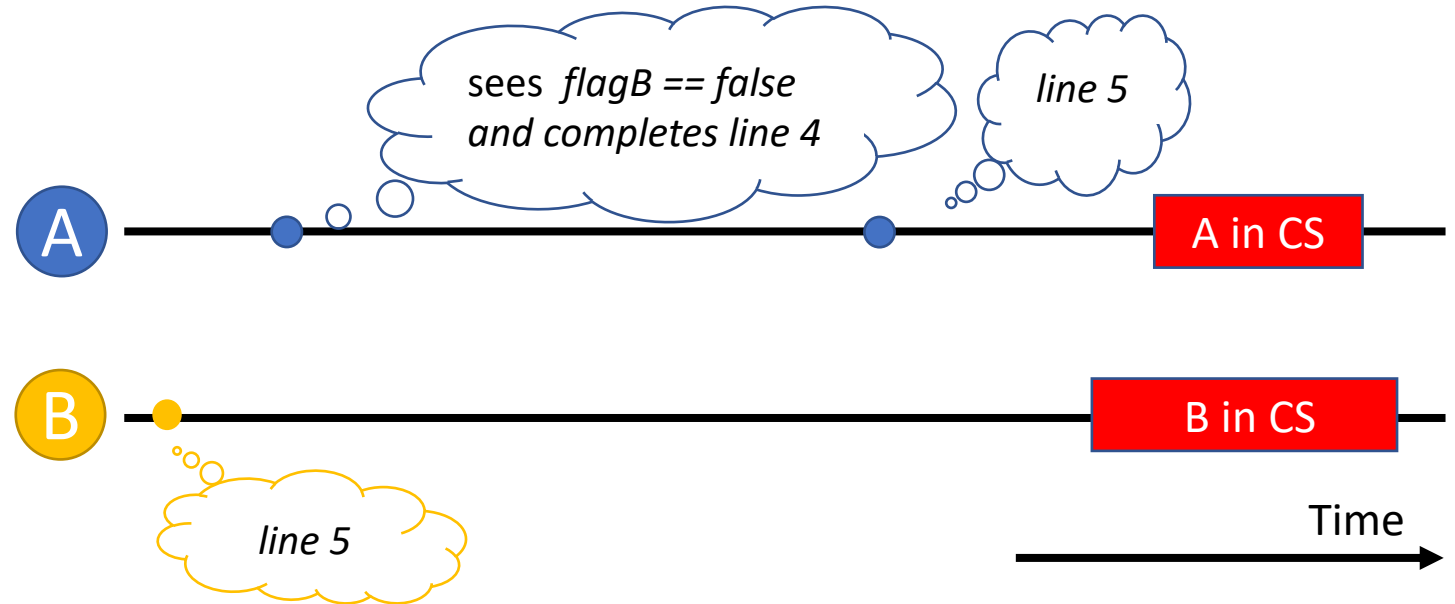


check-competition approach

1. Shared **var** *flagA, flagB: boolean*; // indicate A's, B's intention to enter; initially both *false*
2. Thread A // B is symmetric
3. **repeat**
4. **while** *flagB == true* **do** [nothing] //busywait
5. *flagA* \leftarrow *true*
6. [critical section]
7. *flagA* \leftarrow *false*
8. [remainder section]
9. **forever**

check-competition: mutual exclusion property?

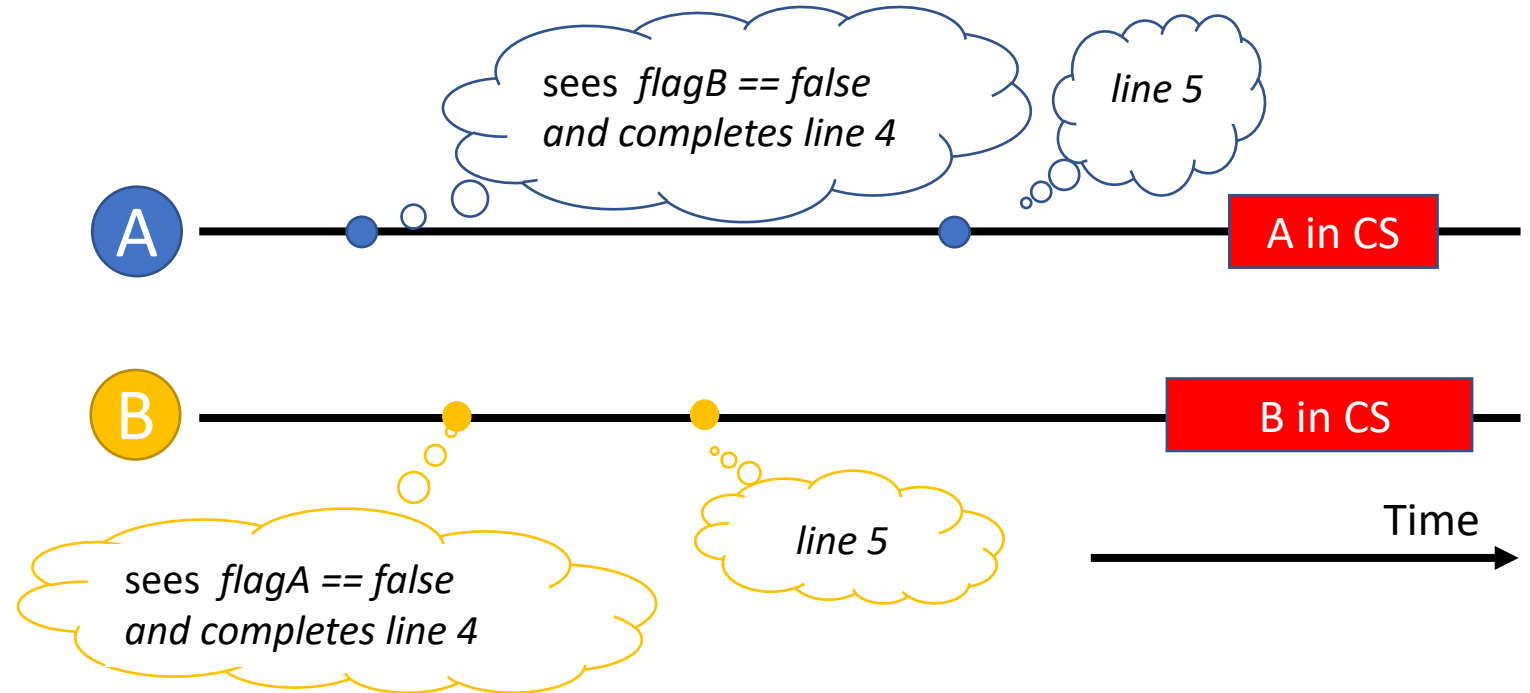
1. Shared **var** *flagA, flagB: boolean*; // indicate A's, B's intention to enter; initially both *false*
2. Thread A // B is symmetric
3. **repeat**
4. **while** *flagB == true* **do** [nothing] //busywait
5. *flagA* \leftarrow *true*
6. [critical section]
7. *flagA* \leftarrow *false*
8. [remainder section]
9. **forever**



check-competition: mutual exclusion property?

```

1. Shared var flagA, flagB: boolean; // indicate A's,
   B's intention to enter; initially both false
2. Thread A // B is symmetric
3. repeat
4.   while flagB == true do [nothing] //busywait
5.   flagA ← true
6.   [critical section]
7.   flagA ← false
8.   [remainder section]
9. forever
  
```

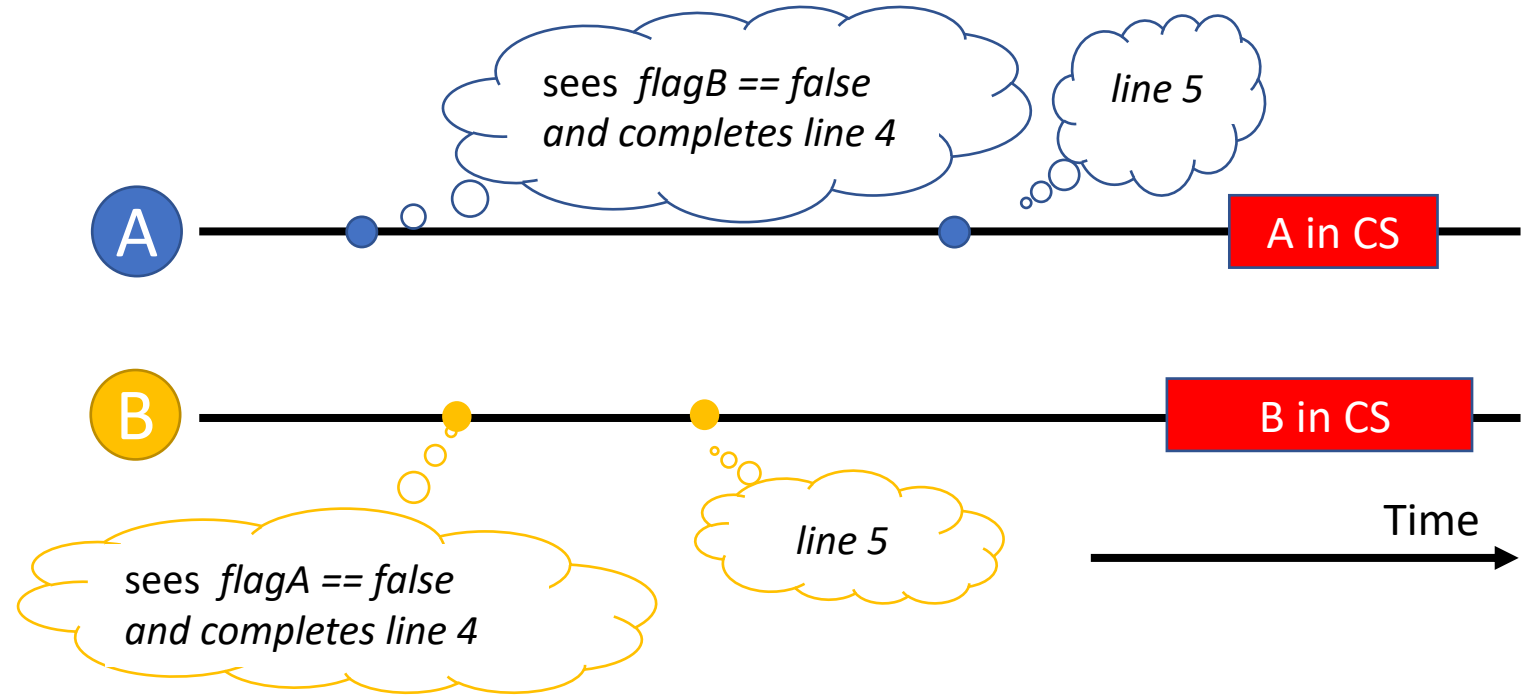


check-competition: mutual exclusion property?

```

1. Shared var flagA, flagB: boolean; // indicate A's,
   B's intention to enter; initially both false
2. Thread A // B is symmetric
3. repeat
4.   while flagB == true do nothing // busywait
5.   flagA ← true
6.   [critical section]
7.   flagA ← false
8.   [remainder section]
9. forever
  
```

NOT VALID



Agenda

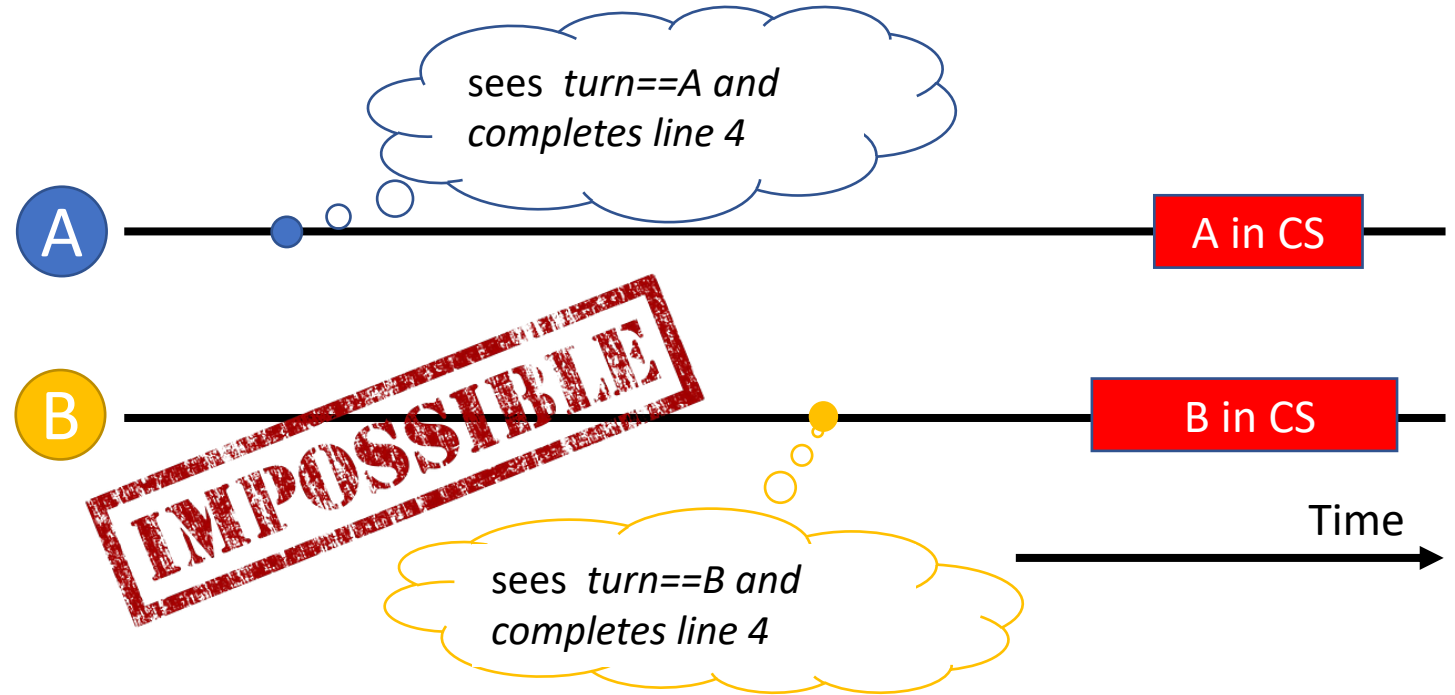
- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques

check-order (*after-you*) approach

1. Shared **var** turn: (idA..idB) // initially e.g. idA
// indicates whose turn it is to enter CS
2. Thread A // B is symmetric
3. **repeat**
4. **while** turn \neq idA **do** [nothing] //busywait
5. [critical section]
6. turn \leftarrow idB
7. [remainder section]
8. **forever**

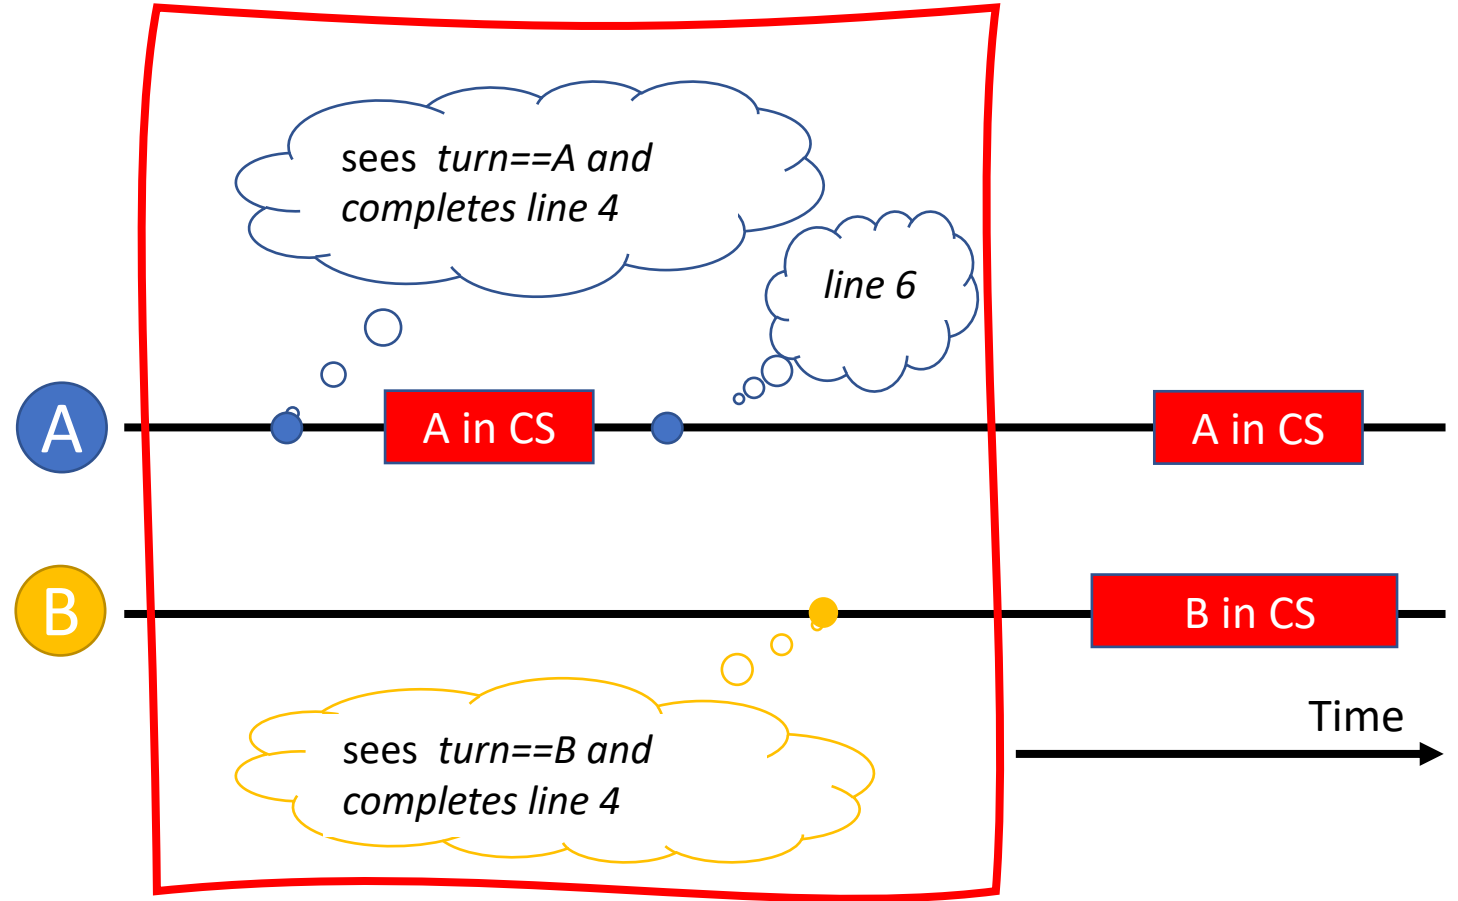
check-order: mutual exclusion property?

1. Shared **var** turn: (idA..idB) // initially e.g. idA
// indicates whose turn it is to enter CS
2. Thread A // B is symmetric
3. **repeat**
4. **while** turn \neq idA **do** [nothing] //busywait
5. [critical section]
6. turn \leftarrow idB
7. [remainder section]
8. **forever**



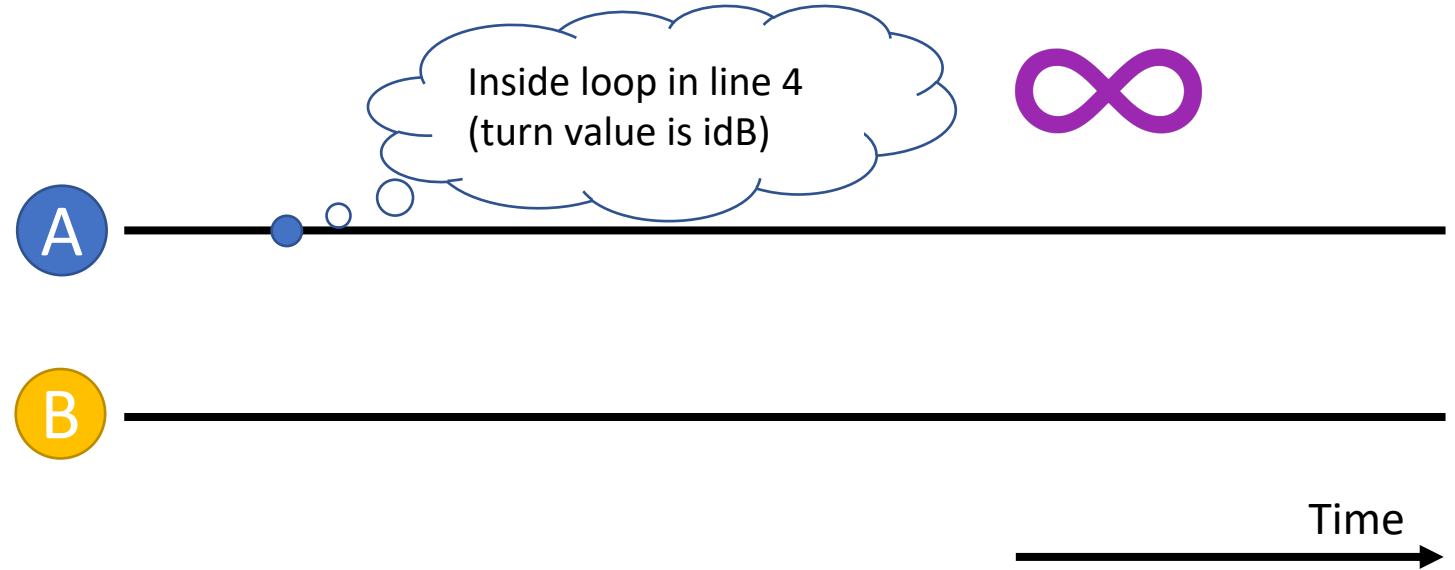
check-order: mutual exclusion property?

1. Shared **var** `turn: (idA..idB)` // initially e.g. `idA`
// indicates whose turn it is to enter CS
2. Thread A // B is symmetric
3. **repeat**
4. **while** `turn ≠ idA` **do** [nothing] //busywait
5. [critical section]
6. `turn ← idB`
7. [remainder section]
8. **forever**



check-order: progress property?

1. Shared **var** turn: (idA..idB) // initially e.g. idA
// indicates whose turn it is to enter CS
2. Thread A // B is symmetric
3. **repeat**
4. **while** turn \neq idA **do** [nothing] //busywait
5. [critical section]
6. turn \leftarrow idB
7. [remainder section]
8. **forever**

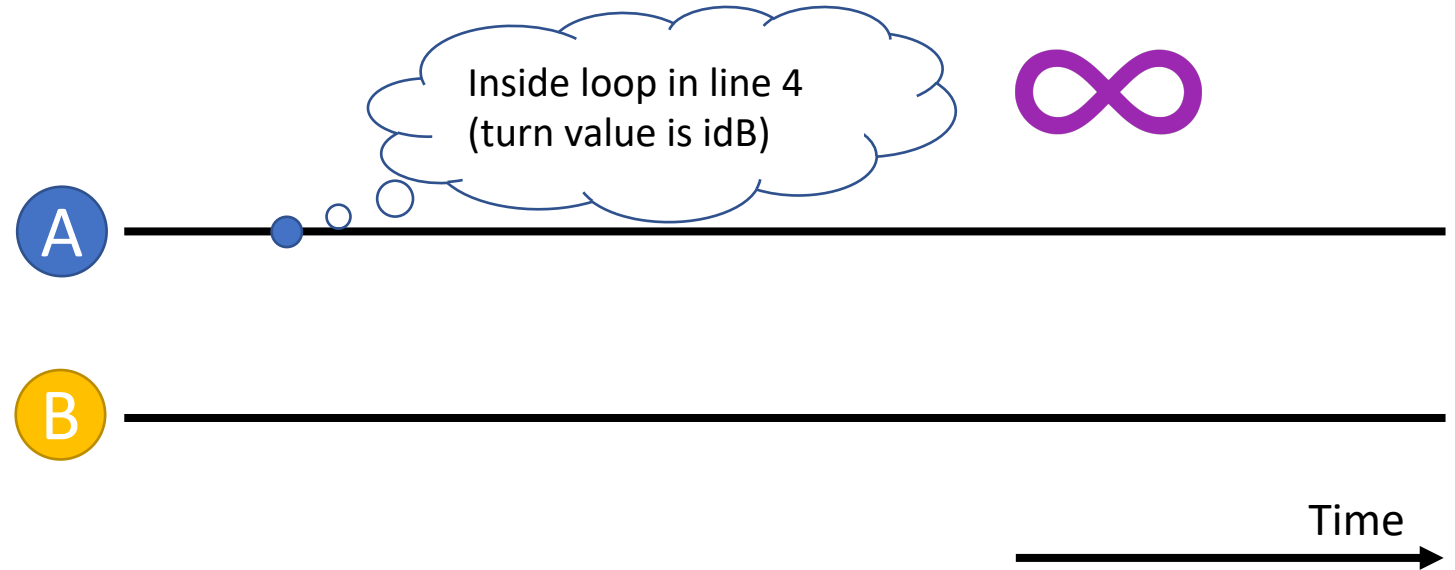


check-order: progress property?

```

1. Shared var turn: (idA..idB) // initially e.g. idA
   // indicates whose turn it is to enter CS
2. Thread A // B is symmetric
3. repeat
4. while turn = idA do nothing // just wait
5. [critical section]
6. turn ← idB
7. [remainder section]
8. forever
  
```

NOT VALID



Agenda

- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- **Peterson's Algorithm**
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques


Peterson's algo (*check order&competition*) approach

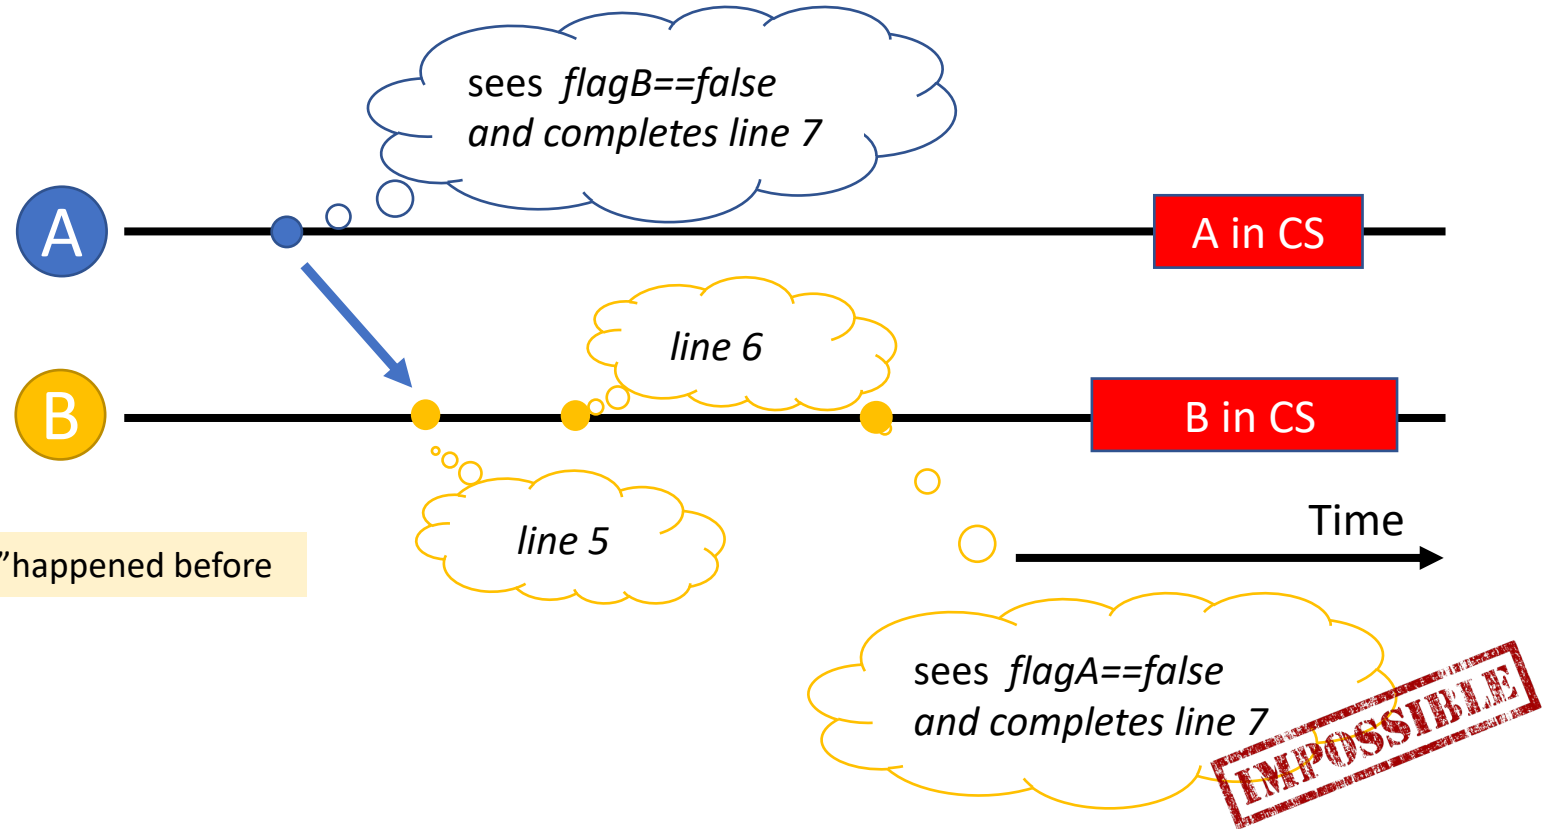
1. Shared var turn: (idA..idB) // initially e.g., idA
2. Shared var flagA, flagB: Boolean // initially false
3. thread A // B is symmetric
4. **repeat**
5. flagA \leftarrow true
6. turn \leftarrow idB
7. **while** (flagB and turn == idB) **do** [nothing]
8. [critical section]
9. flagA \leftarrow false
10. [remainder section]
11. **forever**

Peterson's algo: mutual exclusion property?

```

1. Shared var turn: (idA..idB) // initially e.g., idA
2. Shared var flagA, flagB: Boolean // initially false
3. thread A // B is symmetric
4. repeat
5.   flagA ← true
6.   turn ← idB
7.   while (flagB and turn == idB) do [nothing]
8.   [critical section]
9.   flagA ← false
10.  [remainder section]
11. forever
  
```

-> ,  denote "happened before"

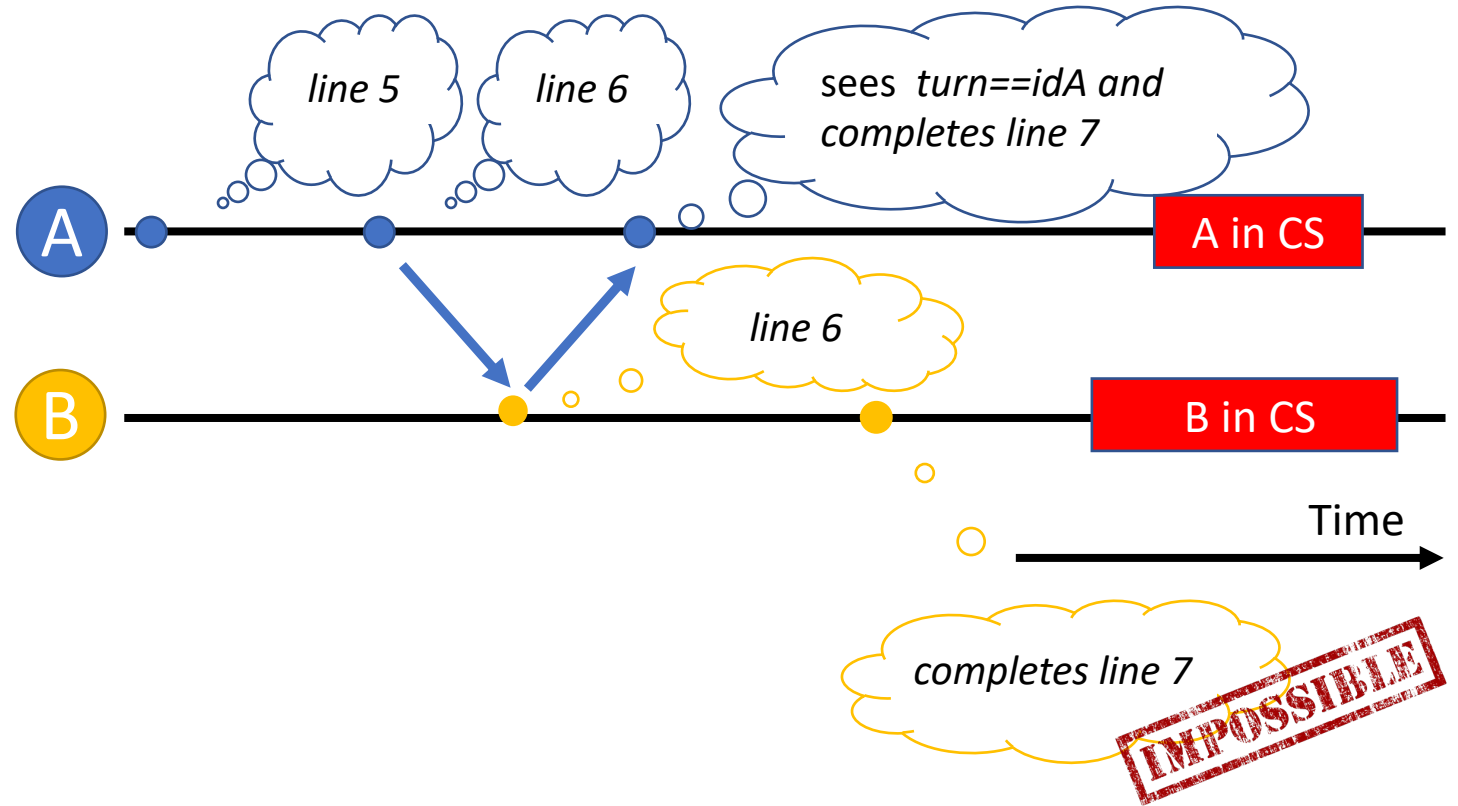


Peterson's algo: mutual exclusion property?

```

1. Shared var turn: (idA..idB) // initially e.g., idA
2. Shared var flagA, flagB: Boolean // initially false

3. thread A // B is symmetric
4.   repeat
5.     flagA ← true
6.     turn ← idB
7.     while (flagB and turn == idB) do [nothing]
8.     [critical section]
9.     flagA ← false
10.    [remainder section]
11.  forever
  
```

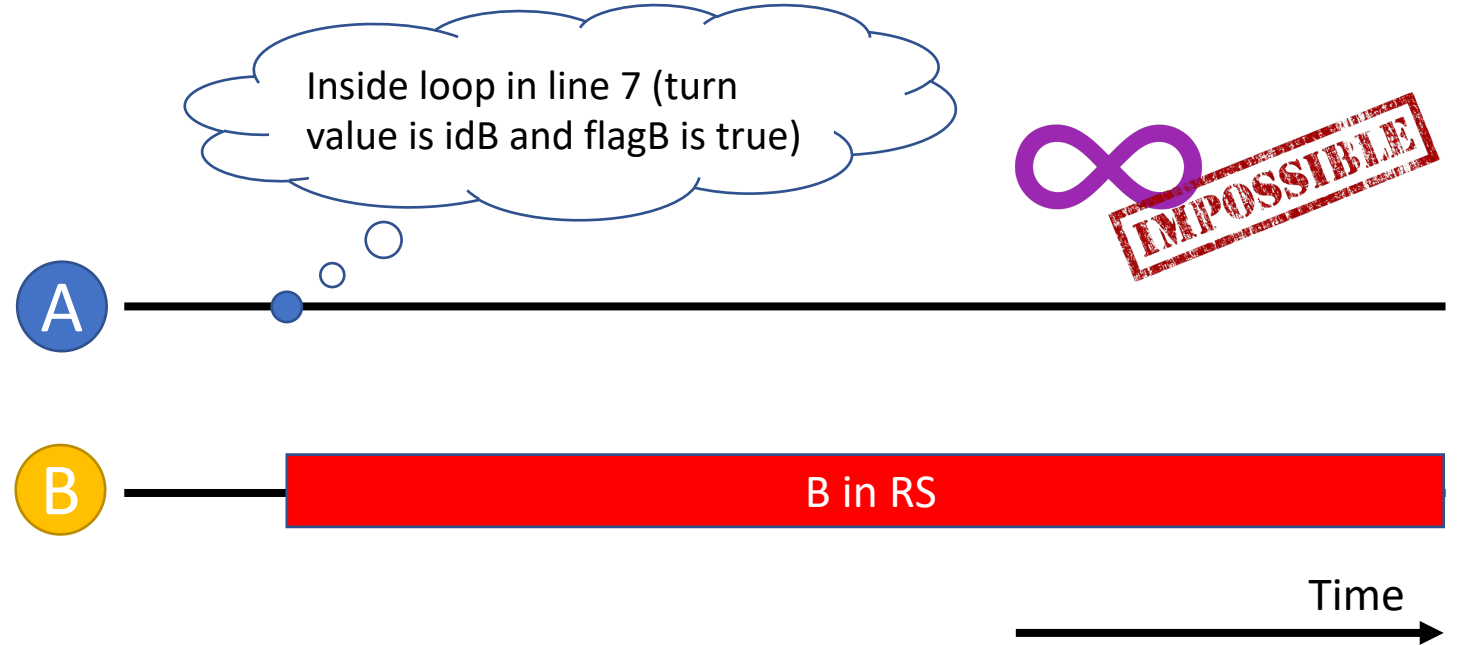


Peterson's algo: progress property?

```

1. Shared var turn: (idA..idB) // initially e.g., idA
2. Shared var flagA, flagB: Boolean // initially false

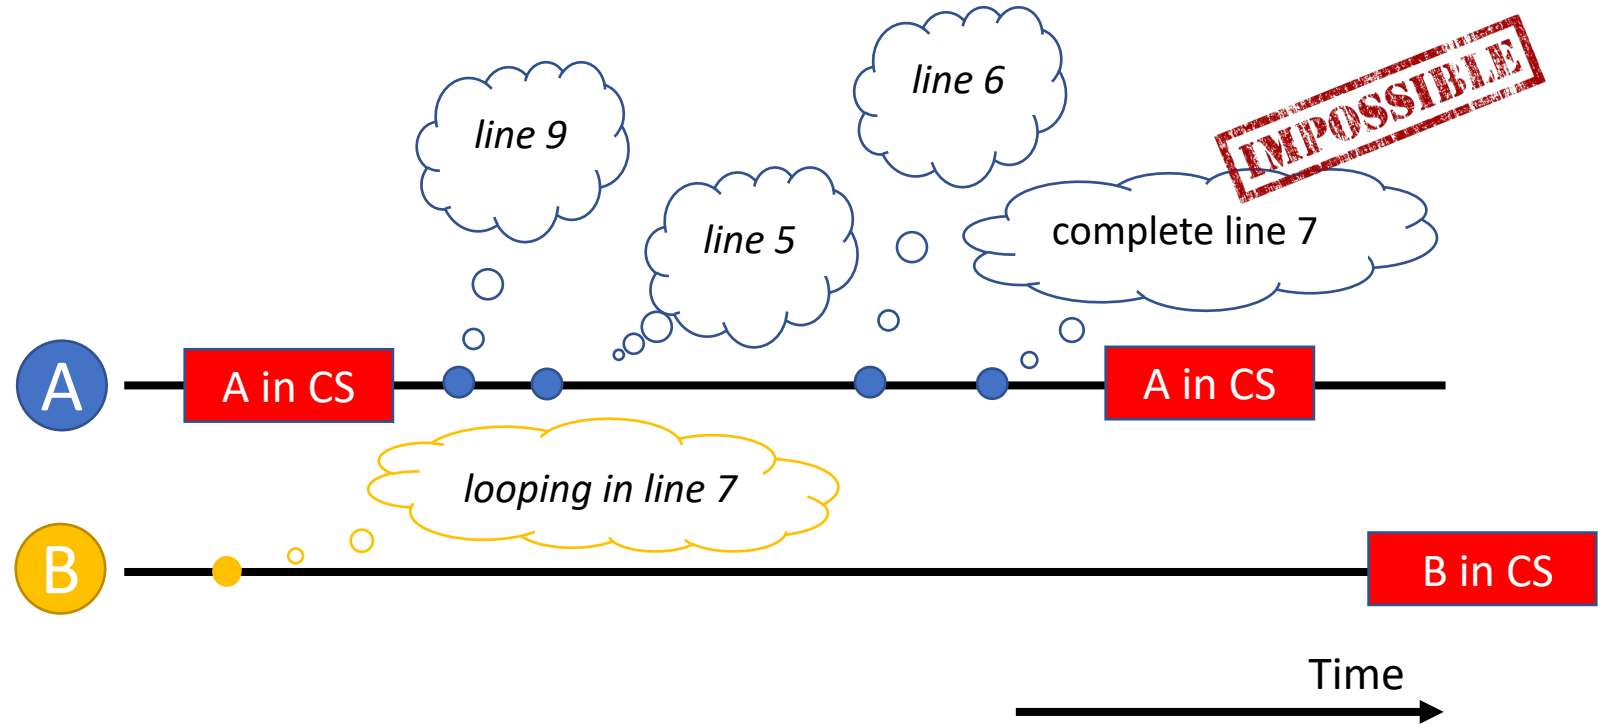
3. thread A // B is symmetric
4.   repeat
5.     flagA ← true
6.     turn ← idB
7.     while (flagB and turn == idB) do [nothing]
8.     [critical section]
9.     flagA ← false
10.    [remainder section]
11.  forever
  
```



Peterson's algo: fairness property?

```
1. Shared var turn: (idA..idB) // initially e.g., idA
2. Shared var flagA, flagB: Boolean // initially false

3. thread A // B is symmetric
4.   repeat
5.     flagA ← true
6.     turn ← idB
7.     while (flagB and turn == idB) do [nothing]
8.     [critical section]
9.     flagA ← false
10.    [remainder section]
11.  forever
```



On counterexamples & proofs by way of contradiction:

- This way of thinking and arguing, when we are presented with a solution to a synchronization problem/race-condition (i.e., try to “attack” it, try to find if it does not work), is practical and useful (and thus very common)
- It is very frequently used when arguing about correctness/incorrectness of proposed solutions. **This way of arguing helps us to:**
 - give a **proof by contradiction that a proposition in focus is *_correct_***, e.g., that mutual exclusion is preserved;
 - see about **proofs by way of contradiction** @ e.g., <http://cgm.cs.mcgill.ca/~godfried/teaching/dm-reading-assignments/Contradiction-Proofs.pdf> p 1-4 or <http://web.stanford.edu/class/archive/cs/cs103/cs103.1152/lectures/02/Small02.pdf> p 45-47
 - OR identify a **counter-example, which implies that the proposition in focus is *incorrect*** (i.e., it *disproves* the proposition)
 - see also <https://mathforlove.com/lesson/counterexamples/> and/or <https://study.com/academy/lesson/counterexample-in-math-definition-examples.html> and/or <http://web.stanford.edu/class/archive/cs/cs103/cs103.1152/lectures/02/Small02.pdf>

Agenda

- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques

Critical section and hardware support: interrupt disabling

Structure of process/thread P

Repeat

interrupt disable

critical section

interrupt enable

remainder section

Forever

Is there a downside?!

- If run in single-processor system: **limited ability to interleave programs**
- **In multiprocessors:** disabling interrupts on one processor/core will not guarantee absence of race conditions by threads running on other processors/cores

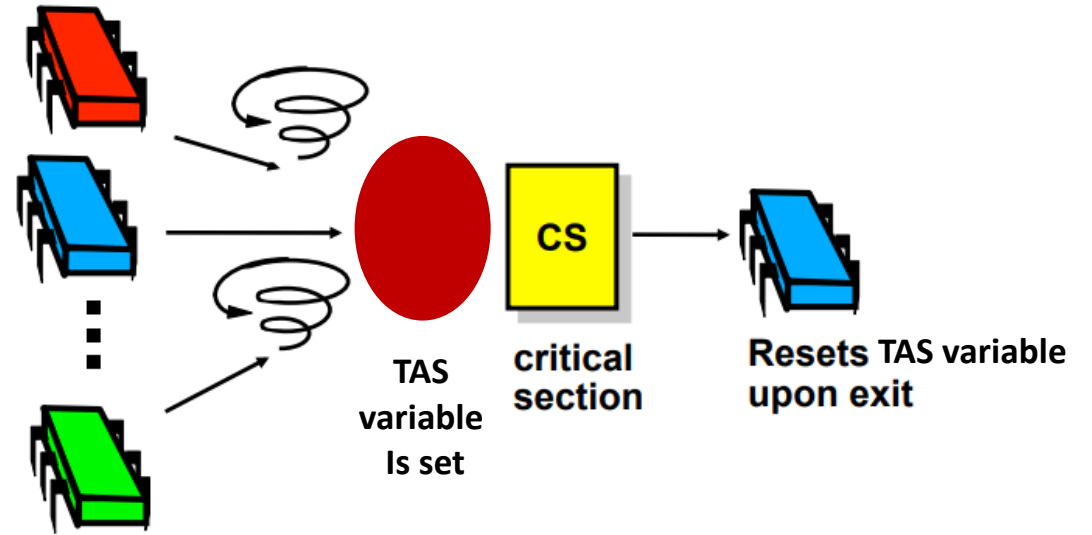
Critical section and hardware support: Read-Modify-Write (RMW) instructions

TestAndSet (aka TAS) Instruction Definition:

```
boolean TestAndSet (boolean *target){
    boolean rv = *target;
    *target := true;
    return rv} // Executed atomically by the HW
```

Shared Boolean **var** *lock*; initialized to *false*

```
repeat
    while (TestAndSet(&lock)) do [nothing];
    [critical section]
    lock := false;
    [remainder section]
forever
```



Homework training: argue as we did for Peterson's algo regarding correctness (mutual exclusion, progress) ; what about fairness?

Critical section and hardware support: Read-Modify-Write (RMW) instructions

TestAndSet (aka TAS) Instruction Definition:

```
boolean TestAndSet (boolean *target) {
    boolean rv = *target;
    *target := true;
    return rv} // Executed atomically by the HW
```

Shared Boolean **var** *lock*; initialized to *false*

```
repeat
    while (TestAndSet(&lock)) do [nothing];
    [critical section]
    lock := false;
    [remainder section]
forever
```

CompareAndSwap (aka CAS) Instruction Definition:

```
int CompareAndSwap(int *V, int expected, int
    new_value) {
    int temp := *V;
    if (temp == expected) then *V := new_value;
    return temp} // Executed atomically by the HW
```

Shared int **var** *lock*; initialized to 0;

```
repeat{
    while CompareAndSwap(&lock,0,1)!=0) do nothing;
    [critical section]
    lock := 0 ;
    [remainder section]
forever
```

A reflection about busy-waiting

- Rather simple, convenient; BUT:
 - Busy-waiting consumes processor time unnecessarily
 - Starvation is possible when *using methods as in the previous slides*
 - Even **Deadlock possible** if spinning in **non-preemptive priority-based scheduling**: example scenario: threads/processes on the same CPU and:
 - low priority thread LP executes in critical section
 - higher priority thread HP needs to access its CS
 - the higher priority thread HP is dispatched (i.e., occupies the processor) and busy-waits for the LP to exit from its critical section...

Critical section + bounded waiting with TaS + busy-waiting for n threads

1. **Shared var** *lock*: boolean // init false;
2. **Shared var** *waiting*[0..*n*-1]: array of boolean // init all false;
3. **do**
4. *waiting*[*i*] := TRUE;
5. **while** (*waiting*[*i*] && **TestAndSet**(&*lock*)) do [nothing]; //busywait
6. *waiting*[*i*] := FALSE;
7. [critical section]
8. *j* = (*i* + 1) % *n*;
9. **while** ((*j* != *i*) && !*waiting*[*j*]) do // find next one waiting, **to help it** (handover the “lock”)
10. *j* := (*j* + 1) % *n*;
11. **if** (*j* == *i*) **then**
12. *lock* := false; // completed one round without handing over; release “lock”
13. **else**
14. *waiting*[*j*] := false; // hand-over “lock”
15. [remainder section]
16. **forever**;

Agenda

- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques


Synchronization with semaphores

Special **variables/data-objects** used for **signaling**

- Accessible via (i.e., API) **atomic Wait and Signal operations**
- If a process is **waiting** for a signal, it is blocked until that **signal** is “sent”

Definition of the **signal()** operation

```
signal(S) {  
    S++; // and unblock from queue if applicable  
} // Executed atomically
```



Definition of the **wait()** operation

```
wait(S) {  
    while (S == 0) do; // here spin; alt. block in a queue  
    S--;  
} // Executed atomically
```

- **Binary semaphore** – value can be only 0 or 1
- **Counting semaphore** – value can vary

Critical section of n threads **using** semaphores

Shared **var** *mutex_sem* : binary semaphore // *init* = 1

Thread *i*

repeat

wait(mutex_sem);

[critical section]

signal(mutex_sem);

[remainder section]

forever

Semaphores as General Synchronization Tools

Q: how to **execute code segment B in T_B only after code segment A has been executed in T_A :**

A: use semaphore *flag*, initialized to 0

T_A	T_B
\vdots	\vdots
opA	<i>wait(flag)</i>
<i>signal(flag)</i>	opB

*T_B will be able to proceed from *wait(flag)* only after *signal(flag)* is executed by T_A*

Agenda

- Motivation
- The Critical Section problem
 - Check-competition approach
- The check-order / after-you approach
- Peterson's Algorithm
- Synchronization and hardware support
 - Read-Modify-Write instructions
- Semaphores
- Other techniques

Other synchronization operations/constructs:

Tools for programmers to solve critical section and synchronization problems (can be implemented using hardware atomic instructions or algorithms such as Peterson's or other such tools)

mutex lock `x`; API:


- `x.acquire()` wait/block if `x` is not available; else make `x` not available
- `x.release()` unblock a waiting/blocked process; else make `x` available
- If implemented with **busy waiting**, called **spinlock**

Condition Variables: `condition x`; API:

- `x.wait ()` – a process that invokes the operation is blocked.
- `x.signal ()` – unblocks one of processes that invoked `x.wait ()` (if any; else, does nothing)

Other high-level synchronization constructs

- **Monitors** (combination of mutex locks and condition variables)
- **Barriers**



3.: Global event synchronisation: an example of n threads synchronised at a barrier with corresponding arrival times and release time of all n threads when the barrier condition is fulfilled.

Linux Synchronization provides system calls for:

- semaphores
- spin locks for **multiprocessors** (for **short critical sections**): why?
 - [hint: wastes CPU time and can cause deadlocks if used with strict priority scheduling on uniprocessor]
- **futex**
 - *try_lock* through 1 invocation of RMW instruction; *if already locked, then process enters blocked queue; why?*
 - *[Balance the cost of entering a queue directly with the cost of spinning on the RMW instruction]*
 - <https://man7.org/linux/man-pages/man2/futex.2.html>

Windows XP Synchronization

- processor-specific **interrupt masks** to protect access to global resources on uniprocessors
- **busy-wait spinlocks** on multiprocessors
 - why not on uniprocessors?
 - [hint: can cause deadlocks if used with strict priority scheduling]
- **dispatcher objects**
 - may act as **mutexes**, **semaphores**, or provide **events** (similar to condition variables)

Pthreads Synchronization

OS-independent; provides:

- mutex **locks with blocking queues** (also fairness in mind)
- condition variables
- Non-portable extensions include:
 - **read-write locks** (we will discuss these in upcoming lectures)
 - **spin locks**

Thread call	Description
Pthread_mutex_init	Create a mutex
Pthread_mutex_destroy	Destroy an existing mutex
Pthread_mutex_lock	Acquire a lock or block
Pthread_mutex_trylock	Acquire a lock or fail
Pthread_mutex_unlock	Release a lock

Thread call	Description
Pthread_cond_init	Create a condition variable
Pthread_cond_destroy	Destroy a condition variable
Pthread_cond_wait	Block waiting for a signal
Pthread_cond_signal	Signal another thread and wake it up
Pthread_cond_broadcast	Signal multiple threads and wake all of them