

The Architectural Change Process

Josef Nedstam

Dept of Telecom, Lund University

visiting: National ICT Australia

josef.nedstam@telecom.lth.se

Even-André Karlsson

Q-Labs

Ideon, SE-223 70 Lund, Sweden

even-andre.karlsson@q-labs.se

Martin Höst

Lund University

Box 118, 221 00 Lund, Sweden

martin.host@telecom.lth.se

Abstract

Software architecture is recognized as a critical factor for successful products, but few have studied how organizations decide on architectural changes. In this paper we study the topic through several case studies. The changes are in all cases changes to the quality attributes of the system, and follow the same general process. We find that architectural changes have aspects of both functional and organizational changes. An architectural change does not only need to be technically sound, but it also needs to be anchored firmly in the organization. This paper describes the general architectural change process, and gives both architects and managers guidelines to balance short-term project goals and long-term organizational goals within this process.

1. Introduction

Software architecture is becoming a well-established field in technical terms, i.e. different types of architectures have been characterized [1]; different useful views of the architecture have been described [2, 3]; and books have been written, covering the general area, e.g. [4, 5, 6]. However, little research has been done on how decisions on architectural changes are made in organizations.

In this work we recognize that any change to a product can be classified anywhere on a continuous scale from solely architectural changes, affecting structure rather than user-observable attributes, to solely functional changes, limited to a single user function. Comparisons made between architectural and functional changes will mostly consider these two extremes, and the changes studied are considered to be close enough to the architectural end of this spectrum to support such arguments.

Architectural changes are often different in nature from functional changes. They can impact larger parts of the product, and they are expensive to implement. In this study we have also observed that architectural changes have a varying degree of organizational impact. Functional changes often originate from a customer demand and are the responsibility of a defined role in a company, i.e. product management. Architectural changes can on the other hand emerge from various sources, and roles are seldom defined to drive such changes.

The process for making decisions regarding functional changes and features has received attention in recent years [7]. Software development processes generally support this well, for instance through the stage-gate approach [8]. When it comes to decisions regarding the software architecture, the architect is often not so well supported, neither for the analysis of the technical solutions, nor for the organizational impact of the change.

The practitioners we cooperated with in this study needed help in managing architectural changes, help we could not provide from surveying literature. We therefore posed an initial research question: How are architectural decisions made? In order to answer this question we have qualitatively studied how several changes to the software architecture have been handled at three software development organizations, and what forces drive the need for changes and control which solutions are decided upon. The research question was therefore refined to: What is the implicit or explicit process of architectural change?

Based on the case studies, we identify such a general process. Detailed investigation of the cases, by comparing the identified process to the process for deciding on functional changes, and theories of organizational change, led to suggested improvements and guidelines for each step of the process. The validity of the process and guidelines has then been analyzed through a workshop session and through a further reference case study.

This introduction is followed by a method description of this study. Section 3 provides the case descriptions, and the suggested process is presented in Section 4, followed by a comparison between cases and process. Section 6 describes the validation, followed by conclusions.

2. Method

We have studied seven architectural changes initiated at three software-developing companies. All companies develop products for a mass-market, and their products have long lifetimes. Therefore their architectures need to support several simultaneous versions of their products, with several releases over an extended period of time.

The field of study, architectural changes, involves extended chains of events, where each chain of events occurs rarely. It is difficult to identify independent and dependent variables, and it is not possible to separate the

context from the instances of the phenomena under study. The research has therefore been carried out in a qualitative fashion [9], rather than in quantitative terms. The research was exploratory in nature, and the overall research question related to architectural decision making.

Qualitative data has been collected in two sets of interviews. The first set was held with architects and system designers at the three companies to collect information about the companies, their products, and their architecture. The interviews were open, and through the three interviews suitable architectural changes were identified. Changes were selected that affected varying aspects of the architectures to a small or large extent. We also had to constrict ourselves to cases that were recent in time, so key personnel and documentation could still be found.

Key persons in those changes were interviewed in a second set of interviews. This set of seven sessions, involving from two to five respondents, was guided by these interview questions:

1. What is the architectural change?
2. Why was the architectural change needed?
3. Who initiated it?
4. How was the associated decision made?

The data was then analyzed to find dimensions, classify and assign values to these, in order to draw qualitative conclusions from the data [9]. The significant dimensions that emerged were the phases of the resulting process, and collected data was categorized and tabulated accordingly. Contributing factors to the success or failure of the change initiatives were then identified, in order to provide the guidelines for the suggested process.

Some important threats to a qualitative research study are *description*, *interpretation*, and *theory* [10]. *Description* refers in this case to the risk of not having a valid description of what has been said at the interviews and during the presentations. Recording and transcribing the interviews, and receiving feedback on the original data from the respondents lowered this risk. *Interpretation* refers to the risk of imposing a preconceived pattern, instead of observing true patterns. By avoiding premature hypothesizing, this risk has been mitigated. *Theory* refers to the risk of not considering alternative patterns and conclusions. We have tried to handle this risk by being many (three) researchers and exposing our findings both to the respondents, and to practitioners in similar situations, described in Section 6.

3. Case Study Descriptions

The three companies involved in this study are industrial partners of the Center for Applied Software Engineering at Lund University (LUCAS). Part of LUCAS is the LUCAS Architecture Academy that is a one-year part time software architecture course for LUCAS partners. This

research is done based on issues that came up in the context of the course. The LUCAS Architecture Academy included a number of sessions where the companies presented their architectural work to each other, and issues in the area were raised and elaborated. This section describes architectural changes at the companies. These descriptions are followed up in Section 5, with a comparison to the suggested change process.

3.1 Company A

Company A develops control system environments for industrial automation, e.g. chemical plants, dairies, etc. The control system environment consists of both a development view, called control builder, and a deployment view, i.e. the controller itself. Within the control builder, controllers can be designed by specifying hardware sensors and actuators, constructing control loops, and connecting variables in those control loops to the hardware devices. A fully specified system can then be compiled and deployed onto a controller in a control system.

Company A typically carries out one large project at a time, involving the entire organization. Each project evolves the same product further by adding features to the control builder, e.g. new editor facilities, and the controller, e.g. new hardware interfaces. Implementation proposals are developed during a feasibility study. Accepted implementation proposals pass a tollgate, in line with the stage-gate approach [8], after which implementation begins. Development is organized into teams, each working on a number of implementation proposals. Work is feature-focused and the organization has no module-responsible and no architects, but instead relies on senior developers to take responsibility for long-term architectural goals. Two changes were studied at the company:

Protocol Framework: Company A recently acquired companies within their domain in order to increase their market share. The controller developed by Company A was intended to replace those companies' products. To support the same customers, the controller therefore had to support a number of legacy protocols for sensors and actuators from those products. This was realized as a problem using the present architecture, as the protocols were intertwined with the rest of the code, and could only be developed at one site, the one studied here. This site only had capacity to develop 1-2 new protocols per project. To be able to develop several protocols a year, Company A decided to develop a generic IO and communication protocol framework. The solution was developed through a pre-study and an implementation proposal, which resulted in a solution that enabled frequent releases of the product with many new or legacy protocols in each release. This would be accomplished by letting other de-

partments of the company develop the protocols they were responsible for, using the protocol framework.

Real-Time Operating System: For a number of years Company A had had discussions about cutting licensing costs on Real-Time Operating Systems (RTOS). A suggestion from local product management at the studied site to develop their own RTOS was rejected by local development management. In parallel, high-level management decided to reduce the number of RTOSs to only one. This would not only lower licensing costs but also provide focus on a common competency regarding RTOSs and tool support, which would standardize and simplify distributed development. Top-level development management initiated a pre-study across all departments of the company. Participants were interviewed regarding their use of, and competencies in RTOSs. The site studied here used one RTOS, but the pre-study led to a recommendation for all departments to switch to another. Eventually the recommendation became a requirement for a project at the studied site. This requirement was postponed by the local organization, while an OS expert prepared a solution with a Virtual Operating System (VOS) layer, which was introduced in a later project.

3.2 Company B

Company B develops platforms for consumer electronic devices. These platforms are sold to external customers who configure the services within the platform to create complete products. The software platform consists of a number of modules, and a middleware layer hides the internal architecture from the customers.

Projects are organized in: a project management group, with product management responsibility; a system-engineering group, with expert groups and function groups responsible for major features within market requirements; and a system realization group, which receives specifications from the system engineering group and develops the platform. The system realization group is divided into a hardware- and a software branch, which are subdivided into development teams, each responsible for a set of modules. The organization has defined roles responsible for each module. These persons work with function groups during specification and development teams during implementation. The company also has a dedicated architecture group that performs most of its work within projects, especially supporting and influencing the system-engineering group. Three changes were studied at the company:

Data Router: During routine reviews the system-engineering group discovered several modules handling data streams in similar ways. These modules could instead use a common data router and thereby save memory. The architecture group developed a design proposal

that was approved, but no resources were provided from the project. Project management did not consider the memory savings to be large enough. Therefore, the solution was implemented by the software architecture group and integrated with a small-scale system on an isolated branch of the code. After inspection this branch was merged with the main track, and the software architecture group initiated documentation and training on the new architectural mechanism. The solution was still not widely accepted, as most modules already had their own implementations of the same functionality.

Hardware Abstraction Layer Split: The bottom layer of the architecture had existed in previous versions of the product, but had not been formally defined, and therefore there had been no clear rules as to how to access the hardware. The hardware was also not encapsulated well enough from the software, leading to unnecessary impact in the software when the hardware changed. The developers working in the lower layers of the product realized the need for a clearer definition of these layers. They proposed a solution that meant removing hardware dependencies from the Hardware Abstraction Layer (HAL) interface, i.e. creating a logical layer on top of the previous HAL. One driving force for introducing this logical layer was that the cost for a product developed from the platform is very dependent on the hardware components used, and therefore these are often changed to provide cheaper solutions. The purpose of the logical layer is to allow such changes without expending effort in the higher layers of the software.

The solution was presented for the system-engineering group and brought to the software architecture group. When the proposed solution was established within the system-engineering and software architecture groups, project management decided to assign resources to the change. The software architecture group introduced new coding rules according to the suggestion and made changes to the architecture descriptions. At the same time, the developers in the HAL prepared by planning the change, before doing the actual implementation when resources were assigned and the architecture was updated.

Include-file Reorganization: The software architecture group had created a flexible structure for the source- and include-files. The design rules that enforced this structure required several files for each component, and when the number of modules grew to around 100, the configuration management system could not handle it any more. The persons responsible for tool support within Company B were in contact with support personnel from the tool supplier, who identified the problem as having too many files in the system. The software architecture group was assigned to create a new structure.

The flexibility provided by the original structure was only needed by a few of the about 100 modules, and these

could continue to use the previous structure. The rest of the modules were given a new structure, which basically involved merging three or four source files into one file. This resulted in a three-to-one reduction of source files.

3.3 Company C

Company C develops software engineering tools. One of their main products is a development tool that consists of a front-end with editors for various types of diagrams and source code, and a back-end for compiling the diagrams into code. Other utilities such as a simulation tool are also part of the development tool.

Company C releases a new version of their product every six months, and successive release cycles overlap. Features are implemented by development teams in an assembly-line fashion, described in [11, 12]. The organization has architects per project but no established line organization for architecture, and module responsibility is assigned to senior experts. Two changes were studied at the company:

Communication Mechanism: New requirements, especially related to new language standards, meant that the old architecture could not support further development. Top-level management therefore decided to create a new product generation. Company C had recently acquired other companies, which developed software engineering tools that were to be integrated into the new product. One of the problems with the new requirements was an increase in the number of diagram editors. The old communication mechanism did not support this increase, but one of the acquired companies had recently solved that problem, using a common object model. A technical discussion led to a consensus of using the new solution, although it meant major architectural changes.

Editor Framework: The editor framework utilized to develop graphical editors was changed as well, using a more generic solution, a decision also taken by consensus in the development project. The drivers for this change were increased reuse of common editor elements, and outsourcing of development throughout the organization. Several other decisions in this change process had to be enforced by the responsible architect, as consensus could not be reached. Both these changes were introduced in the same project.

4. Process Overview

Initial analysis of the cases led to a definition of a general process of architectural change, illustrated in Figure 1. The steps of this process were more or less present in all the seven studied changes, but there was often no conscious enactment of the process.

By analyzing the cases further, we identified problems and opportunities within the different companies, some of which are presented in Section 5. From this information, combined with theories of functional and organizational change, we were able to suggest improvements to this process, which is the form in which the process is described in this section. The relationship between the process and the case studies is shown in Section 5. The purpose of the improved process is to enable organizations to make the right decisions by the right people at the right time. From an employee viewpoint the process should give guidance in the decision process, both for change initiators and decision-makers. Since architectural changes have an impact on organizations they can be compared to the organizational change process. Kotter [13] defines an eight-stage process which describes how to prepare an organization for major change, and how to anchor the change in the organization:

1. Establishing a sense of urgency
2. Creating the guiding coalition
3. Developing a vision and strategy
4. Communicating the change vision
5. Empowering employees for broad-based action
6. Generating short-term wins
7. Consolidating gains and producing more change
8. Anchoring new approaches in the culture

The general process for functional changes involves requirements elicitation, pre-studies, implementation, and related decision-points. It focuses on how an organization should *make decisions*. Kotter's process for large-scale organizational change instead focuses on how to *make changes happen*. In the suggested process the two features are combined. This has basically been done by mapping Kotter's change process onto the functional change framework, such as described in the stage-gate approach [8]. In practice the process therefore has to be adapted to the present functional change framework. The process proposed in this paper contains these steps:

1. **A need emerges:** The process is superceded by a chain of events where need for change emerges or is created,

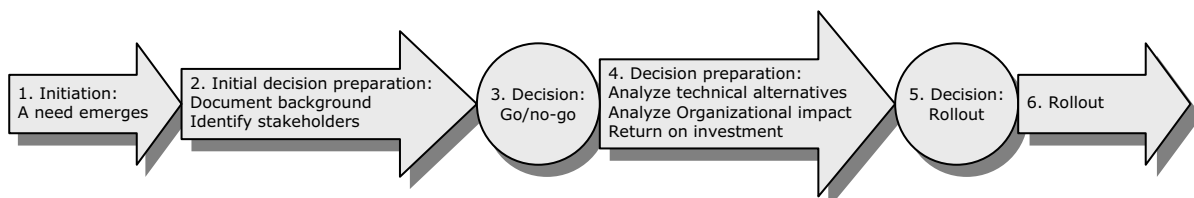


Figure 1. The Process of architectural change

and someone, the change initiator, sees this need and considers it his or her responsibility. This can to some extent be compared to Kotter's *Establishing a sense of urgency*, and to requirements elicitation in a functional change process.

2. **Initial decision preparation:** In this phase the change initiator does preparations with the goal of getting resources to analyze and implement the change.

- Document background: To increase the chance of having an impact on the resolution of the need, the change initiator should document the background of the need, i.e. what products, components or organizational entities are involved, the history behind the need, how it manifests itself, etc.
- Identify stakeholders/decision makers: While documenting the background, stakeholders should be identified. In order to have optimal impact, the change initiator should pay attention to these and especially to the decision makers that will be involved in the following process. This is related to Kotter's *Creating the guiding coalition*.

3. **Decision: Go/no-go:** An initial decision must be made whether the issue at hand is adequate and feasible to treat. Probably, not much effort has been spent before this decision point, e.g. one person's work for hours or days. Work done in the following phase of this process, before a decision on any particular solution or implementation of change, probably requires resources that must be budgeted, e.g. a handful of persons or more, which work for days or weeks. Therefore a person responsible for resources must make a decision whether to go on with this process or not. The formality of this decision-point is controlled by the organization at hand. If the change can be viewed as a normal product requirement or change proposal, it can be treated as such through the ordinary channels: implementation proposal and related decision points. However, if the change is more of a change in the way people work, or a change in an internal quality attribute not leading up to completion of a specific project, the process steps that follow are of a different complexity. The risks of facing opposition are higher and the decision process and preparations must be more thorough.

4. **Decision preparation:** This phase is akin to performing a pre-study or developing an implementation proposal in technical change management. In terms of Kotter's process, it resembles *Developing a vision and strategy*.

- Analyze technical alternatives: When technical alternatives have been proposed, these can be analyzed from an architectural viewpoint in a number of ways [14], i.e. ATAM [15].
- Analyze process and organization impact: When making a technical analysis, the organizational im-

plications are often overlooked, judging from the results of this study. This might lead to unexpected resistance to a change. An organizational analysis is therefore made, based on the initial analysis of stakeholders, in order to assess the impact of the change and prepare the organization for the change. The activity therefore contains parts of Kotter's *Communicating the change vision*.

- Return on investment (ROI): The need that the change satisfies has to have a financial side. An ROI analysis will simplify getting support for the change from top management and management of any project that might implement the change. This activity will support Kotter's *Generating short-term wins*.
5. **Decision: Rollout:** Software projects generally have a tollgate or decision point where it is decided which implementation proposals will be included in the resulting product. The same decision is made in this phase, regarding technical aspects of the architectural change. Organizational changes are however not suitable to implement in a product oriented project, and will therefore need another form of implementation and associated decision.
 6. **Rollout:** This activity involves the implementation of the change. Results from this study show that the rollout of the technical part of the change often is best carried out within an ordinary project, i.e. where most organizational resources often are allocated, but other options exist. The technical implementation has to be synchronized with the rollout of the organizational change, which must be managed by, and given resources from, the line organization. This activity is related to the late phases of Kotter's process: *Consolidating gains and producing more changes*, and *Anchoring new approaches in the culture*.

When comparing to Kotter's process it is important to keep the proper context in mind. Kotter presents a process for long-term organizational changes, which means some phases are of a different scale. Kotter's process also focuses on engaging employees and preparing an organization for a change, and not so much on how to perform the actual change. Since this paper focuses on changes to software architectures, we can use the decision framework common in software projects as a basis for a change process with features of both perspectives.

5. Analysis of Process versus Cases

This section compares the process suggested in Section 4 to the architectural changes described in Section 3. Table 1 gives an overview of the changes in relation to the suggested process, and also includes an architectural change used as a reference and presented in Section 6.

5.1 A Need Emerges

Before the suggested process is initiated a need for a change somehow appears. The reasons for changes in this study have included business decisions to increase market share, lower costs and lead time, but also more technical reasons where the architecture has not been able to support increased complexity and new features.

In Company A the need for the protocol framework was initiated when top management decided to increase the market-share by acquiring other actors in the same domain. Mid-level managers and experts then saw the need for support of legacy protocols found in the newly acquired companies' products. The need for a change of RTOS on local level came from a higher-level need to save licensing costs and focus competencies by reducing the number of RTOSs. The process was initiated by higher-level management and supported by developers at other sites of the organization.

In Company B the introduction of a data router was driven by memory size being an important quality attribute. The opportunity to save memory was discovered by system engineers during routine code-reviews. The need for a HAL split emerged as the company wanted to be able to change hardware components frequently in order to save costs. The developers in the lower levels of the

software initiated the change themselves in order to simplify the frequent changes. The need for an include-file restructuring became apparent, as the configuration management tool did not support the existing structure. The architecture group initiated this change since they were responsible for the include-file structure.

Two related changes were studied at Company C. A new mechanism, allowing editors to work against one system representation, was introduced to increase the number of possible editors. A framework for editor GUIs was introduced in order to increase reuse of common components and enable outsourcing of editor development. Local experts initiated these changes and the technology came from the newly acquired companies. These changes would later lead to a product generation shift.

Change initiators have been identified from all levels of the companies, i.e. managers, experts appointed when the issue came up or as part of their ordinary role, where a special case is the architects themselves, and down to the developers. This can be compared to functional changes where needs often emerge from customers and are managed by marketing or product management.

5.2 Initial Decision Preparation

A decision process that can be initiated by anyone will

Table 1. Overview of studied changes

	1. Initiation	2. Initial Decision Preparation	3. Decision: Go/no-go	4. Decision preparation	5. Decision: Rollout	6. Rollout	Organization impact
Protocol Framework	Need: Support legacy protocols – distribute protocol dev Initiator: Dev mgmt	Background: Effects of decisions not considered Stakeholders: No	Dev mgmt: develop implementation proposal	Tech: Pilot study Org: Improving during process ROI: No	Postpone parts	Part of project	Large, changes of responsibility between units
Real-time operating system switch	Need: Cut cost, standardize development Init: Top mgmt	BG: OK SH: OK	Top mgmt: requirement for next project	Tech: VOS developed on branch Org: Some ROI: No	Merge with current project	Parallel to project	Manageable changes to support contacts
Data Router	Need: Reuse to save memory Init: Sys eng	BG: OK SH: Emerged from established decision process	Sys eng: develop implementation proposal	Tech: Pre-study Org: No ROI: After the fact	No external resources	On isolated branch	Small, but developers not adopting other solutions
Hardware abstraction layer split	Need: Simplify change of HW comp Init: Developers	BG: Development before decision SH: OK	Action already taken	Tech: See 3. Org: Low impact ROI: Obvious	Package	Part of project	New product mgmt opportunities
Include file structure	Need: Re-establish tool support Init: Arch group	BG: OK SH: OK	Internal to architecture group	Tech: 2 structures, 3 rollout alt. Org: OK ROI: Obvious	Project by project	Project by project	Change of initial process steps
Communication mechanism	Need: Support several editors Init: Local experts	BG: Greater impact than documented SH: Some	Top mgmt: start new line of projects	Tech: Prototype project Org: No, ROI: No	Launch new series of projects	New series of projects	Change from line- to project-oriented dev
Editor framework	Need: Distribute dev Init: Local experts	BG: OK SH: OK	As above	As above	As above	As above	As above, and enabling distr
Ref change: Radar framework	Need: Reuse and CM Init: Developers	BG: Not explicitly SH: Not considered	No formal decision	Tech: Dev internal to project Org: No, ROI: No	Already implemented	Never rolled out externally	No existing org to manage new asset

eventually have to be brought before a decision maker. In this phase the change initiator documents background of the issue, and identifies stakeholders and decision makers.

When the need for legacy protocol support had emerged in Company A, local experts analyzed the protocol framework solution in a pre-study. Limited attention was however paid to other departments that were supposed to implement protocols on this framework. Regarding the change of RTOS, the pre-study had been carried out by higher-level management. This resulted in recommendations to change to a single and specified RTOS. The pre-study involved interviews on all company sites.

The introduction of a data router in Company B was prepared by the system-engineering group by marking the places where similar functionality had been found. Current and future clients to the data router were loosely identified but not further analyzed. The only stakeholder that was approached was the architecture group, who would be responsible for developing an implementation proposal. Regarding the HAL split, the developers in that layer prepared a solution themselves, and set up a meeting with the appropriate decision-makers, in this case the system-engineering group. In the case of the include-file restructuring, the initial preparation was made by the tool-vendor's support organization. They concluded that the projects contained too many files. The architecture group was identified as a stakeholder, since they had developed the previous structure. Apart from that, stakeholder identification was not done actively, since the frequent tool failures meant that stakeholders presented themselves.

In Company C the first steps of the product generation shift were taken on many levels, both within the original organization and by developers and managers in newly acquired organizations. Technical discussions were held which lead to the realization that the whole architecture had to be changed. Solutions were gathered from all parts of the organization, and the new architecture was adapted to enable distributed development. Stakeholders and decision-makers were therefore identified adequately.

In the case studies we have seen examples of less successful changes, where too little has been known about the impact of the change. Effects of functional changes are often more limited and customer-oriented. As opposed to architectural changes, functional changes often have resources allocated to this phase, such as product management performing requirements elicitation.

5.3 Decision Point: Go/No-Go

In this activity the first decision to commit resources is made. The right decision maker shall have been defined previously, and process and organizational issues must not be forgotten in this decision.

In Company A, local level management decided that an implementation proposal of the protocol framework should be developed, since the solution would allow for more protocols and more frequent releases of the product. The organizational impact was not given much focus in this decision. Regarding the RTOS switch, top management decided to turn the recommendation into a requirement for the following projects. This requirement was later postponed by the local organization.

In Company B, the system-engineering group decided that the architecture group should develop an implementation proposal of a data router. Regarding the HAL split, the solution was so well prepared by developers that neither the system-engineering group, nor the architecture group had to invest large resources in preparation, and the decision was of little significance. Regarding the include-file structure the architecture group themselves decided that they should develop a solution. Resources spent by the architecture group were considered insignificant in comparison to the resources wasted during tool problems. Organizational impact related to difficulties in rolling out the new structure was considered at this stage.

In Company C, the decision to apply resources to the change process was at a higher level, since it involved starting a whole new line of product-oriented projects. A decision was therefore made by top management to prepare and plan for a first project, which should result in a prototype for the product.

A forum for architecture issues could be helpful when making these decisions. Considering functional changes, organizations sometimes have product management fora, making similar decisions. The problem for the architect is that the decision is one of resources, for which the architect seldom has responsibility. Getting project resources has a benefit since the change can be more easily embraced by that project. It is however not trivial to receive resources from a project manager. The architect or change initiator therefore needs support in building a case for the proposed change.

5.4 Decision Preparation

In the decision preparation phase a small group of people will analyze technical alternatives, process and organizational impact, and ROI. From a company viewpoint this is done to make the right decision, and from an architect or change initiator viewpoint this will help convince people of the need for change. This phase is similar to developing an implementation proposal when making a functional change, and should therefore be adapted to how implementation proposals are handled within the organization. The analysis of technical alternatives can be done in parallel with the analysis of process and organizational impact.

At Company A, the protocol framework was prepared through an implementation proposal, in the same way as a normal requirement. The technical solution was based on expert opinions. The process and organizational impact was considered, and a pilot study was made which involved developing a protocol at another site in the same company. However, there are many developers in the acquired companies that are impacted by this change but have not been involved in these initial phases. The change of RTOS was postponed to a later project, and in the meantime an OS expert prepared a solution involving a VOS layer to allow for several operating systems. One organizational impact was overlooked, as the change meant that new RTOS support contacts had to be established. Regarding ROI, the change of RTOS led to no short-term wins for the local organization.

In Company B, the architecture group developed the data router solution in a pre-study. It was based on already implemented solutions, but the group failed to recognize opposition from project management and developers. ROI was calculated late in the process. Regarding the second change, the developers had already prepared the HAL split so the architecture group only had to prepare changes to architecture documentation and design rules. No quantitative ROI was made but the ability to change components was considered an obvious benefit. Regarding the include-file restructuring, the architecture group found that the flexibility provided by the original structure was only needed in a few modules, and a simpler structure was created for other modules. ROI, or cost, was considered regarding the rollout, since rollout was expensive and did not contribute directly to any product.

In Company C the first project of the new product generation was planned. When making technical decisions many parts of the organization were involved, and consensus in joint forums was the goal. When this could not be reached, the architect responsible for that type of functionality had to make the decision. Organizational impact was not only considered when selecting solutions, but also when distributing development of various modules. This distribution could at least in one case have been better planned, as they ended up with developing a module at one site, which was highly dependent on two other modules at another site, leading to unnecessary problems.

The studied cases have rarely included ROI or business cases for the changes. This might be because technicians often lack knowledge of such issues, and lack insight into the business model of the company. The solution would then be either training and education, or if technically oriented change initiators would receive support from marketing. Another issue in this phase is the involvement of the right people during decision preparation. Finding the right people would in many cases aid in identifying organizational impact at an early stage.

5.5 Decision Point: Rollout

When a feature-oriented implementation proposal is completed, it is generally passed through a tollgate in the project. In this tollgate the project decides which features or implementation proposals shall be included in the upcoming release. The activity described here is similar, but the changes we have studied have had organizational impact. Such changes, and their related decisions, are hard to make in a product-oriented project.

In Company A, the implementation proposal for the protocol framework involved two different types of protocols, and a set of services for these protocols. The decision to implement was made according to the standard project model. Both protocol types were to be implemented in the upcoming project, but part of the services were postponed to later projects. Regarding the change of RTOS, the expert's VOS solution was chosen, and local development management decided to roll it out onto a current project. This project had to start implementation before the VOS was ready, and therefore local development management decided that the VOS team would make relevant modifications of the project's code when the VOS was ready.

In Company B the system-engineering group approved the implementation proposal for the data router, but the architecture group did not receive project resources to implement the proposal. They then decided to implement the data router with their own resources. On the other hand, the HAL split was granted resources by project management, because it had backing from developers, system engineering, and the architecture group. The include-file restructuring was urgent, but difficult to roll out. An attempt to automate rollout failed because standards were not followed throughout the company. A second strategy was to halt development over a number of days, and perform the changes manually. This solution was too costly and eventually appropriate line management decided to roll the new structure out onto newly started projects, letting old projects use the old structure.

Company C decided to launch the series of projects for the new generation of products. Top management took this decision, and the content of each project has slowly been decided throughout the first projects by top management, product management and project management.

One conclusion from this activity is that it might be beneficial to restrict functional content of a new product when introducing major architectural changes. This was adequately done when introducing the IO and communication framework in Company A, as the number of services available to the protocols was restricted in the first release. Company C, however, has had problems deciding on the final content of the first product to be released on the market. Restriction of functional content is a tradeoff

since customers will not accept lower functional content, and the new architecture must be able to support future functional content. Another tradeoff regarding how many future features an architecture should enable concerns the debate of programming for the future or, as XP [16] advocates, programming only for the present.

5.6 Rollout

Implementation of technical aspects of changes is made successfully within product-oriented projects. Implementing technical aspects elsewhere is more problematic, since such implementations are not so easily embraced by developers in projects. The problem is that the process and organizational aspects are often forgotten in product-oriented projects, and there seldom exists a standard procedure for carrying out such changes, as opposed to carrying out a product-oriented project.

In Company A the protocol framework was implemented as part of a product-oriented project, but many departments that were intended to develop protocols have not yet had opportunity to give feedback on the framework. There is therefore still a risk that some departments will object to the framework. The VOS was developed in parallel with a product-oriented project. When the VOS was ready the two projects were merged, and the VOS team had to make remaining modifications.

In Company B the architecture group developed the data router on an isolated branch, which was later merged with the main branch. The problem was that most of the clients to the new data router had already implemented their own solutions, and usage of the router was only recommended, not required. Thus it did not provide the anticipated savings in memory footprint. The HAL split had been well prepared by both developers and architects before decisions were made, and was rolled out in a project. The new include-file structure was rolled out onto one project at a time across the organization. The rollout coincided with an architectural change, which minimized overhead when the key module responsible checked in the new file structure into the tool at project startup.

Company C has implemented their architectural changes in a prototype project, and a product for the market is under development. The main problems have been to settle on feature content, and as previously mentioned, the distribution of work. The company was not prepared for the organizational change they had to do in going from line-oriented development with frequent and regular releases, to project-oriented development.

In the case studies we have seen several examples of changes where the technical part has been assigned to a certain project as a requirement, but postponed to later projects. We have also seen examples where the changes have been performed outside of product-oriented projects,

further decreasing the chance of embracing the change. One of the cases made a satisfactory tradeoff, where the change of operating system was postponed to a later project, but prepared by an expert ahead of the project start.

6. Validation

In this paper we present two results: that there is a general process for architectural change; and that this process can be improved according to our guidelines, both in terms of value for a company, and the impact an architect can have. To study the generalizability of these qualitative results we have improved the validity in two ways.

A workshop was held with architects and system engineers from industry. Some of the companies from the initial study were present, along with some other companies, but none of the individuals were the same. One architectural change in progress was presented from each company, and it was possible to map the status of these changes onto the general change process depicted in Figure 1. However, some of the participants found it difficult doing the mapping by themselves. This still indicates that the presented process is indeed a general description of how an architecture change is performed. During the workshop, the complete process with guidelines and improvements was also presented, and the participants proposed that the concept of feedback should be introduced in the process, which is discussed in the conclusion.

In order to validate guidelines of the proposed process, a case study was made on an architectural change in a company that develops radar systems. This example is presented as the reference change found in Table 1, and shows when the suggested process is not applicable.

That company had developed three similar systems in succession and developers in the later projects decided to develop a framework for the radar tracking functionality, a functionality that was similar between the three different radar sensors found in the system. Due to long project lead-time, over 5 years, the developers could implement this framework without having to rely on higher level decisions. This level of liberty however also meant that acceptance of the framework in other parts of the organization has been very low, especially from management. Although opportunities for further reuse exist, there is no funding or organization to turn the project-internal framework into an organization-wide asset. The situation resembles the HAL split in Company B, where much implementation was done ahead of any decisions to assign resources. In that case the company decision process was however eventually followed. The proposed process of architectural change can therefore be said to fit architectures that are used over time and in many products, while changes to architectures of single products developed within single projects may take place in other ways. Fur-

ther evaluation is however needed in order to package and confirm the utility of the provided guidelines.

The type of systems studied here may also set bounds on the generalizability of these results. On the sliding scale from “sell and forget” products, to information systems based on a database, these systems are closer to the product end of the spectrum. The study has not included web-based systems, which often have an architecture tightly knitted to the middleware they are based on.

7. Conclusions

In the case studies we have seen that the need for architectural change can emerge from various sources, and that various roles, such as managers, architects and developers, may take responsibility for initiating the change. The decisions regarding architectural changes are often carried out in the same way as companies make decisions regarding functional changes, while the implementation of architectural changes may take many forms, such as part of ordinary projects, parallel but separate projects, independent smaller projects, or as new full-scale projects.

We have discovered three major differences between functional changes and architectural changes. First of all, architectural changes are often more complex than functional changes and affect large parts of the product without showing a clear connection to a customer need. Secondly, architectural changes do not only have impact across large parts of the product, but often across the whole organization, and changes of processes and organization are often overlooked and hard to implement in product-oriented projects. Finally, while companies often have mechanisms and resources in place to treat functional changes, such mechanisms are rarely established for architectural changes, and it is also hard to commit resources to activities without clear customer value.

The suggested process is admittedly simplified, but must be seen as a first and feasible step for companies that do not have processes in place for architectural evolution, a situation deemed common from studying the described cases. The simplicity of the process also keeps focus on its main benefits, namely identification of stakeholders and a defined sequence of decisions in order to gain management and organization-wide support for the initiative, and analysis of not only the technical, but also of the organizational implications of the initiative.

With increased architecture process maturity, the notion of feedback could provide further benefits from this architectural change process. In terms of the IDEAL model [19], this would involve adding a ‘Learning’ phase to the end of the suggested process, tying back to the first phases of the process for knowledge reuse in subsequent architectural changes. In the words of the SEI Software Product Line initiative, the suggested process could then

become a Software Product Line Practice Pattern [20], or rather an architecture management pattern.

Acknowledgement

This work was partly funded by The Swedish Agency for Innovation Systems (VINNOVA). National ICT Australia is funded through the Australian Government’s *Backing Australia’s Ability* initiative, in part through the Australian Research Council.

References

- [1] Garlan, D. and Shaw, M., *Software Architecture: Perspectives on an Emerging Discipline*, Prentice Hall, 1996
- [2] Soni, D., Nord, R., Hofmeister, C., “Software Architecture in Industrial Applications”, *ICSE*, Seattle, 1995, pp. 196-210
- [3] Kruchten, P., “The 4+1 Model View of Architecture”, *IEEE Software* 12 (6), IEEE, 1995, pp 42-50
- [4] Bass, L., Clements, P. and Kazman, R., *Software Architecture in Practice*, Addison-Wesley, 1998
- [5] Hofmeister, C., Nord, R. and Soni, D., *Applied Software Architecture*, Addison-Wesley, 1999
- [6] Bosch, J., *Design and Use of Software Architectures*, Addison Wesley, 2000
- [7] Evans, R., Park, S. and Alberts, H., “Decisions not Requirements: Decision-Centered Engineering of Computer-Based Systems”, *ECBS*, Monterey, 1997, pp. 435-442
- [8] Cooper, R. G., “Stage-Gate Systems: A New Tool for Managing New Products”, *Business Horizons*, May-June 1990, pp 44-55
- [9] Patton, M. Q., *Qualitative Evaluation and Research Methods*, 2nd Ed., Sage Publications, 1990
- [10] Robson, C., *Real World Research*, 2nd Ed., Blackwell Publishers Inc., 2002
- [11] Regnell, B., Berenmark, P. and Eklund, O., “A Market-Driven Requirements Engineering Process”, *Journal of Requirements Engineering* 3, Springer-Verlag, 1998, pp 121-129
- [12] Höst, M., Regnell, B., Natt och Dag, J., Nedstam, J., and Nyberg, C., “Exploring bottlenecks in market-driven requirements management processes with discrete event simulation”, *Journal of Systems and Software* 59, Elsevier Science Inc., 2001, pp 323-332
- [13] Kotter, J. P., *Leading Change*, HBS Press, 1996
- [14] Dobrica, L. and Niemelä, E., “A Survey on Software Architecture Analysis Methods”, *IEEE Transactions on Software Engineering* 28, IEEE, 2002, pp 638-653
- [15] Clements, P., Kazman, R. and Klein, M., *Evaluating Software Architectures*, Addison-Wesley, 2002
- [16] Beck, K., “Embracing Change with Extreme Programming”, *IEEE Computer*, October 1999, pp 70-77
- [17] Grant, R. M., *Contemporary Strategy Analysis*, Blackwell Publishers Inc., 1998
- [18] Cleland, D. I., and King, W. R. (eds.), *Project Management Handbook*, 2nd Ed., John Wiley & Sons, 1997
- [19] McFeely, R., *IDEAL: A User’s Guide for Software Process Improvement*, SEI/CMU, 1996
- [20] Clements, P., and Northrop, L., *Software Product Lines: Practices and Patterns*, Addison Wesley, 2001