


TDA357/DIT621 – Databases

Lecture 6 and 7 – Design using Functional Dependencies and normal forms

Jonas Duregård


Another high level design approach

- This week we will look at functional dependencies (FDs) and normal forms
- This is an alternative (and to some degree complementary) approach to ER that we studied last week
 - We start in a domain description and end in a database schema
- Two lectures, and Friday exercises as usual




Normalisation in a nutshell

- Extract a bunch of formal statements from the domain description
- Compute a normalised database schema from the formal statements



- Highly systematic, almost mechanical process
- By a carefully constructed normalization algorithm, the normal form the schema ends up in will satisfy some important properties

Functional Dependencies (FDs)

- A functional dependency is written as $\langle \text{set of attributes} \rangle \rightarrow \langle \text{attribute} \rangle$
- Example: $\text{room time} \rightarrow \text{course}$  Do not confuse with references!
- Pronounced "room and time determines course"
- It is a statement that can be true or false
- A few ways of understanding the meaning of the statement above:
 - If we know room and a time, we can uniquely determine course
 - There can be at most one course value for each (room,time)-pair
 - There exists a partial function f that takes a room and a time and yields a course
- In a domain it might have said something like "courses can book rooms at any free times"

Three ways we can use functional dependencies

- Check if they hold for a specific data set
 - Check if a design ensures they hold for all data sets
 - Express desired properties of a design
-
- I will explain each of these in turn

Functional dependency as a property of data

- One way of formally defining functional dependency $x_1 x_2 \dots \rightarrow y$:
For $R(y x_1 x_2 \dots)$, if two rows agree on $x_1, x_2 \dots$ they must also agree on y
- In other words, there cannot exist two rows where the left-hand side attributes are the same, but the right-hand side attribute differs
 - " $X \rightarrow y$ = rows that agree on X must agree on y "

Two rows "agreeing on x " just means the x -column(s) have the same value

- Which FDs hold for this data?

Table: Bookings

courseCode	name	day	timeslot	room	seats
TDA357	Databases	Tuesday	0	GD	236
TDA357	Databases	Tuesday	1	GD	236
ERE033	Reglerteknik	Tuesday	0	HB4	224
ERE033	Reglerteknik	Friday	0	GD	236

courseCode → name?

- Yes! (TDA357 maps to Databases, ERE033 to Reglerteknik)

day → timeslot?

- No! (Tuesday maps to both 0 and 1)

day timeslot room → courseCode?

- Yes! (There are no rows where all three LHS columns match)

seats → room?

- Yes! 236 for GD, 224 for HB4,
- This might not be intentional given what we know of the domain...

LHS = Left-hand side (of arrow)
RHS = Right-hand side

FDs on small tables

- Relation $R(a,b,c)$ has a single row. Does $a \rightarrow b$ hold for that table?
 - Yes! In fact, all FDs hold for tables with fewer than two rows.
- Consider: For an FD not to hold on a relation, there needs to be two conflicting rows.
- More general: For fd $x_1 x_2 \dots x_n \rightarrow y$, there needs to be two rows that agree on all the x_i values (but not on y) for the FD to not hold.

Thinking about attributes in FD analysis

- Consider this relational schema with courses and teachers:

Teachers(*email*, *tname*)

Courses(*code*, *cname*, *teacher*)

teacher \rightarrow *Teachers.email*

- The domain only has four attributes, teacher emails and names and course codes and names.
- The teacher attribute of Courses and email attribute of Teachers are two names for the same attribute, not two independent attributes

FDs as a properties of designs

- Knowing what it means for an FD to hold for a data set, we can determine if a design (schema) guarantees that it holds for all valid data sets

- Example: Does the schema below guarantee that ...

code \rightarrow cname?

- Yes! (by primary key constraint in courses)

cname \rightarrow code?

- No! (Counterexample: any two courses with the same name)

code \rightarrow email

- Yes! (teacher is just another name for email)

code \rightarrow tname

- Yes! (by primary key + reference)

```
Teachers(email, tname)  
Courses(code, cname, teacher)  
teacher  $\rightarrow$  Teachers.email
```

Bookings (*courseCode*, *name*, *day*, *timeslot*, *room*, *seats*)
(*day*, *timeslot*, *room*) **UNIQUE**

- Does the schema above guarantee ...
- $\text{day timeslot room} \rightarrow \text{courseCode}$
 - Yes (through UNIQUE constraint)
- $\text{day timeslot room courseCode} \rightarrow \text{seats}$
 - Yes (through primary key and/or UNIQUE)
- $\text{room} \rightarrow \text{seats}$
 - No ☹️
- $\text{courseCode} \rightarrow \text{name}$
 - No ☹️

Counterexample of $\text{room} \rightarrow \text{seats}$ and $\text{courseCode} \rightarrow \text{name}$

<u>courseCode</u>	name	<u>day</u>	<u>timeslot</u>	room	seats
CC1	N1	Tuesday	0	R1	0
CC1	N2	Tuesday	1	R1	1

Different timeslot, so no key violation

Creating counterexamples

- A good exercise (and exam question) for FDs is to given a schema with constraints and one or several FDs, show that the FDs are not guaranteed by the schema
- We need to construct a relation that has:
 - For each FD at least two rows that agree on its left hand side
 - ... but disagree on its right hand side
 - ... and respects the constraints of the schema
- We'll need at least two rows, possibly more

- Given the relation $R(\underline{a}, \underline{b}, c)$, construct a counterexample to both $a \rightarrow c$ and $b \rightarrow c$

- This doesn't work because it violates the key constraint:

a1	b1	c1
a1	b1	c2

- This doesn't work because $b \rightarrow c$ holds:

a1	b1	c1
a1	b2	c2

- This is a valid counterexample:

a1	b1	c1
a1	b2	c2
a2	b2	c1

a1 related to both c1 and c2, same for b2

No two rows have the same value for both a and b
--

So far we can ...

- Verify/disprove that an FD holds for a specific data set
- Verify/disprove that an FD always hold for all valid data in a schema

FDs as intention for designs

- Since we can verify that an FD holds for a schema, we can also use them to specify desired properties of our schema
- This is what makes FDs a design tool
- A sentence like "every course has a teacher" can be modelled as the FD $\text{course} \rightarrow \text{teacher}$ (or whatever attributes we use)
 - If this FD does not hold for our design, maybe the design is bad?

Side note: Single/Multiple FDs

Notation: I use lowercase $x/y/z$ for single attributes and uppercase $X/Y/Z$ for attribute sets

- It is common to have multiple attributes on the right-hand side of FDs
 $x\ y\ z \rightarrow a\ b\ c$
- This means exactly the same as these three FDs:
 $x\ y\ z \rightarrow a$
 $x\ y\ z \rightarrow b$
 $x\ y\ z \rightarrow c$
- I find it most useful to think of the first as a convenient way of writing multiple FDs, rather than thinking of it as single FD with multiple attributes
- It is not the same with the left-hand side! $x\ y \rightarrow a$ does not mean $x \rightarrow a$!

Formal properties of FDs

Warning: Things may get slightly mathsy from this point

- FDs have lots of interesting mathematical properties
- I will explain some of the more useful ones:
 - Transitivity
 - Augmentation
 - Reflexivity
- These three are commonly referred to as Armstrong's axioms and they can be formulated in a few different but equivalent ways*

*But the way I formulate them is -of course- the best way

Recall: $X \rightarrow y$ = rows that agree on X must agree on y

Transitivity of functional dependencies

- Functional dependency is a transitive relation
 - This means that if $X \rightarrow Y$ and $Y \rightarrow Z$ then $X \rightarrow Z$
- Note that Y is an attribute set here, so $X \rightarrow Y$ may be multiple FDs with the same LHS
- Proof sketch: Look at any rows that agree on X . Since $X \rightarrow Y$, they must also agree on Y , and since $Y \rightarrow Z$ they must further agree on Z . Thus $X \rightarrow Z$.

Recall: $X \rightarrow y$ = rows that agree on X must agree on y

Augmentation

- If $x_1 x_2 \dots \rightarrow y$, then for all z :
 $z x_1 x_2 \dots \rightarrow y$
- Intuitively: You can add any attributes you want to the LHS of a valid FD and still get a valid FD
 - Think: "knowing an extra attribute never prevent us from finding y "
- Proof sketch: Since all rows that agree on the x s must agree on y , then particularly all rows that agree on z as well as the x s must do so.

Recall: $X \rightarrow y$ = rows that agree on X must agree on y

Reflexivity and trivial FDs

- For all x : $x \rightarrow x$
(x determines itself)
- By augmentation, $X \rightarrow y$ whenever $y \in X$
 - Example: $a b c \rightarrow b$
 - We call these dependencies trivial
 - Rule of thumb: Ignore trivial dependencies
- Proof sketch: Any values that agree on x will agree on x 😊

Example: Deriving functional dependencies

- For any attributes x, y, z, w , and q :
 $x \rightarrow y, z \rightarrow w$, and $y w \rightarrow q$ implies $x z \rightarrow q$
- Proof:
 - $x \rightarrow y$ implies $x z \rightarrow y$ (by augmentation)
 - $z \rightarrow w$ implies $x z \rightarrow w$ (by augmentation)
 - $x z \rightarrow y w$ and $y w \rightarrow q$ implies $x z \rightarrow q$ (by transitivity)
- Note that in the third step we merge $x z \rightarrow y$ and $x z \rightarrow w$ into $x z \rightarrow y w$
(See slide on Single/Multiple FDs)

$$\begin{array}{l} x \rightarrow y \\ z \rightarrow w \\ y w \rightarrow q \\ \hline x z \rightarrow q \end{array}$$

Minimal basis (a.k.a. minimal cover)

- A minimal basis F^- of a set of functional dependencies F is a set equivalent to F but with the following conditions:
 - F^- has no trivial dependencies
 - No dependency in F^- follows from other dependencies in F^- through transitivity or augmentation
- Used for a lot of algorithms and to express a set of FDs in a compact form

Minimal basis

- Suppose we are given this set of FDs (5 total), what is a minimal basis?

$a \rightarrow b$

$b \rightarrow c$

$a d \rightarrow \cancel{b} \cancel{c} \cancel{d}$

- $a d \rightarrow d$ is removed because it is trivial
- $a d \rightarrow b$ is removed because it is implied by $a \rightarrow b$ (augmentation)
- $a d \rightarrow c$ is removed because it is implied by $a \rightarrow b$ and $b \rightarrow c$ (transitivity and augmentation)
- Final set: $a \rightarrow b, b \rightarrow c$

Transitive closure

- The transitive closure X^+ of a set of attributes X , is the set of attributes that can be functionally determined by X
- In other words, $X^+ =$ All attributes y such that $X \rightarrow y$
 - Includes ALL derived functional dependencies
 - Includes trivial dependencies
 - X^+ is closed in the sense that attributes in X^+ never determine attributes outside X^+
- Can be computed by a simple algorithm from any set of FDs:
 - Start with $X^+ = X$ (an under-approximation)
 - Repeat until done: For any FD $Y \rightarrow z$ such that $Y \subseteq X^+$, add z to X^+

Transitive closure, example


$x \rightarrow y$
 $y w \rightarrow q$
 $z \rightarrow w$
 $q \rightarrow x$
 $q r \rightarrow s$

- Given these FDs, compute the closure $\{x,z\}^+$
- Initially we know $\{x,z\} \subseteq \{x,z\}^+$ (from trivial FDs)
- Add y because $x \rightarrow y$ and $\{x\} \subseteq \{x,z\}^+$ $\{x,z,y\} \subseteq \{x,z\}^+$
- Add w because $z \rightarrow w$ and $\{z\} \subseteq \{x,z\}^+$ $\{x,z,y,w\} \subseteq \{x,z\}^+$
- Add q because $y w \rightarrow q$ and $\{y,w\} \subseteq \{x,z\}^+$ $\{x,z,y,w,q\} \subseteq \{x,z\}^+$
- No more FDs add attributes, so $\{x,z\}^+ = \{x,z,y,w,q\}$ is our result
- This proves all these non-trivial FDs:

$x z \rightarrow y$ $x z \rightarrow w$ $x z \rightarrow q$

Closure $\{x,z\}^+$, visually

Red arrow indicates a "non-closed" FD




$x \rightarrow y$
 $y w \rightarrow q$
 $z \rightarrow w$
 $q \rightarrow x$
 $q r \rightarrow s$

$q r \rightarrow s$ is irrelevant
since r is not in $\{x,z\}^+$

$$\{x,z\}^+ = \{x,z,y,w,q\}$$

What it means to be closed

- Study the set $\{x,z\}^+ = \{x,z,y,w,q\}$
- If we "follow the arrow" for any FD on attributes, we end up back in $\{x,z\}^+$.



$x \rightarrow y$
 $y w \rightarrow q$
 $z \rightarrow w$
 $q \rightarrow x$
 $q r \rightarrow s$

Keys and superkeys

- We can define the property of being a key of a relation using FDs
- Intuitively: A set of attributes is a *superkey* if it determines all other attributes
- Formally: The attribute set X is a *superkey* of R if X^+ contains all attributes of R
- X is a (minimal) key if removing any attribute from X makes it a non-superkey
 - Saying only "key" usually means minimal key
 - Each superkey is a superset of at least one minimal key
 - Each key is a superkey (but not the other way around)
 - Adding any attribute to a superkey makes a new superkey

Finding a minimal key

- Find a minimal key for $R(x,y,z,w,q,r,s)$ with the listed FDs
- Start with $K=\{x,y,z,w,q,r\}$, it's trivially a superkey ($K^+=R$)
- Remove one element from K at a time, and compute the closure of K , if it is no longer R , return the removed element
- When no element can be removed, we have a minimal key
- $\{x,z,r\}$ is one such key, since $\{x,z,r\}^+ = R$ and neither $\{x,z\}^+$, $\{x,r\}^+$ nor $\{z,r\}^+$ are equal to R
- $\{q,z,r\}$ is another minimal key, and possible result of the procedure


$x \rightarrow y$
$y w \rightarrow q$
$z \rightarrow w$
$q \rightarrow x$
$r \rightarrow s$

Clever tricks for finding multiple keys

- We know r and z will be in ALL minimal keys, because no FD determines them
(imagine removing r from K , how is it still in K^+ ?)
- This means w and s will be in NO minimal keys
(they are in $\{r,z\}^+$ so in K^+ for any K s.t. $\{r,z\} \subseteq K$)

$$\begin{array}{l} x \rightarrow y \\ yw \rightarrow q \\ z \rightarrow w \\ q \rightarrow x \\ r \rightarrow s \end{array}$$

Normal forms and normalization



Normal forms and normalization

- Normal form is a very important concept in database design
- Identify all the attributes in the domain and place them in one big relation $D(x, y, z, \dots)$, collect FDs, then normalize D to get your design
- Normalizing is a recursive procedure, to normalize relation R :
 - Check if R is already a normal form, if it is we are done
 - Otherwise decompose R into relations R_1 and R_2 and normalize both
- Note: A normal form is not the same as a canonical form, there may be multiple normal forms derived from the same initial domain

BCNF, the Boyce-Codd Normal Form

Probably the most well-known normal form

BCNF Normalisation algorithm

To normalize relation R:

This FD is referred to as a BCNF-violation



Find a non-trivial FD $X \rightarrow y$ such that $X^+ \neq R$ (X is not a superkey)

If there is no such FD you are done, R is already in BCNF

Otherwise decompose R into $R_1(X^+)$ and $R_2(X \cup (R - X^+))$ and normalize them

Note: R is replaced by R_1 and R_2 (so R is not present in the final schema)

Example

1. Find violation
2. Decompose
3. Repeat

courseCode \rightarrow name
room \rightarrow seats
day timeslot courseCode \rightarrow room
day timeslot room \rightarrow courseCode

- Normalise this relation using the FDs above:

~~R~~(courseCode, name, day, timeslot, room, seats)

$R_1(X^+)$

$R_1(\text{courseCode}, \text{name})$

$R_2(X \cup (R - X^+))$

~~R~~₂(courseCode, day, timeslot, room, seats)

Decompose on courseCode \rightarrow name
courseCode⁺ = {courseCode, name}

$R_{21}(X^+)$

$R_{21}(\text{room}, \text{seats})$

$R_{22}(X \cup (R - X^+))$

$R_{22}(\text{courseCode}, \text{day}, \text{timeslot}, \text{room})$

Decompose on room \rightarrow seats
room⁺ = {room, seats}

All of R_1 , R_{21} , and R_{22} are now BCNF!

Wait, why not split on day timeslot course \rightarrow room?

$R_{22}(\text{courseCode}, \text{day}, \text{timeslot}, \text{room})$

day timeslot courseCode \rightarrow room

day timeslot room \rightarrow courseCode

Recall: Find a non-trivial FD $X \rightarrow y$ such that $X^+ \neq R$ (X is not a superkey)

$\{\text{day}, \text{timeslot}, \text{courseCode}\}^+ = \{\text{day}, \text{timeslot}, \text{courseCode}, \text{room}\} = R_{22}$

$\{\text{day}, \text{timeslot}, \text{room}\}^+ = \{\text{day}, \text{timeslot}, \text{room}, \text{courseCode}\} = R_{22}$

Both $\{\text{day}, \text{timeslot}, \text{courseCode}\}$ and $\{\text{day}, \text{timeslot}, \text{room}\}$ are keys!

What about keys?

courseCode \rightarrow name
room \rightarrow seats
day timeslot courseCode \rightarrow room
day timeslot room \rightarrow courseCode

- Keys can be determined using FDs (and closures) after decomposing
- Much of it is already done as part of the algorithm (we found two keys for R_{22} for instance)

R_1 (courseCode, name)

R_{21} (room, seats)

R_{22} (courseCode, day, timeslot, room)
(day, timeslot, room) UNIQUE



Multiple keys: Use one as primary key, the other(s) UNIQUE

What about references?

- In this case it's fairly easy to see that these are sensible references:

$R_1(\underline{\text{courseCode}}, \text{name})$

$R_{21}(\underline{\text{room}}, \text{seats})$

$R_{22}(\underline{\text{courseCode}}, \underline{\text{day}}, \underline{\text{timeslot}}, \text{room})$

$(\text{day}, \text{timeslot}, \text{room}) \text{ UNIQUE}$

$\text{courseCode} \rightarrow R_1.\text{courseCode}$

$\text{room} \rightarrow R_{21}.\text{room}$

- General pattern: When decomposing R , add a reference $X \rightarrow R_1.X$ to R_2
- This will not always work, particularly if R_1 or R_2 is later decomposed ☹

Decomposition of data

No redundancy, no anomalies 😊

Table: Bookings

courseCode	name	day	timeslot	room	seats
TDA357	Databases	Tuesday	0	GD	236
TDA357	Databases	Tuesday	1	GD	236
ERE033	Reglerteknik	Tuesday	0	HB4	224
ERE033	Reglerteknik	Friday	0	GD	236

Table: R₁ (Courses)

<u>courseCode</u>	name
TDA357	Databases
ERE033	Reglerteknik

Table: R₂₂ (Bookings)

<u>courseCode</u>	<u>day</u>	<u>timeslot</u>	room
TDA357	Tuesday	0	GD
TDA357	Tuesday	1	GD
ERE033	Tuesday	0	HB4
ERE033	Friday	0	GD

Table R₂₁ (Rooms)

<u>room</u>	seats
HB4	224
GD	236

Lossless join

- Note that if we join along the references, we get the original table

Table: R_1 (Courses)

courseCode	name
TDA357	Databases
ERE033	Reglerteknik

Table: R_{22} (Bookings)

courseCode	day	timeslot	room
TDA357	Tuesday	0	GD
TDA357	Tuesday	1	GD
ERE033	Tuesday	0	HB4
ERE033	Friday	0	GD

Table R_{21} (Rooms)

room	seats
HB4	224
GD	236

Joins USING (courseCode) and USING (room)

Query: R_1 NATURAL JOIN R_{22} NATURAL JOIN R_{21}

courseCode	name	day	timeslot	room	seats
TDA357	Databases	Tuesday	0	GD	236
TDA357	Databases	Tuesday	1	GD	236
ERE033	Reglerteknik	Tuesday	0	HB4	224
ERE033	Reglerteknik	Friday	0	GD	236

- Means we did not loose any data in the decomposition

Another example

- Normalizing $R(a,b,c,d,e)$
- Decompose R on $a b \rightarrow c$, $\{a,b\}^+ = \{a,b,c,d\}$
 - $R_1(a,b,c,d)$
 $R_2(a,b,e)$
- Further normalize R_1 on $c \rightarrow d$, $\{c\}^+ = \{c,d\}$
 - $R_{11}(c,d)$
 $R_{12}(a,b,c)$
- End result with keys:
 - $R_{11}(\underline{c}, d)$
 $R_{12}(\underline{a}, \underline{b}, c)$
 $R_2(\underline{a}, \underline{b}, \underline{e})$

$a b \rightarrow c$
 $c \rightarrow d$

Note: Not just (a,b,c) !

Try splitting on $c \rightarrow d$ first, for this example that should give the same result

Finding all FDs

- Consider this simple situation with four attributes $R(x,y,z,w)$ and two functional dependencies: $x \rightarrow z$ and $y z \rightarrow w$
- When normalizing R it may be important to know that there is another FD that can be derived from these: $x y \rightarrow w$
- In principle, you should consider all non-trivial derived FDs but sometimes this a large set and it is easy to miss FDs
 - Essentially you have to consider every LHS and compute closures

Example where derived FDs matter

- $R(a,b,c,d,e)$
- Split on $a \rightarrow b$:

$a \rightarrow b$
$b c \rightarrow d$

$R_1(a,b)$

$R_2(a,c,d,e)$

- Now $a c \rightarrow d$ is a BCNF violation in R_2 !
- We cannot entirely ignore $b c \rightarrow d$ even though b is not in R_2 .

A flaw of BCNF

- Consider this example (again): $R(x,y,z,w)$ where $x \rightarrow z$ and $y z \rightarrow w$
- If we decompose on $x \rightarrow z$ ($\{x\}^+ = \{x,z\}$) we get
 - $R_1(\underline{x},z)$ $\{x\}$ is the only key
 - $R_2(\underline{x},\underline{y},w)$ $\{x,y\}$ is the only key
- Both of these relations are in BCNF w.r.t. the given FDs
- But now $y z \rightarrow w$ is not guaranteed by the schema ☹
 - Exercise: Construct a counterexample
- There is a weaker normal form called third normal form (3NF) that does not have this problem, but it has other problems instead...
 - There is no "silver bullet" for design work

Because $x y \rightarrow w$



Yet another issue with BCNF

- This relation has no non-trivial functional dependencies, so is in BCNF:

Table: Courses

course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Reglerteknik	RTB 1	AuthorX	TeacherX
Reglerteknik	RTB 2	AuthorX	TeacherX

Deletion anomaly: Deleting all course books also deletes all teachers

Update anomaly: Changing some value can cause inconsistencies

- The domain said something like "each course has a number of teachers and a number of books with one or more authors" (no FDs at all!)
- The data above says Databases has one book and two teachers, and Reglerteknik has two books and one teacher
- Clearly there is redundancy here, and potential for anomalies

Looks like we need another normal form!

- This one is called the fourth normal form (4NF)
- Since the problematic table had no FDs at all, this form will need some additional source of facts
- We call these facts multivalued dependencies (MVDs)*
- We write $x_1 x_2 x_3 \dots \twoheadrightarrow y_1 y_2 y_3 \dots$
 - Note that both sides are sets of values and we cannot split the RHS

*The term multivalued dependency is really quite unfortunate, but it is what it is

Multivalued dependencies, informally

- For our example, we would say course \twoheadrightarrow teacher
- This means that each course value has a set of teacher values that is independent from all other values (author and book)

Table: Courses

course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Databases	DTCB	Widom	Jonas
Databases	DTCB	Widom	Ana
Reglerteknik	RTB 1	AuthorX	TeacherX
Reglerteknik	RTB 2	AuthorX	TeacherX

Databases has (independently):
One book with two authors
Two teachers

Reglerteknik has two books by the
same author and one teacher

- This is exactly the same as saying course \twoheadrightarrow book author

Multivalued dependencies, formally

- The claim that $X \twoheadrightarrow Y$ holds for relation R means:

For every pair of rows row t and u in R that agree on X we can find a row v s.t:

v agrees with both t and u on X

v agrees with t on Y

v agrees with u on $R - X - Y$ (all attributes not in the MVD)

- Example: $\text{course} \twoheadrightarrow \text{teacher}$

If we remove any row, the MVD won't hold

Table: Courses

course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Databases	DTCB	Widom	Jonas
Databases	DTCB	Widom	Ana

Row t

Row u

Row v:

$v.(book, author) = u.(book, author)$
 $v.teacher = t.teacher$

MVDs, informally

- course \twoheadrightarrow teacher means that every course has a list of teachers that is independent from its list of (book,author)-pairs
- Jonas and Ana are the teachers here, and if Jonas occurs with the row (DTCB, Ullman), Ana must do so as well.

Table: Courses

course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Databases	DTCB	Widom	Jonas
Databases	DTCB	Widom	Ana

If we remove this row, Jonas has a different book list than Ana, and the MVD does not hold

Verifying MVDs on data is hard

- To check if an FD holds: Just group values up by the LHS and check that all rows in each group have the same value for the RHS
- To check if an MVD holds: Check every individual pair of values with identical LHS and search for a row with correct values
- I find a more intuitive way of thinking is this: For $X \twoheadrightarrow Y$, every X needs to have every possible combination of Y and other attributes (R-X-Y)
 - Essentially the rows for a given X must be a cartesian product!
 - If teacher Jonas occurs with one book/autor, it must occur with all book/author combinations for that course
 - This is what makes (book,author) independent from teacher

MVDs and cartesian product

- Note how the rows for databases is the cartesian product of these sets:
 - {(DTCB, Ullman), (DTCB, Widom)}
 - {Jonas, Ana}
- This means that the MVD $\text{course} \twoheadrightarrow \text{teacher}$ holds for these rows (and equivalently, $\text{course} \twoheadrightarrow \text{book author}$)

Table: Courses

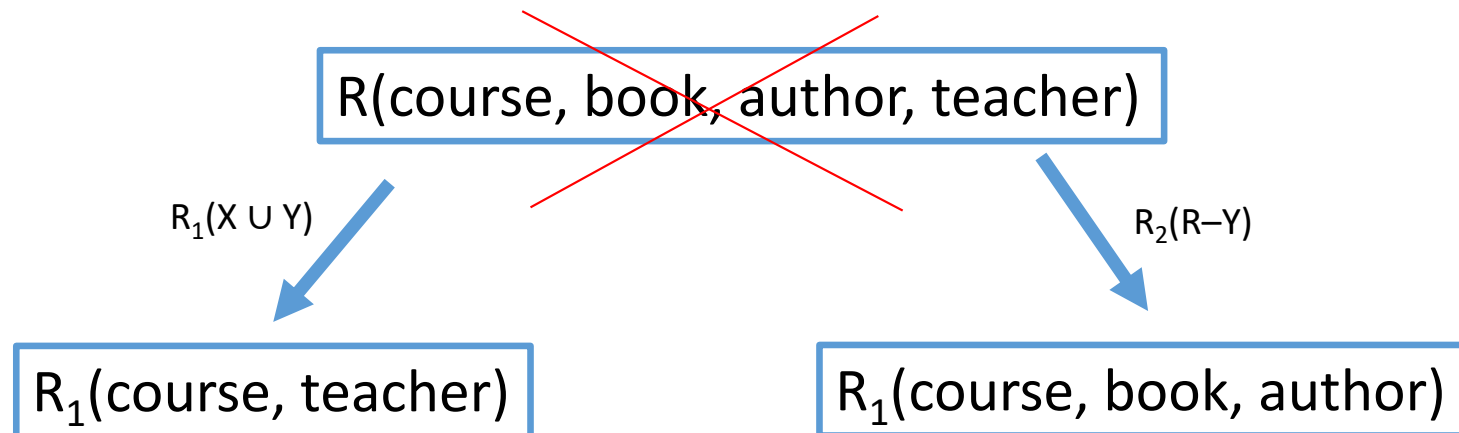
course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Databases	DTCB	Widom	Jonas
Databases	DTCB	Widom	Ana

Fourth normal form

- For a relation R to be in fourth normal:
 - R must be in BCNF
 - For all non-trivial MVDs $X \twoheadrightarrow Y$ on R , X is a superkey of R
- If $X \twoheadrightarrow Y$ and X is not a superkey, we say $X \twoheadrightarrow Y$ is a 4NF violation
- To normalize: Find a violation $X \twoheadrightarrow Y$ and break R into
 - $R_1(X \cup Y)$ ("every attribute in the MVD")
 - $R_2(R - Y)$ ("LHS and every attribute not in the MVD")
 - Then normalize both R_1 and R_2

4NF normalisation

- Normalizing $R(\text{course, book, author, teacher})$ on $\text{course} \twoheadrightarrow \text{teacher}$
 - Here $X=\{\text{course}\}$, $Y=\{\text{teacher}\}$



Normalising the data

Table: Courses

course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Databases	DTCB	Widom	Jonas
Databases	DTCB	Widom	Ana
Reglerteknik	RTB 1	AuthorX	TeacherX
Reglerteknik	RTB 2	AuthorX	TeacherX

Exercise: Find another MVD here?

Table: R₁ (a.k.a. CourseTeacher)

course	teacher
Databases	Jonas
Databases	Ana
Reglerteknik	TeacherX

Table: R₂ (a.k.a. CourseBooks)

course	book	author
Databases	DTCB	Ullman
Databases	DTCB	Widom
Reglerteknik	RTB 1	AuthorX
Reglerteknik	RTB 2	AuthorX

Lossless join

Note that if we join the two tables using course ...

Table: R_1 (a.k.a. CourseTeacher)

course	teacher
Databases	Jonas
Databases	Ana
Reglerteknik	TeacherX

Table: R_2 (a.k.a. CourseBooks)

course	book	author
Databases	DTCB	Ullman
Databases	DTCB	Widom
Reglerteknik	RTB 1	AuthorX
Reglerteknik	RTB 2	AuthorX

 NATURAL JOIN

We get the original table back!

course	book	author	teacher
Databases	DTCB	Ullman	Jonas
Databases	DTCB	Ullman	Ana
Databases	DTCB	Widom	Jonas
Databases	DTCB	Widom	Ana
Reglerteknik	RTB 1	AuthorX	TeacherX
Reglerteknik	RTB 2	AuthorX	TeacherX

Sanity check:
We did not loose
any information


Functional dependencies vs. ER-design

- FDs can find some things that ER cannot find
- ER can find a lot of things that FDs cannot find
 - Most many-to-many relationships cannot be expressed using FDs
 - Sentences like "students can register for courses" do not express any FDs (but possibly some MVDs?)
- The two approaches complement each other, and confirm each other (or sometimes contradict each other which may indicate a problem)
- So doing both an ER-design and a FD analysis may be useful
 - This is what you will do in Task 2


Expressiveness of ER vs FDs

- What does $a b \rightarrow c$ mean formulated in ER?

Maybe this?




Or this?



Pointed arrow also works!

Or even this?



- All of these ER-designs satisfy $a b \rightarrow c$, but differ in what is considered an independent entity

Practical use of FDs combined with ER

- FDs can be used to verify the correctness of an ER-design
 - Is the result in BCNF w.r.t. the dependencies you have identified?
 - Are the primary keys you identified sensible from your FDs?
 - If not there may be an error in your ER-translation or your understanding/modelling of the domain
- Sometimes FDs can be used to patch things up in your ER-design, particularly they are useful for finding secondary keys (UNIQUE constraints)
 - Every (minimal) key of each relation should be either the primary key or unique

Mining attributes (and FDs) from ER-design

- If you already have an ER-design, that may help you determine a useful set of attributes
- Looking at the relational schema is less helpful, because it contains multiple attributes that have different names but are conceptually the same (because of references)
- You can also extract some FDs by studying the diagram/schema, but that sort of misses the point of finding them since you will never find any FDs that can improve your design
 - We want to find FDs that express things our ER-design is missing
 - We should look for FDs in the domain description

Finding functional dependencies

- Determine all attributes
- Discover FD's either by looking at each attribute and ask "what do i need to know to determine this?" and by looking at each fact in the domain description and asking "does this express a dependency?"
- You can find multiple FDs determining the same attribute

Other normal forms

- There is a whole little hierarchy of normal forms

Highly simplified:

- 1NF: basically means "only has actual tables"
- 2NF: 1NF + no FDs from key attributes to non-key attributes
- 3NF: 2NF + no FDs between non-key attributes
- BCNF a.k.a. 3½NF: 3NF + attributes depend only on keys
- 4NF: BCNF + No violating MVDs
- 5NF, 6NF, DK/NF ...: Outside the scope of this course

I expect you to know how to normalize to BCNF and 4NF

1NF



2NF



3NF




BCNF



First normal form – 1NF

- First normal form means that table cells do not contain tables or other complex data structures

col1	stuff						
x	<table><tr><th>a</th><th>b</th></tr><tr><td>a1</td><td>b1</td></tr><tr><td>a2</td><td>b2</td></tr></table>	a	b	a1	b1	a2	b2
a	b						
a1	b1						
a2	b2						
y	<table><tr><th>a</th><th>b</th></tr><tr><td>a3</td><td>b3</td></tr></table>	a	b	a3	b3		
a	b						
a3	b3						



col1	a	b
x	a1	b1
x	a2	b2
y	a3	b3

Second normal form - 2NF

- A table is in second normal form if there is no FD from a set of prime attributes (attributes that are a proper subset of any key) to a non-prime attribute
- Example:

Suppose $R(a,b,c,d)$ has FDs $a \rightarrow b$ and $a \rightarrow c$

Here (a, b) is the only minimal key ($\{a,b\}^+ = R$)

That means a is a prime attribute and d is not, so $a \rightarrow d$ is a 2NF violation

Third normal form - 3NF

- 3NF has this condition on non-trivial FDs:
The LHS must be a superkey, unless the RHS is a prime attribute
- If you remove the "unless"-part, you get BCNF
 - BCNF is stronger than 3NF – some BCNF violations are not 3NF violations
- Example:

Suppose $R(a,b,c,d)$ has FDs $a b \rightarrow c$ and $c d \rightarrow a$

Both FDs are BCNF violations, but neither are 3NF violations since c and a are both prime attributes (since $\{a,b,d\}$ and $\{b,c,d\}$ are minimal keys)

Non-exhaustive checklist for the exam

- You have a general understanding of what an FD or an MVD claims about a relation

You know how to:

- Construct counterexamples to an FD
- Find a minimal basis F^- for a set of FDs F
- Compute the closure X^+ of an attribute set X
- Normalize to BCNF
- Check if an MVD holds and construct counterexamples
- Normalize to 4NF